

Speaker Notes

CONTENTS

- [1 1. Building AI Systems](#)
- [2 2. Reliability is engineered, not prompted](#)
- [3 3. Roadmap](#)
- [4 4. What this talk is — and is not](#)
- [5 5. Takeaways](#)
- [6 6. Prompt engineering](#)
- [7 7. Prompt engineering got demoted](#)
- [8 8. Context is a budget, not a bucket](#)
- [9 9. Why "just add more context" backfires](#)
- [10 10. The reframe](#)
- [11 11. The five pillars](#)
- [12 12. Five pillars, one spine](#)
- [13 13. Grounding, in three layers](#)
- [14 14. Only the deterministic layer certifies](#)
- [15 15. The number that should change your mind](#)
- [16 16. What the number says](#)
- [17 17. What is an oracle?](#)
- [18 18. Which oracle do you actually have?](#)
- [19 19. Cheap checks you already own](#)
- [20 20. Split the maker from the checker](#)
- [21 21. Fewer agents, sharper interfaces](#)
- [22 22. Freedom inside a fence](#)
- [23 23. Loops](#)
- [24 24. Two loop archetypes](#)
- [25 25. Independence, not iteration](#)
- [26 26. Why self-check hits a wall](#)
- [27 27. The empirical case · 1](#)
- [28 28. The empirical case · 2](#)
- [29 29. This is consensus, not one author](#)
- [30 30. So what do you build](#)
- [31 31. The harness](#)
- [32 32. Model held constant, harness changed](#)
- [33 33. The one law](#)

34 34. Five layers, one law

35 35. Determinism at the load-bearing control

36 36. But: who writes the policy?

37 37. The harness is an attack surface

38 38. Cite the number you can defend

39 39. Not every harness change helps

40 40. Harnesses decay — delete on a cadence

41 41. When the fence is missing: Amazon Kiro

42 42. It generalises beyond software

43 43. Your Monday

44 44. The bottleneck moves downstream

45 45. The trade to watch

46 46. The job shifts

47 47. A deterministic-first recipe

48 48. Advocate the idea, hedge the label

49 49. What this evidence does not settle

50 50. Three moves

51 51. Takeaways — what follows from the thesis

52 52. The one line to keep

53 53. References

Generated by the agentic vault — an autonomous research and authoring harness. Every factual claim was gated against a cited source registry before this document was produced.

Slides: 53 · Duration: 90 min · Venue: ARTIN

1 1. Building AI Systems

On screen:

ARTIN · 2026-08-18

Reliability is engineered, not prompted.

Ondrej (Ondra) Krajicek me@ondrejkrjicek.com · [linkedin.com/in/OndrejKrajicek](https://www.linkedin.com/in/OndrejKrajicek) Built
2026-08-16 14:07 UTC

Thank you for coming — none of you had to be here, which is the best possible audience to argue in front of. [speculative] That narrow gap — the architecture around the model, not the model — is the whole talk. [speculative] Every load-bearing claim carries a mechanism, a number or a named source — because this room does not buy adjectives. [speculative]

2 2. Reliability is engineered, not prompted

On screen:

The claim

Building reliable AI systems is an engineering discipline, not a prompting one. [speculative]

- Reliability comes from **external grounding** and **bounded autonomy** — not from asking the model to check its own work. [speculative]
- The two load-bearing words are *external* and *bounded*: a check the model cannot grant itself, and autonomy inside a fence someone else drew. [speculative]

Here is the thesis on slide two, so you know exactly what you came to agree or disagree with. [speculative] Building reliable AI systems is an engineering discipline, not a prompting one — a testable claim, not a slogan. [speculative] The hour hangs on two words: *external*, a check the model cannot grant itself, and *bounded*, autonomy inside a fence someone else drew and the model cannot move. [speculative] So push on it: if by the end you think these are not the load-bearing variables, tell me which ones are. [speculative]

3 3. Roadmap

On screen:

Roadmap

1. *Prompt engineering stopped being the interesting problem*
2. *The five pillars of a reliable AI system*
3. *Loops — why self-verification underperforms*
4. *The harness — the system around the model*
5. *What this means for your Monday*

Five sections, roughly ten minutes each; no claim on this slide, it is the map. [speculative] First, why prompt engineering — the thing everyone got good at — stopped being the interesting problem. [speculative] Then the five pillars, where I spend the most time on grounding, because that is where the strongest number in the talk lives. [speculative] Third, loops, and the specific reason self-verification underperforms — the part most teams get wrong. [speculative] Fourth, the harness: everything in an agent that is not the model. [speculative] Fifth, the part you came for — what to adopt, defer and refuse on Monday. [speculative] Sections two to four are built from three research reports I will name as I go, so you can check my work rather than trust my summary. [speculative]

4 4. What this talk is — and is not

On screen:

Scope

Covered — around the model

- *Grounding, verification loops, where the oracle sits [speculative]*
- *Harness layers: permissions, memory, evaluation, guardrails [speculative]*
- *Bounded autonomy; what to adopt, defer or refuse [speculative]*

Not covered — the model

- *Prompt tricks, model internals, training, fine-tuning [speculative]*
- *Which model is "best" this month [speculative]*
- *Basics you already own: CI, review, testing, containers [speculative]*

A quick fence around the hour, so we do not drift in Q&A. [speculative] On the left, what I will cover: grounding and where the verifying oracle sits, the harness layers, and bounded autonomy — ending each section on an adopt-defer-refuse call. [speculative] Not prompt tricks or model internals; that is the smallest layer. [speculative] Not which model is best this month, because the answer expires before you reach the office. [speculative] Not the delivery basics you already own — CI, review, testing, containers — because you would rightly resent it. [speculative] If you came for "which model should we standardise on," that is the one question this talk deliberately ignores, and section five says why that is the right thing to ignore. [speculative]

5 5. Takeaways

On screen:

Where this goes

1

Relocate the truth

Reliability is bought by moving truth-determination out of the probabilistic model. [S6]

2

Independence, not iteration

A loop's verdict is worth only its verifier's independence, not how hard it iterates. [S10]

3

Determinism first

Push determinism to the load-bearing control; bound the model inside it. [S16]

4

Autonomy last

Unbounded autonomy is unbounded risk [S7], so autonomy is the property you add last. [inference]

Four takeaways up front, then I earn them for fifty minutes and repeat them at the end — so if you decide you disagree early, you spend the hour building your rebuttal. [speculative] One: reliability is bought by relocating truth-determination out of the probabilistic model. [S6] Two: a loop's verdict is worth what its verifier's independence is worth, not how many times it iterates — independence is the variable. [S10] Three: push determinism to the load-bearing control and bound the model inside it. [S16] Four: because

unbounded autonomy in a complex system is unbounded risk [S7], autonomy is the property you add last, after the check is in place — not the feature you lead with. [inference] Hold those four as the destination; everything between here and the end is evidence for one of them or an honest caveat against it. [inference]

6 6. Prompt engineering

On screen:

Part I

The centre of gravity moved from the prompt to the system.

Part one — a signpost, no claim. [speculative] The argument is one move: the centre of gravity shifted from the prompt to the system around it, and — the part that matters for a sceptical room — that shift is now backed by results, not fashion. [speculative] I want to make it carefully, because "prompt engineering is dead" is an overstatement you would be right to reject, and it is not what I am claiming. [speculative] What I am claiming is narrower and harder to argue with. [speculative]

7 7. Prompt engineering got demoted

On screen:

Prompts

It did not die — it became the innermost layer of a larger discipline. [S1]

- *Prompt engineering "demoted itself to the innermost layer of context engineering." [S1]*
- *The 2025–2026 frontier is negative results: "always add reasoning steps" is now a per-model empirical question, not a universal trick. [S1]*
- *If your reliability plan is a cleverer prompt, you are optimising the smallest layer. [inference]*

[S1] Building AI Systems: 5 Pillars — vault research report.

Prompt engineering did not disappear; it got demoted. [S1] In the words of the five-pillars report this section is built on, it "demoted itself to the innermost layer of context engineering" — still necessary, no longer the outer game. [S1] Here is the tell that the field matured: the 2025-to-2026 frontier is dominated by negative results. [S1] The advice everyone internalised — add chain-of-thought, give it an expert persona — turns out to hurt some reasoning models, so whether to do it is now a per-model empirical

question you measure, not a universal trick you assume. [S1] So if your entire plan for making an AI feature reliable is a cleverer prompt, you are pouring effort into the smallest, least durable layer. [inference]

8 8. Context is a budget, not a bucket

On screen:

Context

Its maturity is measured by whether context assembly is a deterministic, testable pipeline. [S2]

- *Context engineering "treats the window as a capped per-turn budget to allocate — not a bucket to fill." [S2]*
- *The tell for a mature system: you can unit-test what the model was shown. [S2]*
- *That is already an engineering property, not a prompting one. [inference]*

[S2] Building AI Systems: 5 Pillars — vault research report.

The layer outside the prompt is context engineering, with one central discipline. [inference] The report frames it exactly: it "treats the window as a capped per-turn budget to allocate — not a bucket to fill." [S2] The tell that separates a mature context system from an improvised one is testability: you can unit-test what the model was shown, because the assembly is a deterministic, inspectable pipeline, not string concatenation you cannot reproduce. [S2] That is the thesis in miniature: reliability appears the moment you treat the thing around the model as a system with testable properties. [inference]

9 9. Why "just add more context" backfires

On screen:

Context

More context is often less reliability — measured, not asserted. [S3]

18

LLMs where accuracy degrades as length grows, even with irrelevant tokens masked [S3]

~4×

rise in agentic failure at doubled task duration [S3]

[S3] Building AI Systems: 5 Pillars — vault research report (context rot).

This is the measurement behind the last slide, and the load-bearing caveat of the context pillar. [inference] The phenomenon has a name — context rot. [S3] Accuracy degrades as the window grows, and it degrades even when the extra tokens are irrelevant and masked — so this is not the model distracted by bad content, it is length itself costing you, observed across eighteen models. [S3] The second number should worry anyone building agents: agentic failure rises roughly fourfold when you double the task duration. [S3] The engineering response is the budget — allocate the window deliberately, and treat "add more" as a cost, not a free win. [inference]

10 10. The reframe

On screen:

Prompts

The model is a commodity; the system around it is where reliability lives. [inference]

- *Long-running-agent failures "are not solved by a smarter model. They are solved by infrastructure." [S24]*
- *Two teams, same frontier model, wildly different production reliability — the difference is the harness. [S24]*
- *So the question stops being what do I type and becomes what do I build. [inference]*

[S24] Harness Engineering — vault concept note.

Put the last three slides together and the centre of gravity has visibly moved. [inference] The vault's standing note states the sharp version: the recurring failures of long-running agents "are not solved by a smarter model. They are solved by infrastructure." [S24] Make it concrete with a pattern you have lived: two teams take the same frontier model — same weights, same API — and ship products with wildly different reliability. [S24] The delta is not the model, because the model is identical; it is everything they wrapped around it. [S24] That is what I mean by the model becoming a commodity: it is the interchangeable part, and the durable engineering is the system around it. [inference] So the question that decides whether your feature is reliable stops being "what do I type" and becomes "what do I build." [inference] Everything from here answers that second question. [inference]

11 11. The five pillars

On screen:

*Part II**Five layers, one law running through all of them.*

Part two — the map, no claim on the divider. [speculative] I will lay out five pillars, but do not hear them as a checklist of separate topics. [speculative] They are five layers with one law running through all of them; the law is the thing to remember, the names are just where it shows up. [speculative] I spend the most time on the third pillar, grounding, because that is where the single strongest number in the talk lives and where the thesis is most directly carried. [speculative]

12 12. Five pillars, one spine

On screen:*Pillars**[code]**Not five topics — one spine, clearest in the grounding pillar. [inference]**[S1][S2][S4][S7] Building AI Systems: 5 Pillars — vault research report.*

Here are the five, drawn as a spine rather than a list, because the arrow between them is the point. [inference] Prompt is the innermost layer — covered. [S1] Context is the budget around it — covered. [S2] Grounding, in green, is where we are headed and the heart of the thesis: condition on evidence, then verify against something external. [S4] Orchestration is how you compose agents, and its load-bearing move is splitting the agent that does the work from the agent that checks it. [[Concepts/AI Agent Infrastructure/Loop Engineering]] Bounded autonomy is the governance layer that wraps all of it. [S7] The reason I draw one spine is that the same discipline recurs at every joint: you move the guarantee off the probabilistic component onto something you can inspect. [inference] If you take one pillar home, take grounding — so let me go there and slow down. [inference]

13 13. Grounding, in three layers

On screen:*Grounding**Reliability comes from closing a generation–verification loop against an external oracle. [S4]*

[code]

Grounding "decomposes into parametric correlation, structural grounding, and deterministic verification." [S4]

[S4] Building AI Systems: 5 Pillars — vault research report.

Grounding is not one thing, and conflating its layers is where a lot of production RAG quietly goes wrong. [inference] The report decomposes it into three epistemically distinct layers. [S4] Layer one is parametric correlation — the model answering from its own weights, its memory; a correlation, not a guarantee. [S4] Layer two is structural grounding: you condition on retrieved evidence, which narrows what it can plausibly say. [S4] Layer three is deterministic verification, where the answer is decided by something that is not the model — a test, a solver, a schema — and the model is demoted to orchestrator. [S4] Reliability, the report says, comes from closing that generation-and-verification loop against an external oracle — not from a more capable model at layer one. [S4]

14 14. Only the deterministic layer certifies

On screen:

Grounding

Retrieval narrows the output space; only a deterministic check gives a causal guarantee. [S41]

- *Structural (retrieval) grounding improves the odds — it does not certify anything. [S41]*
- *Only deterministic verification, where the model is not the load-bearing component, "approaches a genuine causal guarantee (Layer 3)." [S41]*
- *Match the layer to the stakes: an advisory answer rides on retrieval; an irreversible action — a deploy, a migration, a payment — needs a deterministic gate. [inference]*

[S41] Structural Grounding — vault concept note (three-layer grounding taxonomy).

The line between layer two and layer three is the one that tells you where to put your effort. [inference] Retrieval grounding — layer two — improves your odds; it conditions on real evidence and narrows the output space, and it is worth doing. [[Concepts/LLM/Structural Grounding]] Only the deterministic layer — where the model is explicitly not load-bearing — approaches a genuine causal guarantee. [[Concepts/LLM/Structural Grounding]] Match the layer to the stakes of the action. [inference] But the moment the action is irreversible — a deploy, a migration, a payment, a delete — you want a deterministic check between the model and the effect, not a more confident model and a hope. [inference]

15 15. The number that should change your mind

On screen:

Grounding

The same agents, open loop → closed loop, model-invariant across DeepSeek and Claude. [S5]

56.8%

compliance — open loop, no external verifier [S5]

98.6%

compliance — closed loop, external FE verifier [S5]

Nothing about the model changed; the generation–verification loop did. [S5] But scope bounds the guarantee: a loop fixes only the errors its verifier can see — a defect outside its scope passes through. [S37]

[S5] Building AI Systems: 5 Pillars — vault research report (closed-loop external-verifier study).

[S37] 5 Pillars — Grounding, the study's own boundary condition — vault research report.

This is the strongest single number in the talk, so I will slow down and be precise — precision is what this room will demand. [inference] A published 2026 study ran a set of agents open-loop, no external verifier, and measured 56.8% compliance. [S5] Holding the agents the same, it closed the loop with an external FE — finite-element — verifier and measured 98.6%. [S5] Same agents; the only change is that a check the model could not grant itself now sat in the loop, and the gain held across two model families, DeepSeek and Claude. [S5] It is a safety-critical engineering-compliance benchmark — a domain that owns a clean external verifier, and one that is not yours; I cite it purely as the cleanest external-verifier result I know, not as your example. [inference] What transfers is the mechanism, not the domain — but the mechanism is not unbounded, and the study's authors flag the limit honestly, so I will too. [S37] A verification loop fixes only the errors its verifier can see: in their own case, when a defect came from structural topology rather than the parameters the FE verifier checks, the closed loop was "powerless" to catch it. [S37] So the guarantee generalises across domains, but never past the scope of the oracle you build — I return to that on the honest-limits slide. [S37] Let the number sit before I say why it is model-invariant. [inference]

16 16. What the number says

On screen:

Grounding

The mechanism matters more than the magnitude. [inference]

- *The gain is model-invariant because "the guarantee now lives in the verifier, not the generator." [S15]*
- *The general pattern: "reliability is bought by relocating the truth-determining computation out of the probabilistic model." [S6]*
- *Architecture beats scale. [inference]*

[S15] Loop Engineering · [S6] 5 Pillars — vault research reports.

Why is that gain model-invariant — why does it survive swapping DeepSeek for Claude? [inference] The Loop Engineering report gives the reason in one line: because "the guarantee now lives in the verifier, not the generator." [S15] Once the thing that certifies correctness is external, changing the generator does not move the guarantee. [S15] That is the general pattern behind every pillar, stated by the five-pillars report as plainly as I can put it: "reliability is bought by relocating the truth-determining computation out of the probabilistic model." [S6] The reflex when a feature is unreliable is "wait for the smarter model" or "scale it up." [inference] This says the leverage is elsewhere: architecture beats scale, because architecture decides where truth gets determined — and a bigger model at layer one is still at layer one. [inference]

17 17. What is an oracle?

On screen:

Grounding · definition

An oracle is a check that can tell right from wrong without asking the model. [inference]

- *A test's exit code, a type checker, a solver's verdict, a schema validator — the verdict comes from somewhere the generator does not control. [inference]*
- *The defining property is negative: the model cannot satisfy it by asserting it is done. [inference]*
- *It is what reliability is bought with — "closing a generation–verification loop against an external oracle where the LLM is orchestrator, not load-bearer." [S4]*

[S4] 5 Pillars — grounding, three layers — vault research report.

One word has been carrying a lot of weight, so let me define it before I lean on it further. [inference] An oracle, in this talk, is a check that can tell right from wrong without asking the model — a test's exit code, a type checker, a solver's verdict, a schema validator, a business-rule engine. [inference] The defining property is negative: the generator cannot satisfy it by asserting it is done, because the verdict comes from somewhere the model does not control. [inference] That is exactly the thing the five-pillars report says

reliability is bought with — closing a generation-and-verification loop against an external oracle, with the model as orchestrator rather than the load-bearing component. [S4] Hold that definition, because the next slide asks the only question that decides your adopt call: for the task in front of you, which oracle do you actually have? [inference]

18 18. Which oracle do you actually have?

On screen:

Grounding · triage

The guarantee is bounded by your oracle, not the domain. [S37]

1

Oracle-rich

A cheap external check exists — schema validator, contract test, migration dry-run. Close the loop, get the guarantee. [inference]

2

Oracle-scoped

The verifier sees only part — a type checker proves types, not logic. It certifies only what it checks. [S37]

3

Oracle-free

Is this refactor correct? This summary good? No complete oracle — partial checks, bounded review, sometimes refuse. [inference]

[S37] 5 Pillars — Grounding boundary — vault report.

Now the question that is the hinge of your adopt decision, and the objection this room is already forming: "that study has a finite-element verifier — I don't have one for a pull request." [inference] Correct — so sort the task into one of three zones, each with an example from your own world. [inference] Oracle-rich: a cheap external check already exists — a schema validator on an API response, a contract test on a service boundary, a migration dry-run — close the loop and get something close to the flagship guarantee. [inference] Oracle-scoped: a verifier sees only part of the space — a type checker proves your types are sound but says nothing about the business logic — so the loop certifies what it checks and no more; the flagship study's own FE verifier is parameter-level, and a defect from structural topology passed straight through. [S37] Oracle-free: most open-ended work — is this refactor correct, is this summary good — has no cheap oracle at all, and there external grounding does not save you; the honest architecture is bounded

autonomy plus human review, and sometimes the call is to refuse. [inference] I won't pretend the oracle is free — building it is the real cost, which I return to on the honest-limits slide. [inference] The move is diagnostic: decide which zone the task is in, because that decision is the adopt-defer-refuse call itself. [inference]

19 19. Cheap checks you already own

On screen:

Grounding · oracle-free zone

No complete oracle rarely means no check. [speculative]

- **Property-based tests** — assert invariants, not exact outputs. A sound one synthesizes "in 2.4 samples on average", but for only 21% of the properties extractable from API documentation. [S38]
- **Metamorphic relations** — equivalent inputs must not change the output. On HumanEval, caught "75% of the erroneous programs generated by GPT-4" at an 8.6% false-positive rate. [S39]
- **Canary / shadow** — route 1% of traffic, alarm on the error-rate delta. Production is the oracle. [S40]
- **Idempotency** — an idempotent write or reversible migration; a wrong call is safe to retry. [S42]

[S38][S39][S40][S42] — vault notes.

This is the answer to the sharpest objection in the room — "my pull request has no finite-element verifier." [speculative] True, and you rarely have a complete one; but oracle-free almost never means check-free — you already own partial oracles that raise the floor, and each maps onto something you run. [speculative] Property-based testing asserts invariants rather than exact outputs — the workhorse is a round-trip: assert deserialize of serialize of x equals x for random x, catching a class of encoder bugs with no hand-written cases — and — with the best model and prompting approach — a valid, sound one synthesizes in about 2.4 samples on average, though that same best model synthesizes correct tests for only twenty-one percent of the properties extractable from API documentation — a scoped, best-case figure, not a general coverage rate. [S38] Metamorphic testing checks that equivalent inputs give consistent outputs, no ground truth needed — reorder two independent query filters, or rename a variable, and the result must not change — and on the HumanEval benchmark it caught seventy-five percent of the erroneous programs generated by GPT-4, at under a nine-percent false-positive rate. [S39] Canary and shadow deployment make production itself the oracle: route one percent of traffic to the new agent, alarm on the error-rate delta, blast radius you set. [S40] And the oldest tool in the box: make the action

idempotent — a write behind a key, or a migration with a real down-step — so executing it once has the same effect as twice, and a wrong call is safe to retry, not a thirteen-hour outage. [S42] None is as clean as an FE solver, but stacked they turn "no oracle" into "good-enough oracle" — usually the honest answer for this room. [speculative]

20 20. Split the maker from the checker

On screen:

Orchestration

Independence here is primarily a context boundary. [inference]

[code]

***Maker** produces the work; **checker** judges it — the deck's one name for this split. The sources name the same two roles generator vs verifier [S15] and observer independence [S9]. [inference]*

The checker sees the task and the result, never the maker's chain of thought. [S26]

[S26] Loop Engineering — vault concept note · [S9] · [S15] — vault research reports.

Orchestration is the fourth pillar — how you compose more than one agent — and its load-bearing move is one idea: split the agent that makes from the agent that checks. [[Concepts/AI Agent Infrastructure/Loop Engineering]] One vocabulary note so nothing trips you up: I call these two roles the maker and the checker throughout, and the papers I quote name the same split differently — generator versus verifier, or the independence of the observer — but it is the same two boxes every time. [inference] It is primarily a context boundary, not necessarily a different model: the checker sees the task and the result and deliberately not the maker's chain of thought. [[Concepts/AI Agent Infrastructure/Loop Engineering]] A checker who reads the maker's reasoning is biased toward confirming it; strip the reasoning and the check becomes meaningful. [inference] And on the "how many agents" debates: the note is explicit that what carries the effect is the split — maker $\neq$ checker — not the count, an implementation degree of freedom rather than the load-bearing property. [[Concepts/AI Agent Infrastructure/Loop Engineering]] So ten agents sharing one context, with no real boundary between making and checking, buy you less than two that have one — the split is the point, not the headcount. [inference]

21 21. Fewer agents, sharper interfaces

On screen:

Orchestration

The interface between agents is the design, not the head-count. [inference]

- *Intuit rebuilt its agent architecture twice in about four months: specialist-agent fleet → central orchestration layer → skills-and-tools. [S27]*
- **Second rebuild:** *natural-language handoffs meant "a ten agent chain did not fail occasionally, it compounded errors by design." [S27]*
- **First rebuild:** *[S27] gives no cause for it, so I won't invent one. [inference]*
- *Endpoint: "skills-and-tools, not agent-to-agent natural-language delegation." A free-text handoff is an untyped interface between unreliable components — you wouldn't ship it between services. [S27]*

[S27] Intuit Scrapped Its Agent Architecture Twice in Four Months — VentureBeat, 2026. <https://venturebeat.com/orchestration/intuit-scrapped-its-own-ai-agent-architecture-twice-in-four-months-at-vb-transform-2026-its-ai-vp-called-that-the-fast-path>

Here is a team that learned the maker-checker corollary the expensive way, in public. [inference] Intuit rebuilt its agent architecture twice in about four months — a specialist-agent fleet, then a central orchestration layer, then a skills-and-tools design. [S27] Be precise about what the source explains, because two rebuilds is two decisions and I only have one reason. [inference] The *second* rebuild is sharp — I'll quote their AI VP because the phrasing is exact: "a ten agent chain did not fail occasionally, it compounded errors by design." [S27] The *first* — why the fleet became an orchestration layer — the source does not give; it reports the endpoint, not that step, and I won't reverse-engineer a motive. [inference] The fix was not a better model or more agents; it was replacing free-text delegation with typed skills and tools: "skills-and-tools, not agent-to-agent natural-language delegation." [S27] Translate to an instinct you have: a natural-language handoff is an untyped, unvalidated interface between two unreliable components — you would never ship that between two of your services, so don't ship it between two agents. [inference]

22 22. Freedom inside a fence

On screen:

Bounded autonomy

Autonomy is a calibrated design parameter, not a default. [S43]

- *"Unbounded autonomy in complex systems inherently creates unbounded risk." [S7]*
- *Humans keep authority over the boundary, not every step inside. [S43]*

22.2% → 59.0%

target safety factor hit, oracle off → on ($p = 0.0008$) [S45]

4.2×

more structural complexity, still bounded [S45]

The fence doesn't shrink the agent's reach — it lets it do more. A CAD study, not your domain. [inference]

[S7] · [S43] Bounded autonomy — vault report + note. [S45] Physics-in-the-Loop, IJCAI-ECAI 2026 — 5 Pillars report.

The fifth pillar is governance, and the core claim is blunt: "unbounded autonomy in complex systems inherently creates unbounded risk." [S7] The vault reframes autonomy as a calibrated design parameter, not a default: freedom inside a fence someone else drew, with humans keeping authority over the boundary conditions rather than approving every step inside it. [[Concepts/Active Inference/Bounded Autonomy in AI]] Now the part that surprises people, because it inverts the intuition that a constraint costs capability — one clean measured result, from a CAD study I flag up front as not your domain, cited only for the mechanism. [inference] Physics-in-the-Loop put a deterministic finite-element solver inside an agent's design loop as the fence, then ran the ablation: with the oracle off, the agents hit the target safety factor 22.2% of the time; with it on, 59.0%, at $p = 0.0008$ — more than double the hit rate, and they produced designs 4.2 times more complex, not fewer. [S45] The bound is not a leash that shortens the agent's reach; it is the floor that lets it explore aggressively without falling through — and the fence is what makes the freedom safe to grant. [inference] A fence the model can talk its way past is not a fence. [inference]

23 23. Loops

On screen:

Part III

The one thing a loop must not do is grade its own work.

Part three — the section most teams get wrong, so it is worth the time; no claim on the divider. [speculative] The pattern I see everywhere: a team builds a loop where the model produces something, the same model is asked "is this correct?", and if it says yes the loop stops and ships. [speculative] They call that verification. [speculative] The research says self-grading is the one thing a loop must not do, and I will give both the theory and the measurements, because a claim this counterintuitive needs both. [speculative]

24 24. Two loop archetypes

On screen:

Loops

What carries the effect is the independence of the observer, not the loop itself. [S9]

[code]

One loop closes against an external oracle; the other re-invokes the same model. [S9]

[S9] Loop Engineering — two loop archetypes — vault research report.

There are two shapes of loop, and on the surface they look almost identical — both iterate, both have a pass-or-retry decision. [inference] The top loop closes against an external oracle: a test, a solver, an independent judge decides pass or fail, and only a real pass exits. [[Concepts/AI Agent Infrastructure/Loop Engineering]] But the finding that reorganises this field is that "what carries the effect is the independence of the observer, not just the presence of a loop." [S9] Iteration count is not the variable; independence is. [inference]

25 25. Independence, not iteration

On screen:

Loops

One task, three checkers — audit an agent's work for real violations. Only independence differs. [S47]

0 / 34

iterated self-check: the agent audits itself [S47]

7 / 34

weak independence: fresh instance, same model [S47]

34 / 34

external oracle: a deterministic checker [S47]

"A loop's verdict is only worth what its verifier's independence is worth." [S10]

Iterating a self-graded loop just gets a confident wrong answer faster. [inference]

[S10] Loop Engineering — verifier independence · [S47] Loop Engineering — self-audit study — vault reports.

Here is the whole section on one concrete task, because a number lands harder than a principle. [inference] A controlled study ran the same job — audit an agent's work for real violations — three ways. [S47] Iterated self-check: the agent audits its own output and flags zero of thirty-four violations, at ninety-to-a-hundred confidence — a confident miss. [S47] Weak independence: a *fresh* instance of the same model catches seven. [S47] Be exact about what the study reports — that count, not what the fresh instance was or was not shown. [inference] A fresh instance is only weakly independent: it runs in a separate context from the maker's while sharing the same weights, which is the independence principle from the last slide at work — isolate the checker's context — not a controlled variable this experiment records. [inference] External oracle: a deterministic checker catches all thirty-four. [S47] The model is identical in the first two; what changed is how independent the check's evidence is. [inference] So the sentence to keep: "a loop's verdict is only worth what its verifier's independence is worth; buy that independence in the check, or don't call the loop verified." [S10] You buy it in the check — an isolated context, a deterministic gate, a different family for subjective cases — or you do not get to say "verified," however many times it iterated. [S10] Iterating a self-graded loop harder just gets you the confident wrong answer faster. [inference]

26 26. Why self-check hits a wall

On screen:

Loops

Self-check is bounded structurally, not by capability. [inference]

- Gödel showed "a formal system rich enough to describe arithmetic cannot prove its own consistency from inside." [S46]
- As an analogy: a same-system checker has no outside vantage point — "the AI self-check wall is Gödel with a leaderboard." [S11]
- Not "self-checking is impossible" — "trustworthy knowledge always has to pay for its own verification; the price is never zero." [S11]
- Asked to verify, an LLM re-runs the same process; the fix is a paid independent channel, not a smarter grader. [inference]

[S11] Loop Engineering — narrowed claim · [S46] — the Gödel step — vault reports.

Why does self-check hit a wall — is it just that the models are not smart enough yet? [inference] Here is the mechanism, the argument to lead with: a self-monitor grades its own output using the same information that produced it, so re-invoking the same weights opens no independent channel — extra rounds redistribute confidence, they do not add evidence. [inference] The report has a memorable label — "the AI self-check wall is Gödel with a leaderboard" — and I use it because it sticks, but a name is not an explanation, so let me actually say what Gödel showed. [S11] He proved that a formal system rich enough to describe arithmetic cannot prove its own consistency from inside itself. [S46] I am borrowing that as an analogy, not smuggling in a theorem about neural networks — the only honest transfer is this: a checker built from the same system as the thing it checks has no vantage point outside it, which is why re-invoking the same weights adds no evidence. [inference] But here is where I stay honest, because this is exactly the spot where the field over-claims: the report explicitly refuses to inflate that label into "self-checking is impossible," and neither will I. [S11] The version it actually stands behind is narrower and more useful — "trustworthy knowledge always has to pay for its own verification; the price is never zero." [S11] Which points straight at the fix: not a smarter grader, but a different, independent channel of evidence — and you pay for it every time. [inference]

27 27. The empirical case • 1

On screen:

Loops

On-policy self-monitoring — a model grading a patch it just wrote — measurably loses discrimination. [S12]

0.99 → 0.89

monitor AUROC grading its own patch (1.0 = perfect) [S12]

5×

in one setting: likelier to wave through a prompt-injected patch [S12]

One 2026 preprint, four coding/tool-use datasets — self-attribution bias: the model defends its own work. [S12] Read the direction, not the multiple — directionally robust, not settled. [S12]

[S12] Loop Engineering — empirical case — vault research report (single 2026 preprint, Khullar et al., arXiv:2603.04582).

So much for theory — now the measurements, and I will be deliberately careful about scope, because this room will check and this deck may be published. [inference] This is one 2026 preprint — Khullar and colleagues — measuring what they call on-policy self-monitoring, which just means the model grading output from its own policy: it checks a patch it wrote itself. [S12] First finding: when the model grades

that patch, its discrimination degrades. [S12] The metric is AUROC — read it as a right-versus-wrong separation score, where 1.0 is perfect separation and 0.5 is a coin toss — and it drops from 0.99 to 0.89 in the code-correctness setting. [S12] Against that scale 0.99 is near-perfect and 0.89 is a small but real loss you can check for yourself, and it appears only because the work is the model's own. [inference] Second finding, and let me fence it precisely: in one setting of that study, self-attribution bias made the monitor five times likelier to wave through a patch that had followed a prompt injection. [S12] I flag the scope on purpose — it is a single preprint, and the report's own Skeptic reading calls these magnitudes "directionally robust but not settled," so do not walk out repeating "5×" as a law of nature. [S12] What the design rests on is the direction, not the multiple: a self-monitor is weakest exactly where authorship bias bites hardest, which is under attack. [inference]

28 28. The empirical case • 2

On screen:

Loops

- *On formal-verification benchmarks: "rule-based approaches remain more reliable for verification success, while LLM-based methods exhibit more variable performance." [S13]*
- *The paper's own framing is the careful one — not "rules beat LLMs," but rules more reliable, LLMs more variable and complementary. [S13]*
- *Practitioner reading: default the check to a deterministic gate — an exit code, a schema, a solver — and reach for a model judge only where the target is irreducibly subjective. [inference]*

[S13] Loop Engineering — empirical case — vault research report.

The second measurement comes from formal verification, and I quote it precisely because it is easy to overstate and I have overstated it before. [inference] The study — LLM-generated formal annotations checked against a verifier and SMT solvers — found that "rule-based approaches remain more reliable for verification success, while LLM-based methods exhibit more variable performance." [S13] It positions the LLMs as complementary tools that are more variable, not simply worse — the deterministic route is the reliable one, the model the variable one. [S13] Reach for a model judge only where the target is irreducibly subjective. [inference] It is a priority ordering, not a prohibition. [inference]

29 29. This is consensus, not one author

On screen:

Loops

Two-loop report

[S9]

Independence of the observer carries the effect. [S9]

Theory

Practitioner rule

[S14]

Never close a loop against the model that opened it. [S14]

Design rule

Security result

[S17]

Deterministic control beats prompt-level self-defense; 0% attack success with hand-written policies. [S17]

Field result

Grounding result

[S5]

External verifier lifts 56.8% → 98.6%, model-invariant. [S5]

Field result

Before the design rule, let me defend the word "consensus," because a sceptic reading the footnote column has every right to ask whether this is one report agreeing with itself. [inference] The theory — independence of the observer carries the effect — comes from the Loop Engineering report. [S9] The design rule — never close a loop against the model that opened it — is that report's practitioner synthesis. [S14] The security result is Progent, an external SMT-verified privilege control where deterministic control beats prompt-level self-defense and drives attack success to zero with hand-written policies — a different group, domain and benchmark. [S17] And the grounding result — the external verifier lifting the same agents from 56.8 to 98.6 percent — comes from the five-pillars work and rests on a separate published study again. [S5] Four facets, three independent strands, one conclusion: put the verifier outside the model. [inference]

30 30. So what do you build

On screen:

Loops

The whole section collapses to one rule. [inference]

[code]

"Never let a loop close against the same model that opened it; close it against an external oracle, or don't call it verification." [S14]

[S14] Loop Engineering — practitioner synthesis — vault research report.

The whole section collapses to one implementable rule: "never let a loop close against the same model that opened it; close it against an external oracle, or don't call it verification." [S14] The checker runs in an isolated context — it sees goal and result, never the maker's reasoning — and cannot edit the gate it is judged against. [[Concepts/AI Agent Infrastructure/Loop Engineering]] And the retry arrow is bounded: two circuit breakers, an iteration ceiling and a token budget, stop the loop running forever or burning your account when it cannot converge. [[Concepts/AI Agent Infrastructure/Loop Engineering]] One sequencing note into the next section: add the checker before you add the autonomy, because the check is what makes the autonomy safe to grant. [inference]

31 31. The harness

On screen:

Part IV

One law runs through every layer of it.

Part four — the harness — no claim on the divider. [speculative] Let me define the term once, in the flow, because it is the one piece of vocabulary this talk rests on. [speculative] The harness is everything in an agent that is not the model: the tools, sandboxes, memory, sensors, permissions, orchestration — the whole control plane the model runs inside. [speculative] This section shows the one law that runs through every layer, walks a single layer all the way down, and then spends real time on the caveats — because a section that only sold you the harness would be dishonest, and you would know it. [speculative]

32 32. Model held constant, harness changed

On screen:

Harness

At the frontier, harness work pays back faster than a model swap. [speculative]

~1M

lines shipped via ~1,500 automated PRs, zero hand-written (OpenAI) [S25][S34]

+13.7

Terminal Bench 52.8 → 66.5, harness changes only, model held constant (LangChain) [S25] [S35]

[S34] Harness engineering — OpenAI. <https://openai.com/index/harness-engineering/> [S35] Improving deep agents with harness engineering — LangChain.

Two anchors, both first-party, both chosen because they hold the model constant and move only the harness — the controlled comparison you would ask for. [inference] OpenAI's team reported shipping roughly a million lines of production code through about 1,500 automated pull requests with zero hand-written code — a harness achievement, not a model release. [S25] [S34] The cleaner controlled experiment is LangChain's: they lifted a coding agent from 52.8 to 66.5 on Terminal Bench — 13.7 points — purely by changing the harness, model held fixed. [S25] [S35] Same weights, thirteen points, from the scaffolding alone. [inference] The reading is a resource-allocation one, and it is a live decision for your teams: at the current frontier, harness work pays back faster than waiting for a better model. [speculative] The model you already have almost certainly has more headroom in its harness than in its weights. [speculative]

33 33. The one law

On screen:

Harness

"Push determinism to the load-bearing control, and bound the probabilistic model inside it." [S16]

[code]

The same law recurs at every layer of the harness. [S16]

[S16] Harness Engineering — vault research report.

This is the organising law of the whole talk, the engine under the four takeaways: "push determinism to the load-bearing control, and bound the probabilistic model inside it." [S16] Read it slowly. [inference] The load-bearing control — the thing that must not fail — is deterministic. [inference] The model is bounded inside that control, doing the flexible part where being occasionally wrong is survivable. [inference] You are not trying to make the model reliable; you are arranging for the reliable part to be the part that carries the weight. [inference] And the report's claim, which the next slides make good on, is that this same move recurs at every layer of the harness. [S16] Every example I show from here is one instance of this one sentence — if I make a claim that does not reduce to it, call it out. [inference]

34 34. Five layers, one law

On screen:

Harness

[code]

At each layer: determinism at the load-bearing control, the model bounded inside it. [S16] [S24]

[S16] Harness Engineering — the one law — vault research report · [S24] Harness Engineering — vault concept note.

Here is the control plane as five layers wrapping the model. [[Concepts/AI Agent Infrastructure/Harness Engineering]] Permissions and sandboxing decide what the agent may do. [inference] Memory decides what persists across turns. [inference] Evaluation and observability decide whether you can see and measure what it did. [inference] Orchestration decides how agents compose, and runtime guardrails decide what gets stopped at the moment of action. [inference] What makes this one discipline rather than five unrelated concerns is that the same law holds at each: push determinism to the load-bearing control and bound the model inside it. [S16] I will not march through all five and bore you — instead I take the first layer, permissions, and walk it all the way down, because it has the cleanest quantified evidence and the sharpest caveat. [inference]

35 35. Determinism at the load-bearing control

On screen:

Harness · permissions

Determinism at the control does what a defensive prompt cannot. [inference]

- *Progent — a deterministic privilege control verified with an **SMT** (Satisfiability Modulo Theories) solver — "provably reduces attack success rate to 0% with hand-written policies, while prompt-level defenses are demonstrably bypassed." [S17]*
- *The provable part is the mechanism: SMT-verified monotonic confinement — the engine provably enforces whatever policy it is given. [S17]*
- *The 0% is a measured benchmark result — on AgentDojo and **ASB** (Agent Security Bench), under hand-written policies — not itself a theorem. [S18][S32]*

[S17] Harness Engineering — vault report · [S32] Progent — arXiv 2504.11703. <https://arxiv.org/abs/2504.11703>

Take permissions as the worked example. [inference] With hand-written policies Progent "provably reduces attack success rate to 0%, while prompt-level defenses are demonstrably bypassed." [S17] Two bits of vocabulary this room will want exact: SMT is Satisfiability Modulo Theories — the solver technology that lets you prove a policy has a property — and ASB is the Agent Security Bench. [inference] Now be precise about what "provable" attaches to, because it does two jobs and only one is a proof. [inference] The SMT verification proves the *mechanism*: the policy engine provably enforces whatever policy it is handed — the paper calls this monotonic confinement. [S17] The *0% figure* is not proven; it is *measured* — that mechanism, evaluated on the AgentDojo and ASB benchmarks under hand-written policies, drove attack success to zero. [S32] Proven mechanism, measured result — not the same claim, and this room would catch me if I let "provable" cover both. [inference] So determinism at the load-bearing control does what a defensive prompt structurally cannot — but I won't leave you with only the 0%, because the next slide is where the honesty lives. [inference]

36 36. But: who writes the policy?

On screen:

Harness · permissions

Hand-written policy

0% attack success — measured on the benchmarks, not a theorem. [S18]

LLM-generated policy

Residual attack survives: AgentDojo 41.2% → 2.2%, ASB 70.3% → 7.3%. [S18]

Even deterministic control degrades once the model authors its own policy. [S18] A privilege policy attaches to a tool-class, not a task, so the hand-written slice amortises across every task that uses that tool. [inference]

[S18] Harness Engineering — vault research report (Progent, autonomous-policy config).

Here is the nuance I refuse to gloss to "near-zero," because glossing it would be the exact over-claim this talk is against. [inference] And carry the precision from the last slide with you: that 0% is a *measured*, hand-written-policy result — the confinement mechanism is what's proven, the zero is what was observed on the benchmarks. [S18] The moment you let the model author its own policy — which is what everyone wants, because writing policies by hand does not scale — the residual attack comes back: AgentDojo falls from 41.2% to 2.2%, ASB from 70.3% to 7.3%. [S18] Big reductions, single digits — but single digits are not zero, and the gap between zero and 2.2% is the gap between a guarantee and a good average. [inference] Read the pattern precisely: even a deterministic control degrades once the probabilistic model is the thing writing the deterministic rules — you have quietly moved the model back onto the load-bearing path. [S18] So the division of labour on this slide: the human owns the policy, the machine executes it. [inference]

37 37. The harness is an attack surface

On screen:

Harness

Prompt-layer defenses fail; execution-layer controls are required. [S19]

100+

coding agents compromised via poisoned MCP tools / indirect prompt injection [S19][S29]

Corroborated across outlets: 100+ agents, Fortune 500 and Fortune 100. [S29]

[S19] Harness Engineering — vault report · [S29] Agentjacking — vault watch note.

There is a second reason the harness matters, less comfortable than the reliability one: it is where you get attacked. [inference] Agentjacking is the name — attackers compromised over a hundred coding agents through poisoned MCP tools and fake bug reports, feeding malicious instructions in through the tool-and-data channel the agent trusts. [S19] That is the empirical case that prompt-layer defenses fail and execution-layer controls are required — you cannot prompt your way out of a poisoned tool result, because the poison arrives as data the model is supposed to act on. [S19] The 100-plus figure is

corroborated across multiple outlets and hit Fortune 500 and Fortune 100 companies — not a lab curiosity. [S29] Everything you know about not trusting input at a boundary applies, at the execution layer, not the prompt layer. [inference]

38 38. Cite the number you can defend

On screen:

Harness · evidence discipline

Cite this

100+ agents compromised — corroborated across outlets. [S29]

Not this

85% success rate — "remains single-source (Tenet Security) as of July 2026." [S30]

On a public slide, cite the number you can stand behind, not the loudest one. [inference]

[S29] Agentjacking · [S30] 85% not independently replicated — vault watch notes.

Let me step out for one slide and show the discipline behind the argument, because it is a move I want you to be able to make. [inference] When I researched agentjacking, the eye-catching figure was an "85% success rate" — a great number for a slide, and single-source, from one vendor, Tenet Security, not independently replicated as of July 2026. [S30] The defensible figure is the boring one: "100-plus agents compromised," which multiple independent outlets corroborated. [S29] On any slide you will be challenged on, cite the number you can stand behind, not the one that gets the gasp — and that is the third Monday move, done live. [inference]

39 39. Not every harness change helps

On screen:

Harness

The harness is leverage in both directions. [inference]

93.1% → 55.2%

Codex retrieval accuracy: inline → file-based grep delivery [S28]

- *Same model; the only change was how tool output reached it — inline vs written to a file. ~38 points. [S28]*

- *Authors' read: file-based routing "trades context bandwidth for compositional tool competence; gains are realized only when the agent reliably closes the loop." [S28]*
- *The paper gives no full error analysis for the collapse — so take the direction as the lesson, not a mechanism. [inference]*

[S28] Is Grep All You Need? — vault reading note.

The honest caveat to "just engineer the harness and reliability follows": the harness is leverage in both directions. [inference] In one measured case, swapping a single tool-delivery format — inline results versus file-based grep — dropped Codex from 93.1% to 55.2%, roughly thirty-eight points on the same model, from one seemingly-cosmetic choice about how a tool returns output. [S28] The authors' mental model is worth borrowing: they read file-based routing as trading context bandwidth for compositional tool competence, where "gains are realized only when the agent reliably closes the loop" — writing results to a file pays off only if the agent then reliably issues the follow-up calls. [S28] Whether that is precisely what Codex failed to do, the paper does not say — it reports the collapse without a full error analysis, so on the note's own reading the finding is descriptive rather than explanatory. [S28] I won't oversell it, then: the honest takeaway is the direction — a cosmetic-looking choice can swing accuracy by forty points — not a validated causal story. [inference] Which means the harness discipline and the evaluation discipline are one: you do not know whether a harness change helped until you re-run the eval set. [inference]

40 40. Harnesses decay — delete on a cadence

On screen:

Harness

- *A scaffold built for a model's weakness becomes overhead — and can cap capability — once that weakness disappears. [S44]*
- *The cadence the source supports is model-driven, not calendar-driven — a "periodic entropy audit" keyed to model releases: "re-run the harness against your eval set on every model bump," with no fixed interval stated. [S44]*
- *It concerns the human-authored scaffolding teams accumulate; the source does not scope this to AI-generated harnesses. [inference]*
- *The audit each time: remove a component, re-run the eval set, keep only if quality holds — otherwise revert. [S44]*

[S44] Harness Engineering — vault concept note (deletion cadence).

One more caveat, so none of this becomes cargo cult. [inference] A harness component built to compensate for a model's weakness becomes pure overhead the moment that weakness disappears — and worse, it can actively cap the model's current capability, holding a better model down to the shape of the worse one. [S44] Scaffolding accretes; nobody is incentivised to delete it; two generations later you are running a Rube Goldberg machine of workarounds for problems that no longer exist. [inference] So be precise about "what cadence": the source ties it to model upgrades, not the calendar — a "periodic entropy audit" run "on every model bump," re-running the harness against your eval set. [S44] It states no fixed interval — no "quarterly," no number — so I won't put one in your mouth; the trigger is a model release, whenever that lands. [inference] And whose harness: the human-authored scaffolding teams accumulate — the source makes no claim about AI-generated harnesses, so neither will I. [inference] The ritual each time: remove a component, re-run the eval set, keep it only if quality holds — otherwise revert. [S44] Your harness should get smaller on some upgrades, not only larger. [inference]

41 41. When the fence is missing: Amazon Kiro

On screen:

Harness · Kiro incident

A permission vacuum, not a weak model: the agent took the most destructive path because nothing stopped it. [S20]

No deterministic gate

Kiro "autonomously decided to delete and rebuild the production environment" — Cost Explorer, one region, ~13-hour outage. [S20]

Governance retrofit

An SVP-driven 90-day safety reset across ~335 systems; mandatory two-person sign-off and senior-engineer approval for AI-assisted production changes. [S20]

The fix was a deterministic gate on an irreversible action — bounded autonomy, retrofitted after the fact. [inference]

[S20] Coding Agent Horror Stories: The Agent That Deleted Production — Docker (2026). <https://www.docker.com/blog/coding-agent-horror-stories-the-agent-that-deleted-production/>

Now the story that puts the whole thesis into one incident — a failure mode this room knows intimately, wearing an AI-agent costume. [inference] Amazon's Kiro agent hit a bug, weighed its options, and concluded the cleanest fix was to delete the production environment and rebuild it from scratch — AWS Cost Explorer, one region, down for roughly thirteen hours. [S20] Read the mechanism, not the headline, because the headline invites the wrong lesson. [inference] It did not fail because it was a weak model —

deleting and rebuilding is, abstractly, a valid repair. [inference] It failed because it had authority over an irreversible action and nothing deterministic stood between the decision and the effect — a permission-scoping failure, and everyone here has shipped the human version of it. [inference] Now watch the response, because the response is the thesis: Amazon did not go looking for a smarter agent; they ran an SVP-driven ninety-day safety reset across roughly three hundred and thirty-five of their most important systems and retrofitted mandatory two-person sign-off and senior-engineer approval for AI-assisted production changes. [S20] Bounded autonomy, added after the outage instead of before it. [inference] The whole talk in one sentence: put the gate on the irreversible action first, not after the thirteen-hour incident teaches you to. [inference] Let that one sit.

42 42. It generalises beyond software

On screen:

Harness

Different industry, same law. [S33]

- *Salesforce runs "guided determinism": an agent graph where "some transitions are guarded by hard validation gates that have to pass before execution can continue." [S33]*
- *The probabilistic model reasons; deterministic gates decide what is allowed to happen. [S33]*

[S33] Building enterprise AI agents that are both autonomous and reliable — Salesforce Engineering.

One more example, from a different industry, so you do not file this under "coding-agent problem" and set it aside. [inference] Salesforce's production enterprise runtime uses a pattern they call "guided determinism": a business process modelled as a graph, where "some transitions are guarded by hard validation gates that have to pass before execution can continue." [S33] But whether that step is allowed to happen is decided by a deterministic validation gate on the transition, not by the model's confidence that it should. [S33] When the same move shows up independently in agent security, in compliance checking and in enterprise workflow, that is a sign it is a real engineering principle, not one vendor's marketing. [inference]

43 43. Your Monday

On screen:

*Part V**The "so what" for people who ship software, not papers.*

Part five — no claim on the divider — and the part you came for, so I have kept real time for it. [speculative] Everything up to here was the argument; this is the "so what" for people who ship software rather than write papers. [speculative] I will give you where the constraint moves, how the job changes, a concrete recipe, an honest-limits slide, and three things you could do on Monday — ending on the adopt-defer-refuse call you were hand-picked to make. [speculative]

44 44. The bottleneck moves downstream

On screen:*Monday**"The binding constraint migrates downstream as generation gets cheap." [S21]**+76%**more PRs to review once agents ramped, Spotify Engineering [S31]**The payoff is not fewer humans; it is redirecting human effort to review and verification. [inference]**[S21] Harness Engineering — vault report · [S31] Coding is no longer the constraint — Spotify Engineering.*

Start with where the work goes, because this reframes the business case. [inference] As generation gets cheap, "the binding constraint migrates downstream" — off producing the artifact and onto understanding, validating, integrating and owning it. [S21] Here it is in a real number: Spotify's own engineering organisation reported 76% more pull requests to review once agents ramped. [S31] If you sold agents to leadership as headcount reduction, this number is the correction, and better to have it before the reorg than after. [inference]

45 45. The trade to watch

On screen:

Monday

- *Cheap generation moves the constraint to review — it does not remove it. [S21]*
- *More PRs mean more surface area; quality debt surfaces downstream if review cannot keep pace. [inference]*
- *The answer is the whole talk: deterministic gates and independent verification, so review scales without a human reading every line. [speculative]*

[S21] Harness Engineering — vault research report.

So the trade to watch as you adopt this, stated as a risk not a promise. [inference] Cheapening generation moves the binding constraint downstream to review and integration; it does not delete the constraint, it relocates it. [S21] And if review capacity does not keep pace with the new volume, the quality debt does not disappear — it surfaces downstream, later, more expensively, in production. [inference] The way out is not "review harder" or "hire an army of reviewers." [inference] It is the whole talk applied to your own pipeline: deterministic gates and independent verification, so review scales without a human reading every line the agents produced. [inference] Make the machine check what the machine can check, and spend your scarce human attention on the judgement calls a gate cannot make. [inference]

46 46. The job shifts

On screen:

Monday

"The engineer's value is defined not by what the agent generates but by the control plane the engineer builds to constrain, verify, and absorb it." [S23]

That moves you up the stack, not out of it. [speculative]

[S23] Harness Engineering — practitioner synthesis — vault research report.

Which changes the job description — in your favour, I think, though I know that is the anxious question under all of this. [inference] The report's framing: "the engineer's value is defined not by what the agent generates but by the control plane the engineer builds to constrain, verify, and absorb it." [S23] It becomes the control plane: the gates, the checks, the verification, the architecture that takes probabilistic output and makes it safe to run. [S23] The senior engineers here have done the human version of that for years; this is the same skill, aimed at a new source of unreliable output. [inference]

47 47. A deterministic-first recipe

On screen:

Monday

- *"Make the deterministic component load-bearing, make the model the orchestrator around it, and bound the whole with calibrated human oversight." [S8]*
- *Close every loop against an external oracle — a test, a solver, an independent gate — never the model that produced the work. [S14]*
- *Deterministic first; model judge only where the target is genuinely subjective. [inference]*

[S8] 5 Pillars — practitioner synthesis · [S14] Loop Engineering — vault research reports.

If you take one reusable recipe home, take this — the whole architecture in three clauses. [inference] "Make the deterministic component load-bearing, make the model the orchestrator around it, and bound the whole with calibrated human oversight." [S8] Deterministic part carries the weight; model drives; human owns the fence. [inference] And in every loop inside that architecture, close against an external oracle — a test, a solver, an independent gate — never against the model that produced the work. [S14] The priority ordering holds: deterministic checks first because they are reliable and cheap, a model judge only where the target is genuinely, irreducibly subjective. [inference] Notice the recipe says nothing about which model — on purpose, because that is the part that does not expire when the next model ships. [inference]

48 48. Advocate the idea, hedge the label

On screen:

Monday

- *The vocabulary is contested — "harness," "loop," "graph" engineering. [inference]*
- *"Advocate the engineering idea; hedge the compound noun." [S26]*
- *The durable idea — verifier-gated, externally-grounded, bounded autonomy — survives whatever the term becomes. [speculative]*

[S26] Loop Engineering — vault concept note.

A short caveat on vocabulary, so you are not caught out when someone pushes on the words. [inference] The labels in this space — harness engineering, loop engineering, graph engineering — are genuinely unsettled, and each was contested within weeks of being coined. [inference] So the vault's guidance, which I pass on directly: "advocate the engineering idea; hedge the compound noun." [S26] Do not go to war for the term — you will lose, because the term will change. [inference] Go to war for the durable idea underneath: verifier-gated loops, external grounding, bounded autonomy. [inference] That idea survives whatever the industry calls it next quarter, and it is what you are actually adopting — sell the mechanism, not the buzzword. [inference]

49 49. What this evidence does not settle

On screen:

Honest limits

-

Model-authored policies

[S18]

LLM-written policies leave residual attack. [S18]

-

Harness can backfire

[S28]

One tool-format swap cost ~38 points. [S28]

-

Widening surface

[S36]

GhostJacking extends agentjacking (Aug 2026). [S36]

-

Verifier scope

[S37]

Certifies only what it can see. [S37]

-

No-oracle zone

[S39]

No complete oracle; partial checks only. [S38][S39]

The most credible slide in a technical talk is the one about what the evidence does not settle, so here are five honest limits — I would rather raise them than have you raise them in Q&A. [inference] One: deterministic control is not a free zero once the model writes the policy — the residual attack comes back. [S18] Two: not every harness change helps, and one wrong tool-format choice cost thirty-eight points, so measure changes against an eval set. [S28] Three: the attack surface is widening, not narrowing — GhostJacking, disclosed this month, extends agentjacking by poisoning trusted logs and alerts. [S36] Four, and this one lands on my strongest number, so I raise it before you do: even that closed loop only fixes the errors its verifier can see — the study's own authors flag that a defect from structural topology rather than parameters left the parameter-level loop "powerless," so the guarantee is bounded by the scope of the oracle you build, not by the domain it came from. [S37] And five, the honest one a hostile expert will press: most of your production targets — is this refactor correct, is this summary good — have no cheap external oracle at all, so where none exists the honest architecture is bounded autonomy plus human review, and sometimes the right call is to refuse. [inference] Knowing which zone you are standing in — oracle-rich, oracle-scoped, or oracle-free — is itself part of the job. [inference]

50 50. Three moves

On screen:

Monday

1

Add a check

Find one self-graded loop; give it an independent check. [S14]

2

Gate one action

Put one deterministic gate on an irreversible action — a merge, a deploy, a write. [S8]

3

Check one figure

Take one number you cite about AI; confirm it is corroborated, not single-source. [S30]

None of these needs a better model; all of them are engineering. [inference]

Three concrete moves you could start on Monday morning — small enough to do, pointed enough to matter. [inference] One: find one loop where the model grades its own work, and give it an independent check — a test, a schema, an isolated-context checker — just one, as proof to yourselves that it changes the failure rate. [S14] Two: put one deterministic gate in front of one irreversible action — a merge, a deploy, a write to production — the Kiro fix, applied before your Kiro incident. [S8] Three: take one number you personally repeat about AI, and check whether it is corroborated or single-source — the thing I did with the 85% figure. [S30] Which is the point: the reliability was never going to come from the model. [inference]

51 51. Takeaways — what follows from the thesis

On screen:

In closing

1

Relocate the truth

Reliability is bought by relocating truth-determination out of the model. [S6]

2

Independence wins

Independence — not iteration — is the load-bearing variable. [S10]

3

Determinism first

Push determinism to the load-bearing control; bound the model inside it. [S16]

4

Autonomy last

Unbounded autonomy is unbounded risk [S7], so autonomy is what you add last. [inference]

The same four takeaways I opened with, now earned rather than asserted — this is the summary, and notice it is what follows from the thesis, not a recap of section titles. [inference] If reliability is bought by relocating truth-determination out of the model, then it follows that independence of the verifier — not how hard you iterate — is the load-bearing variable. [S6] [S10] It follows that you push determinism to the load-bearing control and bound the model inside it, at every layer of the harness. [S16] And it follows that because unbounded autonomy is unbounded risk [S7], autonomy is the property you add last, after the

check is in place. [inference] Four consequences, one thesis. [inference] If you disagree with the conclusions, the honest place to attack is the thesis they descend from — and that is exactly the argument I want to have next. [inference]

52 52. The one line to keep

On screen:

One line

"The model is a commodity; the harness is the product." [S22]

Reliability is engineered, not prompted. [speculative]

[S22] Harness Engineering — practitioner synthesis — vault research report.

If you forget everything else, keep this one line: "the model is a commodity; the harness is the product." [S22] The model is the commodity — it changes every few months, everyone has the same ones, and it is not where your durable advantage lives. [inference] The harness is the product — the control plane, the gates, the verification, the architecture — the engineering asset that makes a stochastic model behave well enough to trust in production. [S22] Which brings us back to where we started: reliability is engineered, not prompted. [inference] That is the claim I opened with; I have given you the evidence for it and the honest limits against it, and now I would genuinely like you to tell me where it breaks — because you are exactly the room that would know. [inference] Thank you. [inference]

53 53. References

On screen:

[S1] Building AI Systems: 5 Pillars — comprehensive report — vault research report.

[S2] 5 Pillars — context engineering — vault research report.

[S3] 5 Pillars — context rot — vault research report.

[S4] 5 Pillars — grounding, three layers — vault research report.

[S5] 5 Pillars — closed-loop external-verifier compliance result (56.8 → 98.6) — vault research report.

- [S6] 5 Pillars — relocating truth-determination — vault research report.
- [S7] 5 Pillars — bounded autonomy — vault research report.
- [S8] 5 Pillars — practitioner synthesis — vault research report.
- [S9] Loop Engineering — two loop archetypes — vault research report.
- [S10] Loop Engineering — verifier independence — vault research report.
- [S11] Loop Engineering — "Gödel with a leaderboard," narrowed to "the price is never zero" — vault research report.
- [S12] Loop Engineering — self-attribution bias (scoped: 1 preprint, "in one setting") — vault research report.
- [S13] Loop Engineering — rule-based more reliable, LLM more variable — vault research report.
- [S14] Loop Engineering — practitioner synthesis — vault research report.
- [S15] Loop Engineering — guarantee in the verifier — vault research report.
- [S16] Harness Engineering — the one law — vault research report.
- [S17] Harness Engineering — Progent, hand-written policy — vault research report.
- [S18] Harness Engineering — Progent, autonomous-policy residual — vault research report.
- [S19] Harness Engineering — agentjacking, execution-layer controls — vault research report.
- [S20] Coding Agent Horror Stories: The Agent That Deleted Production (Amazon Kiro) — Docker (2026). <https://www.docker.com/blog/coding-agent-horror-stories-the-agent-that-deleted-production/>
- [S21] Harness Engineering — binding constraint migrates downstream — vault research report.
- [S22] Harness Engineering — model is a commodity, harness is the product — vault research report.
- [S23] Harness Engineering — the control plane the engineer builds — vault research report.
- [S24] Harness Engineering — vault concept note.
- [S25] Harness Engineering — OpenAI ~1M lines / LangChain +13.7 — vault concept note.
- [S26] Loop Engineering — advocate the idea, hedge the noun — vault concept note.
- [S27] Intuit Scrapped Its AI Agent Architecture Twice in Four Months — VentureBeat, VB Transform 2026. <https://venturebeat.com/orchestration/intuit-scrapped-its-own-ai-agent-architecture-twice-in-four-months-at-vb-transform-2026-its-ai-vp-called-that-the-fast-path>
- [S28] Is Grep All You Need? How Agent Harnesses Reshape Agentic Search — vault reading note.

- [S29] Over 100 AI Coding Agents Taken Over Via Agentjacking — vault watch note.
- [S30] 85% Agentjacking Success Rate Not Independently Replicated — vault watch note.
- [S31] Coding Is No Longer the Constraint — Spotify Engineering. <https://engineering.atspotify.com/2026/6/code-with-claude-coding-is-no-longer-the-constraint>
- [S32] Progent: Securing AI Agents with Privilege Control — arXiv 2504.11703. <https://arxiv.org/abs/2504.11703>
- [S33] Building Enterprise AI Agents That Are Both Autonomous and Reliable — Salesforce Engineering. <https://engineering.salesforce.com/building-enterprise-ai-agents-that-are-both-autonomous-and-reliable/>
- [S34] Harness Engineering: Leveraging Codex in an Agent-First World — OpenAI. <https://openai.com/index/harness-engineering/>
- [S35] Improving Deep Agents with Harness Engineering — LangChain. <https://blog.langchain.com/improving-deep-agents-with-harness-engineering/>
- [S36] ThreatsDay: GhostJacking AI Attacks — The Hacker News (Aug 2026). <https://thehackernews.com/2026/08/threatsday-ghostjacking-ai-attacks.html>
- [S37] 5 Pillars — Grounding, the closed-loop boundary condition (verifier scope bounds the guarantee) — vault research report.
- [S38] Property-Based Testing for LLMs — arXiv 2307.04346, via vault note Alternative Methods to LLM Evaluations. <https://arxiv.org/abs/2307.04346>
- [S39] Metamorphic Prompt Testing for LLMs — arXiv 2406.06864, via vault note Alternative Methods to LLM Evaluations. <https://arxiv.org/html/2406.06864v1>
- [S40] Canary & Shadow Deployment for LLM Apps — vault note Alternative Methods to LLM Evaluations.
- [S41] Structural Grounding — vault concept note (three-layer grounding taxonomy).
- [S42] Idempotency in Distributed Systems — vault concept note.
- [S43] Bounded Autonomy in AI — vault concept note.
- [S44] Harness Engineering — deletion cadence — vault concept note.
- [S45] Physics-in-the-Loop: Validated CAD Engineering Design — Berger et al., IJCAI-ECAI 2026 (arXiv 2605.19717), via 5 Pillars report. <https://arxiv.org/abs/2605.19717>
- [S46] Loop Engineering — the theoretical case (Gödel / Tarski) — vault research report.
- [S47] Loop Engineering — two loop archetypes (self-audit study) — vault research report.

Reference card — pointers only, no new claim. [inference] Every figure on these slides traces to a numbered entry in sources.md, and the vault reports carry the deeper source chains. [inference] The handout goes a layer deeper on each. [inference]