Claude on the web: habits, projects, and checking

Contents

Generated by the agentic vault — an autonomous research and authoring harness. Every factual claim was gated against a cited source registry before this document was produced.

> ***Reading this document:*** `[S<n>]` *cites the numbered source list at the end.* `[inference]` *marks a claim that follows from those sources but is not stated by them.* `[speculative]` *marks the author's judgement, offered as judgement.*

1 The one question

Every disappointing answer you have had from Claude is one of two failures, and they have different fixes. Either the content is right and the writing is wrong — too long, too formal, aimed at the wrong reader. Or the writing is fine and the content is wrong, because Claude was never given the price list, the positioning, or the thing we never promise before it ships. That two-way split is my own framing rather than a result from any source below; I use it because each half has a different fix, and picking the wrong half is why people re-word a prompt that needed a document. [speculative]

The first is a wording problem and you fix it in the prompt. The second is a facts problem and no rewording touches it: wording can only select among what Claude already holds, and cannot add a fact it never had. That half of the split is mine too, and no source below measures it. [speculative] Anthropic's own engineering guidance frames the goal as "finding the smallest possible set of high-signal tokens that maximize the likelihood of some desired outcome" [S8] — the operative word is *smallest*, and read against this question the answer is the right document rather than every document. [inference]

So: **short of wording, or short of facts?** The four habits below are the wording half. Projects are the facts half. [speculative]

2 Four prompting habits

2.1 1. Brief Claude the way you brief a new colleague

Anthropic's own instruction is "Think of Claude as a brilliant but new employee who lacks context on your norms and workflows", and it supplies a test that costs nothing: "Show your prompt to a colleague with minimal context on the task and ask them to follow it. If they'd be confused, Claude will be too" [S3].

Three things go into almost every prompt: who will read the output, what it has to achieve, and what must not appear in it. The third is the one people skip, and it is the one that keeps a delivery date out of a draft. [inference]

This transfers with no new vocabulary, which is the point: the prompt has to carry exactly what you would have to tell a colleague who has never seen the file. [speculative]

2.2 2. Role prompting is a register control, not an accuracy control

Wharton Generative AI Labs tested domain-tailored expert personas against a no-persona baseline across six models. The two test sets, GPQA Diamond and MMLU-Pro, are standard multiple-choice exams used to score AI models on graduate-level science and general knowledge. Expert personas produced "no significant positive differences" over a no-persona baseline [S1]; *significant* here means the difference could not be told apart from chance.

Two limits on carrying that result to Claude. No Claude model is among Wharton's six [S1]. And persona effects do not transfer across model families: on one labelling task, 11 of 21 personas significantly *improved* Mistral-Medium's label accuracy while 16 of 21 significantly *degraded* Qwen3-32B's [S24]. So reading Wharton's null result as a fact about Claude is my judgement rather than a measurement. [speculative] What the sources here do establish is narrower and still enough to act on: no persona reliably buys accuracy — Wharton found no significant gain across six models, and where personas did shift accuracy the direction depended on the model. [inference]

The one *measured* style effect in my own notes points the opposite way from the folklore — with an important caveat about what it measured. Across 1,140 open-ended questions, 38 expert roles and six domains, a model-based judge — another AI model scoring the answers — rated those written under an expert role as deeper in expertise and less clear than answers written without one. The note's own wording: "hybrid role retrieval raised expertise depth from 3.638 to 3.923 (Cohen's $d = 0.54$, $p < 0.001$) while *lowering* clarity from 4.896 to 4.550 ($p < 0.001$)" [V1].

Two pieces of jargon in that sentence, and neither is needed to use the result. *$p < 0.001$* means chance is an implausible explanation for the difference. *Cohen's $d = 0.54$* is an effect size — the difference expressed in standard deviations. Calling 0.54 *moderate* is my gloss by the conventional bands, not a word my note or the study attaches to the number. [inference] What the sentence does not state is the scale those ratings run on, so read the direction — depth up, clarity down — rather than the decimals. [inference] Hybrid role retrieval is also a retrieval technique, not a plain "act as an expert" instruction, so this is the nearest evidence rather than a measurement of the same thing. [inference]

The limit is in the same note, and it matters: in aggregate the two effects roughly cancelled, with no significant overall difference from the no-role baseline [V1]. The two figures are also not directly comparable — different rated constructs, and the study reports an effect size for the depth change only. What survives is narrow: the one measured style effect includes a rated clarity cost, on the same model-judged scale as the depth gain [V1]. Where a text will be read once, by someone who cannot ask a follow-up question, a rated clarity drop is the half of that trade I would not spend. [speculative] Use role prompting when you want a particular register; do not use it hoping for a more correct answer. [speculative]

Note the limit on this: Wharton measured accuracy on objective questions, not readability, and the depth/clarity study is a separate experiment whose ratings came from another AI model rather than from people [V1]. Neither establishes that role prompting is useless — they establish that its benefit is not accuracy. [inference]

2.3 3. Let Claude ask, then restate everything in one message

You do not need a perfect prompt from a standing start. Ask for the questions first: *before you write anything, ask me the five questions you need answered.* No source below measures that move; it is my own practice. [speculative]

The second move is the one people drop, and dropping it makes the habit worse than not using it. Once you have the questions, answer them in a **single** message rather than one reply per turn. Laban et al. simulated over 200,000 conversations and measured an "average performance drop of 39%" when the same information arrived spread across turns instead of in one prompt, and reported that "when LLMs take a wrong turn in a conversation, they get lost and do not recover" [S2] — that is, an early error was not corrected by anything later in the conversation, so the rest of the exchange built on it. [inference]

That study is an arXiv preprint using simulated rather than human conversations, so the exact figure should be read as a direction rather than a constant [inference]. The direction is enough: elicit, then consolidate.

2.4 4. One deliberate check step between the draft and the send button

This is the habit to keep if you keep only one. The mechanism is simple and it is not a matter of the model being immature: "An LLM has no way to distinguish between a 'hallucination' and the correct answer — in both cases, it 'makes up' the answer by relying on statistics to guess the next word" [S10]. There is no internal signal that fires when the output is wrong, so a confident tone carries no information about correctness. [inference]

The check therefore has to come from outside, and it has to be small enough that you will actually do it every time. Pick the single fact that would cost the most if it were wrong — the price, the certification, the deadline — and confirm it against the source. Not by asking Claude again in a fresh chat, which is the same mechanism twice. [inference]

MIT Sloan reported a field experiment with 140 Accenture research professionals in which "intense highlighting techniques improved error and omission detection but increased task completion time, with the medium friction condition striking the right balance between accuracy and efficiency" [S11]. Note what was varied there: the intensity of colour-coding applied to the output, not how many claims the reader chose to check. In this deck's terms the finding is a direction rather than a procedure — a check heavy enough to slow you down catches more, so the version worth running every time is one deliberate check on the single fact whose being wrong would cost most. [speculative]

2.5 Why habit 4 is in the takeaways rather than in a footnote

I am carrying a risk into this session, and it is fair to state it plainly — with the limit on the evidence first: nobody has measured the effect below on a room that has just been trained, so carrying it into this session is my step rather than the study's. [speculative] Fernandes et al., with 500 participants, found the Dunning-Kruger pattern *vanishes* under LLM use: "We found that when it comes to AI, the DKE vanishes. In fact, what's really surprising is that higher AI literacy brings more overconfidence" [S13].

A PRISMA systematic review of 28 studies of non-IT laypersons found "a consistent pattern of overestimation of chatbot capabilities due to limited understanding", naming specifically "the attribution of human-like language comprehension" and the assumption "that the systems can retrieve information in real time" [S12].

The overconfidence finding is measured [S13]; that you will produce more after this session is my expectation rather than a finding. [speculative] Habit 4 is the only one of the four that puts a person between the draft and the reader, which is why it is a takeaway and not a caveat. [inference]

3 Projects — the facts half

3.1 What a project is, and why it beats retyping

Anthropic defines projects as "self-contained workspaces with their own chat histories and knowledge bases" [S4]. The reason to want one is narrower than that definition: the background paragraph you retype into every new chat goes in once as the project's *instructions*, the documents you keep re-uploading go in once as *project knowledge*, and every chat started inside the project begins already briefed. [inference]

3.2 The three shared projects

Prompt help. For when you know what you want and cannot get Claude to produce it. Not for when the answer is factually wrong — a better-worded prompt does not supply a missing fact. That split is my own framing rather than a result from any source below. [speculative]

This distinction is worth labouring. Deel, rolling Claude out to 7,000 staff, reported "people were working inside a Claude Project when what they actually needed was a Skill... That's a config problem that looks like a capability problem from the inside" [S22] — a *Skill* here being the third of the three personalization places described in the next section, not the ordinary noun — a different pair of features from the one above, but the same shape of error. [inference] From the user's seat it feels like the tool cannot do the job; in fact the wrong one of the three places is open. [inference]

Personalization. It is not one switch. Anthropic's page carries three sections — *Instructions for Claude*, *Project instructions* and *Skills* [S5]. The page states the reach of the first two directly: instructions for Claude are "account-wide settings" and "will be applied to all of your conversations with Claude", while project instructions "only apply to chats within that project" [S5]. It does not state a reach for skills at all — it describes them by function, as what to use "when you want to customize how Claude formats and delivers its responses" [S5], so the account-wide-versus-per-project rule below applies to the first two and not to the third. [inference]

One naming trap is worth carrying into the room, and it is the vendor's rather than the reader's: the page's closing summary says "Use profile instructions for account-wide settings", while its heading and the settings screen both say *Instructions for Claude* [S5]. Same setting, two names. Go by the label on the screen. [inference]

The rule that keeps this tidy: put a preference at the widest scope where it is *always* true, and never wider. A tone rule that belongs to one campaign, set account-wide, will follow you into every unrelated conversation. [inference] It is the same shape as company policy → team norm → today's brief. [inference]

Set one preference, then immediately re-ask a question you asked yesterday — one whose answer would visibly change if the preference had applied. If the answer does not change, the likeliest cause is that the preference is not in scope for that chat, because it went into the wrong one of the three places. [inference] Nothing I have seen in the product announces that a preference did not apply, so an unverified preference goes on being trusted — my own observation of it, not something the documentation states. [speculative]

LinkedIn drafts. Two numbers set the constraint. Pangram Labs analysed 1,002,627 posts and found "LinkedIn was the most AI-saturated platform, where more than 40% of longform posts flagged as fully AI-generated" [S14]. Originality.ai, on 2,726 posts, found "the average likely-AI-generated post received 45% less engagement than a Likely Original post" [S15].

Both figures come from companies selling AI-detection tooling, and both rest on their own classifiers being right, so they are a direction rather than a measurement to bet on [inference]. The direction is enough to change how the project is used: take the structure and the first draft from it, and keep the specifics — the customer, the number, the thing that actually happened — in your own words.

3.3 What not to put in a shared project

"Be deliberate about what goes into shared Projects because the instructions and documents are visible to everyone with access" [S19]. Assume everything you upload is readable by everyone the project is shared with. Whether your conversations inside a shared project are equally visible is not stated by any source here, and is worth confirming before you rely on either answer. [speculative]

The test: would you send this file to everyone the project is shared with? If not, it belongs in a private project instead. [inference]

Two operational hazards, both real. "Two people editing the same Project will overwrite each other" — there is no merge [S19]. And stale files never announce themselves: "They upload stale material, forget it exists, and then wonder why Claude repeats old details" [S20].

4 Build your own

4.1 Four steps

1. **Create it** from the projects panel, and name it after a recurring job — *replies to inbound enquiries* — rather than a topic such as *sales*. No source below recommends that naming: it is my own practice, and the reason is that one project carries one set of instructions and one set of files, which can only be right for one repeating job. [speculative]
2. **Paste the background paragraph** in as the project's instructions. [inference]
3. **Upload two or three documents**, not twenty. Anthropic's own guidance is "the smallest possible set of high-signal tokens" [S8], and two or three files is what I take that to mean for a project of this kind [speculative]. Project files are capped at "File size: 30MB per file", which differs from the chat upload limit on the same page [S6].
4. **Share it — if the plan allows.** Anthropic lists "Collaboration and sharing (Team and Enterprise plans only)", with two permission levels, "Can view" and "Can edit" [S4]. Confirm what your own account carries before promising a colleague access; free accounts are also capped at five projects [S4].

4.2 Instructions describe behaviour; documents carry content

"Write instructions that describe behavior, not content… 'Here is a summary of our Q3 strategy' is content, which belongs in a document, not in instructions" [S19]. People invert this constantly, because pasting content into an instructions box feels like configuring something. [speculative]

The reason it is not cosmetic: the project changes what it does underneath you, at a threshold the vendor names. There is a limit to how much text Claude can read at one go — the vendor calls it the context window. As your uploaded files approach that limit, the project stops reading everything you uploaded and starts searching those files and reading back the parts it finds. Anthropic's name for that mode is RAG, short for retrieval-augmented generation: "When your project knowledge approaches the context window limit, Claude will automatically enable RAG mode to expand your project's capacity by up to 10x", and "RAG activates automatically when needed. No setup or configuration is required" [S7]. Nothing on screen tells you the switch happened — that is my own observation of the product, not something the documentation states. [speculative]

The guidance is the same one from the opening — "the smallest possible set of high-signal tokens" [S8]. More documents is therefore not more accurate answers, and the practical form of that is: when a document is superseded you replace it, rather than adding the new one beside the old. [inference]

4.3 Starting from a slide template

The exercise works and the obvious first move does not. Anthropic's supported document formats are "PDF, DOCX, CSV, TXT, HTML, ODT, RTF, EPUB, JSON, XLSX" [S6]. PPTX is not on that list, so the template file itself will not upload. [inference]

Two workarounds: export the template to PDF and upload that, or type the slide structure into the project instructions directly [inference]. The second is often better, because a fixed slide sequence is behaviour, and behaviour belongs in instructions [S19].

What comes back is the words for each slide against your standing structure. The branded file is still assembled in PowerPoint, which is fine: the writing is the part that took the time. [speculative]

4.4 An owner and a review date

Custom GPTs were the same idea one product generation earlier. OpenAI reported in January 2024 that "users have already created over 3 million custom versions of ChatGPT" [S16]. Two years later an independent count puts roughly "159,000 are public and active in the store" [S17] — about one in twenty, which is my arithmetic across the two counts rather than a ratio either source reports [inference]. That counts public store listings rather than working assistants, and comes from an analysis site with no published counting method, so it is an order of magnitude rather than a figure [inference].

Creation is free and maintenance is not — that gap is my own read of the two figures above, and neither source states it. [speculative] One practitioner write-up describes the failure and names the fix: "They upload stale material, forget it exists, and then wonder why Claude repeats old details. The fix is an owner and a review date" [S20]. That source is a blog with no measurement behind it, so this is practitioner experience rather than evidence. Put a name and a date in the project's own instructions, so the next person can see how much to trust it. [inference]

5 Why I keep the session short and protect the practice time

Eight of the twenty-five minutes are hands-on, and that is a design decision rather than a shortage of material. Correlation One's write-up of its Claude adoption programmes quotes one client team near-verbatim — "access is not the problem — the bottleneck is fluency" — and says in its own voice that employees "can't transfer a generic demo onto their actual Tuesday" [S18]. The first of those is one team's self-report rather than a pattern the consultancy measured across its rollouts, so it is evidence that the failure happens somewhere, not evidence of how often. [inference]

Fluency on its own is an evaluative noun with no measure attached, so take that team's own naming of what was missing rather than the word: most of them "didn't know which tool or surface to use for which task, how to query effectively across their document environment, or how to verify an output before relying on it" [S18]. Three deficits, and I take them in that order: I answer *which surface* with the wording-or-facts question and the three shared projects, *how to ask* with the four habits, and *how to verify* with habit 4. [inference] I have measured nothing about whether that transfer happens; the only test available to you is whether one of your own recurring jobs runs differently next week. [speculative]

The ordering — demonstrate each habit, then practise — is also deliberate. For novices, worked examples and example-problem pairs "would lead to lower cognitive load during learning and higher learning outcomes than PS" (problem-solving alone) [S21]. That is the study's predicted direction for novices, which is what this room is; it says nothing about a room of daily users. [inference]

Most of the room is not starting from zero. One 2025 survey reports "employees at 90% of companies regularly use AI tools at a personal level" [V2]. The note carrying that figure files it under its own contested findings, not its verified ones, in a report whose methodology was publicly attacked — so it is a rough shape rather than a count, and it is quoted here only to say that most of the room already uses something unofficially. [inference] So access is not what I am fixing — most of the room already has it. What I am fixing is that the background paragraph gets retyped into every new chat instead of being written once into a project's instructions. [inference]

6 What I am not covering, and why

Multi-step document production. Anthropic sells a separate product for it: "Cowork is built for non-coding knowledge work — research, analysis, document creation, and other multi-step tasks" [S9]. If that is the shape of your work, start with Cowork rather than with anything here. [inference]

I also skip per-project memory, usage limits, and the account's data-retention settings. They change what Claude has available and who else can see it; they do not change how reliable a given answer is, and reliability is what the four habits and the check step are for. [speculative]

If you want to read further, the field files all of this under *prompt engineering*. The standard reference is Schulhoff et al.'s *The Prompt Report*, a systematic review of 1,565 papers that catalogues 58 text-based prompting techniques [S23]. That is the phrase to search on.

7 References

- **[S1]** *Playing Pretend: Expert Personas*, Prompting Science Report 4 — Wharton Generative AI Labs, December 2025. Industry report; credibility High. https://gail.wharton.upenn.edu/research-and-insights/playing-pretend-expert-personas/
- **[S2]** Laban et al., *LLMs Get Lost In Multi-Turn Conversation*, arXiv:2505.06120 (2025). Academic paper (preprint); credibility Medium. https://arxiv.org/abs/2505.06120
- **[S3]** Anthropic, *Be clear and direct* (prompt engineering documentation). Reference; credibility High. https://docs.claude.com/en/docs/build-with-claude/prompt-engineering/be-clear-and-direct
- **[S4]** Anthropic Help Center, *What are Projects?* Reference; credibility High. https://support.claude.com/en/articles/9517075-what-are-projects
- **[S5]** Anthropic Help Center, *Understanding Claude's personalization features*. Reference; credibility High. https://support.claude.com/en/articles/10185728-understanding-claude-s-personalization-features
- **[S6]** Anthropic Help Center, *Upload files to Claude*. Reference; credibility High. https://support.claude.com/en/articles/8241126-upload-files-to-claude
- **[S7]** Anthropic Help Center, *Retrieval augmented generation (RAG) for projects*. Reference; credibility High. https://support.claude.com/en/articles/11473015-retrieval-augmented-generation-rag-for-projects

- **[S8]** Anthropic, *Effective context engineering for AI agents*. Blog / editorial; credibility High. https://www.anthropic.com/engineering/effective-context-engineering-for-ai-agents
- **[S9]** Anthropic, *Claude Cowork* (product page). Reference; credibility High. https://www.anthropic.com/product/claude-cowork
- **[S10]** IBM, *What are AI hallucinations?* Reference; credibility High. https://www.ibm.com/think/topics/ai-hallucinations
- **[S11]** MIT Sloan, *To help improve the accuracy of generative AI, add speed bumps*. News; credibility High. https://mitsloan.mit.edu/ideas-made-to-matter/to-help-improve-accuracy-generative-ai-add-speed-bumps
- **[S12]** *Universal Access in the Information Society* (2026), PRISMA review of 28 studies of non-IT laypersons. Academic paper; credibility High. https://link.springer.com/article/10.1007/s10209-026-01323-w
- **[S13]** Aalto University on Fernandes et al., *Computers in Human Behavior* 175:108779. News; credibility Medium. https://www.aalto.fi/en/news/ai-use-makes-us-overestimate-our-cognitive-performance
- **[S14]** Pangram Labs, *AI in your feed* (1,002,627 posts). Industry report; credibility Medium. https://www.pangram.com/blog/ai-in-your-feed
- **[S15]** Originality.ai, analysis of 2,726 LinkedIn posts. Industry report; credibility Medium. https://originality.ai/blog/ai-content-published-linkedin
- **[S16]** OpenAI, *Introducing the GPT Store*, 10 January 2024. Primary source; credibility High. https://openai.com/index/introducing-the-gpt-store
- **[S17]** DigitalApplied, *GPT Store and custom GPTs: business guide 2026*. Blog / editorial; credibility Low. https://www.digitalapplied.com/blog/gpt-store-custom-gpts-business-guide-2026
- **[S18]** Correlation One, *Enterprise Claude adoption: 7 real problems and the solutions that worked*. Blog / editorial; credibility Medium. https://www.correlation-one.com/blog/enterprise-claude-adoption-7-real-problems-and-the-solutions-that-worked
- **[S19]** EnterpriseDNA, *Claude Projects tutorial*. Blog / editorial; credibility Low. https://enterprisedna.co/resources/guides/guide-claude-projects-tutorial
- **[S20]** OpenClawReady, *Claude Projects for teams*. Blog / editorial; credibility Low. https://openclawready.com/blog/claude-projects-for-teams
- **[S21]** van Gog, Kester & Paas, *Effects of worked examples, example-problem, and problem-example pairs on novices' learning*, Contemporary Educational Psychology. Academic paper; credibility High. https://www.sciencedirect.com/science/article/abs/pii/S0361476X1000055X

- **[S22]** Deel, *Non-technical AI adoption*. Blog / editorial; credibility Medium. https://www.deel.com/deel-works/non-technical-ai-adoption
- **[S23]** Schulhoff et al., *The Prompt Report: A Systematic Survey of Prompting Techniques*, arXiv:2406.06608 (2024). 1,565 papers reviewed under PRISMA; taxonomy of 58 text-based prompting techniques. Academic paper; credibility High. https://arxiv.org/abs/2406.06608
- **[S24]** Yang, Hechtbauer, Khalilov, Brinkmann, Schmitt & Feldhus, *Persona Prompting as a Lens on LLM Social Reasoning*, EACL 2026 (main), pp. 1152-1170, doi: 10.18653/v1/2026.eacl-long.52. Three LLMs, 21 personas, three datasets of varying subjectivity. Academic paper; credibility High. https://aclanthology.org/2026.eacl-long.52/
- **[V1]** *Role Prompting* — author's working notes. Summarised in the appendix below.
- **[V2]** *State of AI in Business 2025 (MIT NANDA) — Overview* — author's working notes. Summarised in the appendix below.

8 Appendix — sources from the author's working notes

The talk cites material from a private research vault. Each entry below summarises one such note and lists every external source it draws on.

8.1 V1 — Role Prompting (Persona Prompting) in Large Language Models

Role prompting — assigning a model a system-message identity or expert persona to steer style, tone or reasoning — has a documented but narrower value than its popularity suggests, according to this note's synthesis of 2024–2026 empirical literature. Its central finding is a task-dependent trade-off: personas help alignment-oriented tasks (tone, style, human preference) but tend to damage accuracy on knowledge- and logic-dependent tasks. The quoted evidence, from a large-scale study of open-ended generative answers, states: "Across 1,140 open-ended questions, 38 expert roles and six domains, hybrid role retrieval raised expertise depth from 3.638 to 3.923 (Cohen's d = 0.54, p < 0.001) while *lowering* clarity from 4.896 to 4.550 (p < 0.001)." So an expert persona measurably deepened perceived expertise while measurably reducing clarity — a genuine two-sided shift, not a pure gain. The note reports this trade-off is domain-conditional (personas won on advisory questions in medicine and psychology; baseline prompting won on conceptual/explanatory questions in finance, legal, science and technology), and that in aggregate the two effects roughly canceled, with no significant overall difference from baseline.

The note is a literature synthesis, not original experimentation, and flags several limits. The quoted result rests on a single generation model, a synthetic benchmark and an LLM judge that its own authors warn may favor verbosity and structure — the source paper asks that it be read as controlled comparative evidence, not a definitive human-preference measurement. The note does not claim personas are uniformly harmful: it records a contested paper finding personas "usually lead to positive or non-significant performance changes," and shows persona effects vary sharply by model, so a persona validated on one backend cannot be assumed to transfer. It treats a supporting preprint about a persona causing a model to refuse coding tasks as its weakest evidence — unreviewed, uncited, unusual institutional provenance — and instructs it be read only as illustration. Left explicitly open: whether instance-level persona selection can become cheap and portable enough for production beyond small benchmarks; why persona effects reverse on rigid-structure reasoning tasks; and why models are so sensitive to irrelevant persona details, an unexplained ~30-point accuracy swing.

Referenced sources

- Prompting Science Report 4 — Playing Pretend: Expert Personas Don't Improve Factual Accuracy — Wharton Generative AI Labs
- Expert Personas Improve LLM Alignment but Damage Accuracy: Bootstrapping Intent-Based Persona Routing with PRISM — Hu, Rostami & Thomason, USC
- When "A Helpful Assistant" Is Not Really Helpful: Personas in System Prompts Do Not Improve Performances of Large Language Models — Zheng et al., Findings of EMNLP 2024
- Principled Personas: Defining and Measuring the Intended Effects of Persona Prompting on Task Performance — EMNLP 2025 (main)
- Telling an AI model that it's an expert makes it worse — The Register (2026)
- Persona Prompting as a Lens on LLM Social Reasoning — Yang, Hechtbauer, Khalilov, Brinkmann, Schmitt & Feldhus, EACL 2026
- One Persona, Many Cues, Different Results — Weeber et al., ACL 2026
- Research Shows Where Persona Prompting Works And When It Backfires — Search Engine Journal (2026)
- A Systematic Analysis of the Impact of Persona Steering on LLM Capabilities — Chen, Wang, Xie, Feng & Liu
- Persona is a Double-Edged Sword: Rethinking the Impact of Role-play Prompts in Zero-shot Reasoning Tasks — Kim, Yang & Jung, Seoul National University; Findings of IJCNLP-AACL 2025
- Persona Switch: Mixing Distinct Perspectives in Decoding Time — Kim, Yang & Jung, Findings of EACL 2026

- When Does Persona Prompting Actually Help? A Retrieval and Metric Analysis of Expert Role Injection in LLMs — Xiao, Liu, Zhou, Wu, He, Lin & Xie
- The Librarian Who Refused to Code: Model-Dependent Identity Enactment in LLM Code Generation — Salomon, Israel, Safruti & Shaked, Bluefields Indian and Caribbean University
- Persona Prompts and Other Sources of False Refusal in Language Models — WiNLP 2025
- Beyond Static Personas: Situational Personality Steering for Large Language Models — Wei, Li, Wang & Deng, USTC / Singapore Management University; Findings of ACL 2026

8.2 V2 — MIT NANDA's 2025 report on the "GenAI Divide" in enterprise AI adoption

The note reviews a July 2025 MIT NANDA draft report, "The GenAI Divide: State of AI in Business 2025," by Aditya Challapally, Chris Pease, Ramesh Raskar, and Pradyumna Chari, drawn from 52 executive interviews, 153 survey responses, and 300+ implementation reviews. Among its findings, which the note flags as contested, is a "Shadow AI" phenomenon: the report states that "employees at 90% of companies regularly use AI tools at a personal level," despite only 40% of companies having purchased official LLM subscriptions. As reported by Fortune, this sits alongside the report's separate, more heavily disputed headline claim that 95% of generative AI pilots fail to reach production with measurable P&L impact, producing what the note calls a paradox — widespread individual usage coexisting with limited organizational transformation. The note's synthesis leans toward reading this as evidence of an integration gap rather than a capability gap: "the problem may not be AI capability but organizational integration challenges," with adoption proceeding through informal, employee-led channels faster than top-down enterprise deployment.

The note does not claim to know what employees are actually doing with these tools, how the 90% figure was operationally measured, or whether personal usage translates into any measurable output. It flags the same methodological concerns that apply to the report as a whole — a small, non-random sample (300 public initiatives), no publicly released raw data, and a possible conflict of interest given MIT NANDA's ties to agentic-AI infrastructure (MCP, A2A) it may be positioned to promote. The note explicitly distinguishes the report's own "GenAI Divide" framing (a gap between organizations achieving production deployment and those stuck in pilots) from the media's blunter "95% failure" headline, and lists the question of how to define "success" — production deployment versus regular use versus measurable ROI — as an unresolved, contested debate rather than something the report settles. It likewise leaves open whether the shadow-usage pattern represents the fastest enterprise technology adoption in corporate history or simply a measurement gap, without adjudicating between the two.

Referenced sources

- The GenAI Divide: State of AI in Business 2025 (MIT NANDA, v0.1 draft, Aditya Challapally, Chris Pease, Ramesh Raskar, Pradyumna Chari): https://mlq.ai/media/quarterly_decks/v0.1_State_of_AI_in_Business_2025_Report.pdf
- aigl.blog analysis of the report: https://www.aigl.blog/state-of-ai-in-business-2025/
- Futuriom, "Why We Don't Believe MIT NANDA's Weird AI Study": https://www.futuriom.com/articles/news/why-we-dont-believe-mit-nandas-werid-ai-study/2025/08
- Fortune, "MIT report: 95% of generative AI pilots at companies are failing": https://fortune.com/2025/08/18/mit-report-95-percent-generative-ai-pilots-at-companies-failing-cfo/
- McKinsey, "The state of AI" (State of AI 2025): https://www.mckinsey.com/capabilities/quantumblack/our-insights/the-state-of-ai
- Fortune, "An MIT report that 95% of AI pilots fail spooked investors, but the reason why those pilots failed is what should make the C-suite anxious": https://fortune.com/2025/08/21/an-mit-report-that-95-of-ai-pilots-fail-spooked-investors-but-the-reason-why-those-pilots-failed-is-what-should-make-the-c-suite-anxious/
- Fortune, "Shadow AI economy" reporting on the MIT NANDA findings: https://fortune.com/2025/08/19/shadow-ai-economy-mit-study-genai-divide-llm-chatbots/
- pivot-to-ai.com, "MIT on enterprise AI: 95% failure, 5% spam": https://pivot-to-ai.com/2025/08/21/mit-on-enterprise-ai-95-failure-5-spam/
- Menlo Ventures, "2024: The State of Generative AI in the Enterprise": https://menlovc.com/2024-the-state-of-generative-ai-in-the-enterprise/
- Sify, "95% companies failing with AI: an MIT NANDA report misread by all": https://www.sify.com/ai-analytics/95-companies-failing-with-ai-an-mit-nanda-report-misread-by-all/