

Claude on the web: habits, projects, and checking

Four habits and three projects already in your account are enough to do real work with Claude this week; building one of your own is what makes it repeatable rather than a one-off.

Ondrej (Ondra) Krajicek

me@ondrejkrjicek.com · linkedin.com/in/OndrejKrajicek

Built 2026-09-07 03:40 UTC

Generated by the agentic vault — an autonomous research and authoring harness. Every factual claim was gated against a cited source registry before this document was produced.

Agenda — four habits, three shared projects, one thing you build yourself

- 1 Four prompting habits** Brief it [S3]; role: wording not accuracy [S1] [V1]; answer in one message [S2]; check a fact [S10]
 - 2 Three shared projects** Prompt Engineering, Personalization, LinkedIn posts
 - 3 Your own project** Create and share it [S4]; owner and review date [S20]
 - 4 Eight minutes hands-on** Practise on your own recurring job [S18]
- Not covered: multi-step document work — Anthropic sells Cowork [S9].

[S9] Claude Cowork — Anthropic. <https://www.anthropic.com/product/claude-cowork> [S18] Enterprise Claude adoption — Correlation One. <https://www.correlation-one.com/blog/enterprise-claude-adoption-7-real-problems-and-the-solutions-that-worked>

WRITE THESE DOWN

Takeaways — four sentences, before we start

- | | | |
|---|------------------------------|--|
| 1 | Wording or facts | Rewrite the prompt, or add a document |
| 2 | Role prompting | No accuracy gain: six models, two standard exams [S1] |
| 3 | Ask, then consolidate | Spread over turns: 39% lower average performance than one message [S2] |
| 4 | One check step | No signal fires when Claude is wrong [S10] |

[S1] *Playing Pretend* — Wharton GAIL (2025). <https://gail.wharton.upenn.edu/research-and-insights/playing-pretend-expert-personas/> [S2] *LLMs Get Lost In Multi-Turn Conversation* — Laban et al. (2025). <https://arxiv.org/abs/2505.06120> [S10] *AI hallucinations* — IBM. <https://www.ibm.com/think/topics/ai-hallucinations>

You have retyped the same background paragraph into a new chat this week

The same paragraph — what we sell, what we never promise — pasted into every new chat, because Claude starts each one knowing nothing about us.

Already using AI
"Employees at 90% of companies regularly use AI tools at a personal level" — contested [V2]

What is missing
"Access is not the problem — the bottleneck is fluency" [S18]

"Fluency" is three gaps [S18]: **which tool** to use, **how to ask**, and **how to check** the answer. This talk is those three habits.

[S18] Enterprise Claude adoption — Correlation One. <https://www.correlation-one.com/blog/enterprise-claude-adoption-7-real-problems-and-the-solutions-that-worked> [V2] State of AI in Business 2025 — author's notes.

One question before every prompt: short of wording, or short of facts?

Short of wording

True, but reads wrong — too long, too formal, wrong reader. Fix it in the prompt.

The test: could a smart colleague who has never seen our files write this answer? If yes, it is a **wording** problem — fix the prompt. If no, Claude is **short of facts** it was never given, and no rewording invents them. Then the fix is a document you upload, not a better sentence [S8].

Short of facts

Fluent and wrong, because Claude never had our pricing, our positioning, our customer. Fix it with a document.

[S8] *Effective context engineering* — Anthropic. <https://www.anthropic.com/engineering/effective-context-engineering-for-ai-agents>

Habit 1 — brief Claude the way you brief a new colleague

Anthropic: "a brilliant but new employee who lacks context on your norms and workflows" [S3] — capable, and holding no facts about us.

Anthropic's test costs nothing: give the prompt to a colleague who does not know the task [S3]. Whatever they would have to ask you is what the prompt is missing.

Who reads it	A procurement manager who has never bought from us
What it is for	A reply that wins a meeting, not a sale
What to avoid	No pricing, no dates

[S3] Be clear and direct — Anthropic. <https://docs.claude.com/en/docs/build-with-claude/prompt-engineering/be-clear-and-direct>

Habit 2 — "act as an expert" does not make Claude more correct

Role prompting means naming a role — an "expert persona", in Wharton's words [S1].

no significant

accuracy gain from role prompting —
six models, two standard exams (GPQA
Diamond, MMLU-Pro) [S1]

So telling Claude it is a senior engineer does not make it more correct — Wharton's pre-registered analysis found no significant gain over no role at all [S1]. What it moves is the wording, measured next.

[S1] *Playing Pretend* — Wharton GAIL (2025). <https://gail.wharton.upenn.edu/research-and-insights/playing-pretend-expert-personas/>

What adding a role changed: rated depth up, rated clarity down, no overall gain

Another AI model scored the answers to 1,140 questions [V1].

**sounds
more
expert**

expertise depth 3.638 → 3.923,
Cohen's d = 0.54 [V1]

**reads less
clearly**

clarity 4.896 → 4.550, $p < 0.001$ [V1]

No overall gain: the two effects roughly cancelled [V1]. My note gives no rating scale: read direction, not size. The study picked the role automatically; you type yours [V1]. That the trade survives typing is my judgement. Use role prompting for wording, never for accuracy.

Habit 3 — let Claude ask, then restate everything in one message

1 **Ask for the questions** "Before you write anything, ask me the five questions you need answered."

2 **Answer in one message** Not five replies across five turns

Why restate it all at once: drip-feeding — "write a post", then "make it shorter", then "add our pricing", then "change the tone" — is the trap. Across 200,000+ simulated conversations, spreading the same information over turns instead of one message cost a **39% average performance drop**, and one early wrong turn was never recovered [S2].

[S2] LLMs Get Lost In Multi-Turn Conversation — Laban et al. (2025), arXiv:2505.06120. <https://arxiv.org/abs/2505.06120>

Habit 4 — one deliberate check between the draft and the send button

Claude produces a true sentence and a hallucination by the same mechanism [S10] — nothing signals a wrong answer. Check the costliest fact.

- Read the number off the source page, not from Claude.
- Ask the person who owns the fact, not Claude.
- Delete any sentence you cannot trace to a page or person.

How much? An MIT Sloan study had 140 professionals check AI output: flag every claim and they caught the most errors but were slowest; flag only the risky few and they caught nearly as much, far faster [S11]. So check the costliest fact — not zero, not an audit.

[S10] AI hallucinations — IBM. <https://www.ibm.com/think/topics/ai-hallucinations> [S11] Add speed bumps — MIT Sloan. <https://mitsloan.mit.edu/ideas-made-to-matter/to-help-improve-accuracy-generative-ai-add-speed-bumps>

The four habits — and a study where higher AI literacy came with more overconfidence

- | | | |
|---|------------------------------|----------------------------------|
| 1 | Brief it | Reader, purpose, what to avoid |
| 2 | Role for wording | Not for accuracy [S1] |
| 3 | Ask, then consolidate | One message, not five turns [S2] |
| 4 | One check step | The costliest fact [S10] |

Across 500 participants, "higher AI literacy brings more overconfidence" [S13] — never measured on a room leaving a training session, so carrying it here is my step. Habit 4 alone puts a person between draft and reader.

[S13] AI and overconfidence — Aalto University (2026). <https://www.aalto.fi/en/news/ai-use-makes-us-overestimate-our-cognitive-performance>

A project holds the background paragraph, so you stop retyping it

Anthropic's definition: "self-contained workspaces with their own chat histories and knowledge bases" [S4].

Without one
You paste who we are and what
we never promise into every new
chat, slightly differently each time.

With one
Write it once as instructions,
upload the documents once, and
every chat inside starts already
briefed.

This is the facts half of the wording-or-facts question — not a better way to word things, but a place to put what Claude could not otherwise have.

[S4] What are Projects? — Anthropic Help Center. <https://support.claude.com/en/articles/9517075-what-are-projects>

Shared project 1 — Prompt Engineering, for when you are short of wording

Where

Open your Projects list; sharing needs Team or Enterprise [S4]

Use it when

You cannot get Claude to word what you want. Not when facts are wrong

What comes back

Treat it as a draft prompt, not as finished work

In a 7,000-person rollout, "people were working inside a Claude Project when what they actually needed was a Skill" [S22]. Read from inside, the wrong surface looks like a tool that cannot do the job.

[S4] What are Projects? — Anthropic. <https://support.claude.com/en/articles/9517075-what-are-projects> [S22] Non-technical AI adoption — Deel. <https://www.deel.com/deel-works/non-technical-ai-adoption>

Personalization is three separate places, not one switch

Anthropic's page names three, differing by how far each reaches: Instructions for Claude, Project instructions, Skills [S5].

- | | | |
|---|--------------------------------|---|
| 1 | Instructions for Claude | Account-wide: "applied to all of your conversations" [S5] |
| 2 | Project instructions | Only "apply to chats within that project" [S5] |
| 3 | Skills | Format and tone: "how Claude formats and delivers" [S5] |
- Put a preference at the widest place where it is always true, never wider. A tone rule for one campaign, set account-wide, follows you into unrelated chats.

[S5] Personalization features — Anthropic Help Center. <https://support.claude.com/en/articles/10185728-understanding-claude-s-personalization-features>

Shared project 2 — Personalization: set one preference, prove it changed something

- 1** **Open Personalization** All three places: Instructions for Claude, Project instructions, Skills [S5]
- 2** **Write one line** "Plain English, under 150 words, no bullets unless I ask."
- 3** **Test it now** Re-ask yesterday's question and read the difference
If the answer does not change, check scope [S5]: Instructions for Claude applies to all your conversations, Project instructions only to chats in that project, Skills to how a reply is formatted and delivered. Thirty seconds, instead of assuming it works.

[S5] *Personalization features — Anthropic Help Center.* <https://support.claude.com/en/articles/10185728-understanding-claude-s-personalization-features>

Shared project 3 — LinkedIn posts, into the feed Pangram measured as the most AI-saturated

40%+

of longform LinkedIn posts flagged fully
AI-generated, 1,002,627 analysed [S14]

−45%

engagement on likely-AI posts versus
likely-original, 2,726 posts [S15]

Both numbers come from companies selling AI detection, so treat them as a direction, not a measurement to bet on. The direction is enough: take the structure and the first draft from the project, and keep the specific — the customer, the number, the thing that happened — in your own words.

[S14] AI in your feed — Pangram Labs. <https://www.pangram.com/blog/ai-in-your-feed> [S15] AI content on LinkedIn — Originality.ai. <https://originality.ai/blog/ai-content-published-linkedin>

What not to put in a shared project

A shared project's instructions and documents are "visible to everyone with access" [S19].

Fine to upload

Published material, product pages, the tone-of-voice guide, an approved case study.

Not in a shared project

A customer contract, a quotation, a named complaint.

Two more hazards: two people editing one project "will overwrite each other" [S19], and stale files never announce themselves [S20]. If you are not certain your organisation's data policy covers customer material in Claude, ask before the exercise.

[S19] Claude Projects tutorial — EnterpriseDNA. <https://enterprisedna.co/resources/guides/guide-claude-projects-tutorial> [S20] Claude Projects for teams — OpenClawReady. <https://openclawready.com/blog/claude-projects-for-teams>

Build your own project in four steps

- 1 **Create it** "Replies to inbound enquiries", not "sales"
- 2 **Paste the paragraph** The background you retype, as instructions
- 3 **Upload two or three files** Not twenty. 30MB per project file [S6]
- 4 **Share it, if your plan has sharing** "Team and Enterprise plans only" [S4]
Check your own plan before promising a colleague access [S4].

[S4] What are Projects? — Anthropic Help Center. <https://support.claude.com/en/articles/9517075-what-are-projects> [S6] Upload files to Claude — Anthropic Help Center. <https://support.claude.com/en/articles/8241126-upload-files-to-claude>

Instructions describe behaviour; documents carry the content

"Write instructions that describe behavior, not content" — a Q3 strategy summary is content, so it is a document [S19].

Why that is not cosmetic: as your project knowledge approaches how much text Claude can read at one go, the project stops loading everything and starts searching instead [S7]. Nothing on screen announces the switch.

Anthropic's guidance is the smallest possible set of high-signal information [S8]. So more documents is not more accurate answers: replace a superseded document rather than leaving it beside the new one.

[S7] RAG for projects — Anthropic Help Center. <https://support.claude.com/en/articles/11473015-retrieval-augmented-generation-rag-for-projects> [S19] Claude Projects tutorial — EnterpriseDNA. <https://enterprisedna.co/resources/guides/guide-claude-projects-tutorial>

Starting from a slide template — and the file type that will stop you

Accepted

"PDF, DOCX, CSV, TXT, HTML, ODT, RTF, EPUB, JSON, XLSX" [S6]

Not on that list

PPTX. The template file itself will not upload

Instead

Export the template to PDF, or type the slide structure into the instructions

What comes back is the words for each slide against your standing structure. You still assemble the branded file in PowerPoint, which is fine: the writing is the part that took the time.

[S6] Upload files to Claude — Anthropic Help Center. <https://support.claude.com/en/articles/8241126-upload-files-to-claude>

Three million custom GPTs were built; about one in twenty is still listed publicly

3,000,000+

custom GPTs built on OpenAI's earlier version of this feature [S16]

~159,000

still public and active in the GPT Store two years later [S17]

Creating one is free; keeping it correct is not, and nothing tells you when a project has gone stale. One practitioner write-up names the fix, and it is small — "an owner and a review date" [S20], written into the project's own instructions.

[S16] *Introducing the GPT Store* — OpenAI (2024). <https://openai.com/index/introducing-the-gpt-store> [S17] *GPT Store business guide 2026* — DigitalApplied. <https://www.digitalapplied.com/blog/gpt-store-custom-gpts-business-guide-2026>

Hands-on — pick one and do it now

- 1 **Rewrite one real prompt** Reader, purpose, what to avoid
- 2 **Set one preference, test it** Re-ask yesterday's question
- 3 **Create one project** Paste the paragraph you retype
Use your own work, not mine: employees "can't transfer a generic demo onto their actual Tuesday" [S18]. For novices, van Gog and colleagues predicted worked examples produce lower cognitive load and better learning than problem-solving alone [S21].

[S18] Enterprise Claude adoption — Correlation One. <https://www.correlation-one.com/blog/enterprise-claude-adoption-7-real-problems-and-the-solutions-that-worked> [S21] Worked examples — van Gog et al. <https://www.sciencedirect.com/science/article/abs/pii/S0361476X1000055X>

Takeaways — what survives to next Tuesday

- | | | |
|---|------------------------------------|--------------------------------------|
| 1 | Wording or facts | Before every prompt |
| 2 | Role prompting | No accuracy gain [S1]; wording only |
| 3 | Ask, then one message | Consolidate; do not dribble [S2] |
| 4 | One check step, every draft | Nothing signals a wrong answer [S10] |
| 5 | An owner and a review date | In the project instructions |
- One practitioner: unowned projects answer from stale files [S20].

[S10] AI hallucinations — IBM. <https://www.ibm.com/think/topics/ai-hallucinations> [S20] Claude Projects for teams — OpenClawReady. <https://openclawready.com/blog/claude-projects-for-teams>

References · 1/5

- **[S1]** Playing Pretend: Expert Personas, Prompting Science Report 4 — Wharton Generative AI Labs (2025). <https://gail.wharton.upenn.edu/research-and-insights/playing-pretend-expert-personas/>
- **[S2]** LLMs Get Lost In Multi-Turn Conversation — Laban et al. (2025), arXiv:2505.06120. <https://arxiv.org/abs/2505.06120>
- **[S3]** Be clear and direct — Anthropic. <https://docs.claude.com/en/docs/build-with-claude/prompt-engineering/be-clear-and-direct>
- **[S4]** What are Projects? — Anthropic Help Center. <https://support.claude.com/en/articles/9517075-what-are-projects>
- **[S5]** Understanding Claude's personalization features — Anthropic Help Center. <https://support.claude.com/en/articles/10185728-understanding-claude-s-personalization-features>

References · 2/5

- **[S6]** Upload files to Claude — Anthropic Help Center. <https://support.claude.com/en/articles/8241126-upload-files-to-claude>
- **[S7]** Retrieval augmented generation (RAG) for projects — Anthropic Help Center. <https://support.claude.com/en/articles/11473015-retrieval-augmented-generation-rag-for-projects>
- **[S8]** Effective context engineering for AI agents — Anthropic. <https://www.anthropic.com/engineering/effective-context-engineering-for-ai-agents>
- **[S9]** Claude Cowork — Anthropic. <https://www.anthropic.com/product/claude-cowork>
- **[S10]** What are AI hallucinations? — IBM. <https://www.ibm.com/think/topics/ai-hallucinations>
- **[S11]** To help improve the accuracy of generative AI, add speed bumps — MIT Sloan. <https://mitsloan.mit.edu/ideas-made-to-matter/to-help-improve-accuracy-generative-ai-add-speed-bumps>

References · 3/5

- **[S12]** Chatbot use by non-IT laypersons: a PRISMA review of 28 studies — Universal Access in the Information Society (2026). <https://link.springer.com/article/10.1007/s10209-026-01323-w>
- **[S13]** AI use makes us overestimate our cognitive performance — Aalto University on Fernandes et al. (2026). <https://www.aalto.fi/en/news/ai-use-makes-us-overestimate-our-cognitive-performance>
- **[S14]** AI in your feed — Pangram Labs. <https://www.pangram.com/blog/ai-in-your-feed>
- **[S15]** AI content published on LinkedIn — Originality.ai. <https://originality.ai/blog/ai-content-published-linkedin>
- **[S16]** Introducing the GPT Store — OpenAI (2024). <https://openai.com/index/introducing-the-gpt-store>

References · 4/5

- **[S17]** GPT Store and custom GPTs: business guide 2026 — DigitalApplied. <https://www.digitalapplied.com/blog/gpt-store-custom-gpts-business-guide-2026>
- **[S18]** Enterprise Claude adoption: 7 real problems and the solutions that worked — Correlation One. <https://www.correlation-one.com/blog/enterprise-claude-adoption-7-real-problems-and-the-solutions-that-worked>
- **[S19]** Claude Projects tutorial — EnterpriseDNA. <https://enterprisedna.co/resources/guides/guide-claude-projects-tutorial>
- **[S20]** Claude Projects for teams — OpenClawReady. <https://openclawready.com/blog/claude-projects-for-teams>

References · 5/5

- **[S21]** Effects of worked examples, example-problem, and problem-example pairs on novices' learning — van Gog, Kester & Paas. <https://www.sciencedirect.com/science/article/abs/pii/S0361476X1000055X>
- **[S22]** Non-technical AI adoption — Deel. <https://www.deel.com/deel-works/non-technical-ai-adoption>
- **[S23]** The Prompt Report: A Systematic Survey of Prompting Techniques — Schulhoff et al. (2024), arXiv:2406.06608. <https://arxiv.org/abs/2406.06608>
- **[V1]** Role Prompting — author's working notes. Summarised in the printed appendix.
- **[V2]** State of AI in Business 2025 (MIT NANDA) — Overview — author's working notes. Summarised in the printed appendix.

The handout PDF carries the full reference list, and a one-page summary of every [V] source with its own external references.

TAKE IT WITH YOU

Everything from this talk

[Slides \(PDF\)](#)

[Handout \(PDF\)](#)

[Speaker notes \(PDF\)](#)

[Everything \(.zip\) — readable by an AI agent](#)

ondrasek.github.io/talks

me@ondrejkrasicek.com