

Claude on the web: habits, projects, and checking

Four habits and three projects already in your account are enough to do real work with Claude this week; building one of your own is what makes it repeatable rather than a one-off.

Ondrej (Ondra) Krajicek

me@ondrejkrjicek.com · linkedin.com/in/OndrejKrajicek

Built 2026-09-11 06:13 UTC

Generated by the agentic vault — an autonomous research and authoring harness. Every factual claim was gated against a cited source registry before this document was produced.

Agenda — four habits, three shared projects, one thing you build yourself

- 1 Four prompting habits** Brief it [S3]; role: wording not accuracy [S1] [V1]; answer in one message [S2]; check a fact [S10]
 - 2 Three shared projects** Prompt Engineering, Personalization, LinkedIn posts
 - 3 Your own project** Create and share it [S4]; owner and review date [S20]
 - 4 Eight minutes hands-on** Practise on your own recurring job [S18]
- Not covered: multi-step document work — Anthropic sells Cowork [S9].

[S9] Claude Cowork — Anthropic. <https://www.anthropic.com/product/claude-cowork> [S18] Enterprise Claude adoption — Correlation One. <https://www.correlation-one.com/blog/enterprise-claude-adoption-7-real-problems-and-the-solutions-that-worked>

WRITE THESE DOWN

Takeaways — four sentences, before we start

- | | | |
|---|------------------------------|---|
| 1 | Wording or facts | Reword, or add a document |
| 2 | Role prompting | No significant accuracy gain: GPQA Diamond, MMLU-Pro exams [S1] |
| 3 | Ask, then consolidate | Drip-feeding: "write a post", "add our pricing" — 39% average performance drop vs. one message [S2] |
| 4 | One check step | Nothing signals a wrong answer [S10] |

[S1] *Playing Pretend* — Wharton GAIL (2025). <https://gail.wharton.upenn.edu/research-and-insights/playing-pretend-expert-personas/> [S2] *LLMs Get Lost In Multi-Turn Conversation* — Laban et al. (2025). <https://arxiv.org/abs/2505.06120> [S10] *AI hallucinations* — IBM. <https://www.ibm.com/think/topics/ai-hallucinations>

You have retyped the same background paragraph into a new chat this week

The paragraph you retype into every new chat — what we sell, what we never promise — because Claude starts each one blank.

Already using AI
"Employees at 90% of companies regularly use AI tools at a personal level" — contested [V2]

One client team
"Access is not the problem — the bottleneck is fluency" [S18]

That team's "fluency" is three gaps [S18]: **which tool, how to ask, how to check.**

[S18] Enterprise Claude adoption — Correlation One. <https://www.correlation-one.com/blog/enterprise-claude-adoption-7-real-problems-and-the-solutions-that-worked> [V2] State of AI in Business 2025 — author's notes.

One question before every prompt: short of wording, or short of facts?

Short of wording

True, but reads wrong — too long, too formal, wrong reader.

Short of facts

Fluent and wrong: Claude never had our pricing, our positioning, our customer.

The test: could a colleague who has never seen our files write this answer? If yes, it is a **wording** problem — fix the prompt. If no, Claude is **short of facts**, and no rewording invents them. The fix is the right document — Anthropic's rule is the smallest high-signal set, not every file you have [S8].

[S8] *Effective context engineering* — Anthropic. <https://www.anthropic.com/engineering/effective-context-engineering-for-ai-agents>

Habit 1 — brief Claude the way you brief a new colleague

Anthropic: "a brilliant but new employee who lacks context on your norms and workflows" [S3] — capable, and holding no facts about us.

Anthropic's test costs nothing: give the prompt to a colleague who does not know the task [S3]. Whatever they would have to ask you is what the prompt is missing.

Who reads it A procurement manager who has never bought from us

What it is for A reply that wins a meeting, not a sale

What to avoid No pricing, no dates

[S3] Be clear and direct — Anthropic. <https://docs.claude.com/en/docs/build-with-claude/prompt-engineering/be-clear-and-direct>

Habit 2 — "act as an expert" bought no accuracy gain in the models Wharton tested

Role prompting: naming a role, Wharton's "expert persona" [S1].

no significant

accuracy gain from role prompting —
six models, two knowledge exams:
GPQA Diamond, MMLU-Pro [S1]

No Claude model is among the six [S1], and persona effects differ sharply by model: on one labelling task, 11 of 21 personas raised one model's accuracy, 16 of 21 lowered another's [S24]. Expecting the same of Claude is my read. A role moves wording — measured next.

[S1] *Playing Pretend* — Wharton GAIL (2025). <https://gail.wharton.upenn.edu/research-and-insights/playing-pretend-expert-personas/> [S24] *Persona prompting, social reasoning* — Yang et al. (2026). <https://aclanthology.org/2026.eacl-long.52/>

What adding a role changed: rated depth up, rated clarity down, no overall gain

Another AI model scored answers to 1,140 questions [V1].

**sounds
more
expert**

rated expertise depth 3.638 → 3.923 —
a rise, not a fluke [V1]

**reads less
clearly**

rated clarity 4.896 → 4.550 — a fall, not
a fluke [V1]

No overall gain: the two effects roughly cancelled [V1]. My note gives no rating scale: read direction, not size. The study picked roles automatically; you type yours [V1] — the trade surviving that is my judgement. So: role for wording, never accuracy.

Habit 3 — let Claude ask, then restate everything in one message

- 1 **Ask for the questions** "Before you write anything, ask me the five questions you need answered."
- 2 **Answer in one message** Not five replies across five turns

Why one message: drip-feeding — "write a post", then "make it shorter", then "add our pricing" — is the trap. Across 200,000+ simulated conversations, spreading the same information over turns instead of one message cost a **39% average performance drop**, and one early wrong turn was never recovered [S2].

[S2] LLMs Get Lost In Multi-Turn Conversation — Laban et al. (2025). <https://arxiv.org/abs/2505.06120>

Habit 4 — one deliberate check between the draft and the send button

Claude produces a true sentence and a hallucination by the same mechanism [S10] — nothing signals a wrong answer.

- Read the number off the source page, or ask whoever owns it — not Claude.
- Delete what you cannot trace to a page or person.

MIT Sloan, 140 professionals: colour-coding AI output so it had to be read, not skimmed. The heaviest colour-coding caught more errors and omissions but took longer; the medium setting balanced detection against time [S11]. So: the costliest fact, not an audit.

[S10] AI hallucinations — IBM. <https://www.ibm.com/think/topics/ai-hallucinations> [S11] Add speed bumps — MIT Sloan. <https://mitsloan.mit.edu/ideas-made-to-matter/to-help-improve-accuracy-generative-ai-add-speed-bumps>

The four habits — and a study where higher AI literacy came with more overconfidence

- | | | |
|---|------------------------------|----------------------------------|
| 1 | Brief it | Reader, purpose, what to avoid |
| 2 | Role for wording | Not for accuracy [S1] |
| 3 | Ask, then consolidate | One message, not five turns [S2] |
| 4 | One check step | The costliest fact [S10] |

Across 500 participants, "higher AI literacy brings more overconfidence" [S13] — never measured on a room leaving a training session, so carrying it here is my step. Habit 4 alone puts a person between draft and reader.

[S13] AI and overconfidence — Aalto University (2026). <https://www.aalto.fi/en/news/ai-use-makes-us-overestimate-our-cognitive-performance>

A project holds the background paragraph, so you stop retyping it

Anthropic's definition: "self-contained workspaces with their own chat histories and knowledge bases" [S4].

Without one
You paste who we are and what
we never promise into every new
chat, slightly differently each time.

With one
Write it once as instructions,
upload the documents once, and
every chat inside starts already
briefed.

This is the facts half of the wording-or-facts question — not a better way to word things, but a place to put what Claude could not otherwise have.

[S4] What are Projects? — Anthropic Help Center. <https://support.claude.com/en/articles/9517075-what-are-projects>

Shared project 1 — Prompt Engineering, for when you are short of wording

Where

Open your Projects list; sharing needs Team or Enterprise [S4]

Use it when

You cannot get Claude to word what you want. Not when facts are wrong

What comes back

Treat it as a draft prompt, not as finished work

Deel's 7,000-person rollout reported people working inside a Project when they needed a Skill — another Claude feature, covered next [S22]. From inside, the wrong place looks like a tool that cannot do the job.

[S4] What are Projects? — Anthropic. <https://support.claude.com/en/articles/9517075-what-are-projects> [S22] Non-technical AI adoption — Deel. <https://www.deel.com/deel-works/non-technical-ai-adoption>

Personalization is three separate places, not one switch

Anthropic's page names three: two differ by how far each reaches, the third by what it does [S5].

- | | | |
|---|--------------------------------|--|
| 1 | Instructions for Claude | Account-wide: "all of your conversations" [S5] |
| 2 | Project instructions | Only "apply to chats within that project" [S5] |
| 3 | Skills | Formatting: "how Claude formats and delivers" [S5] |

For the first two: put a preference at the widest place where it is always true, never wider. A campaign tone rule set account-wide follows you into holiday chats.

[S5] Personalization features — Anthropic Help Center. <https://support.claude.com/en/articles/10185728-understanding-claude-s-personalization-features>

Shared project 2 — Personalization: set one preference, prove it changed something

- 1** **Open Personalization** All three places: Instructions for Claude, Project instructions, Skills [S5]
- 2** **Write one line** "Plain English, under 150 words, no bullets unless I ask."
- 3** **Test it now** Re-ask yesterday's question and read the difference
If the answer does not change, check which of the three it went in [S5]: Instructions for Claude covers all your conversations, Project instructions only that project's chats, and Skills set formatting, not reach. Thirty seconds, instead of assuming it works.

[S5] *Personalization features — Anthropic Help Center.* <https://support.claude.com/en/articles/10185728-understanding-claude-s-personalization-features>

Shared project 3 — LinkedIn posts, into the feed Pangram measured as the most AI-saturated

40%+

of longform LinkedIn posts flagged fully
AI-generated, 1,002,627 analysed [S14]

−45%

engagement on likely-AI posts versus
likely-original, 2,726 posts [S15]

Both numbers come from companies selling AI detection, so treat them as a direction, not a measurement to bet on. The direction is enough: take the structure and the first draft from the project, and keep the specific — the customer, the number, the thing that happened — in your own words.

[S14] AI in your feed — Pangram Labs. <https://www.pangram.com/blog/ai-in-your-feed> [S15] AI content on LinkedIn — Originality.ai. <https://originality.ai/blog/ai-content-published-linkedin>

What not to put in a shared project

A shared project's instructions and documents are "visible to everyone with access" [S19].

Fine to upload

Published material, product pages, the tone-of-voice guide, an approved case study.

Not in a shared project

A customer contract, a quotation, a named complaint.

Two more hazards: two people editing one project "will overwrite each other" [S19], and stale files never announce themselves [S20]. Check what your organisation's data policy says about customer material in Claude before the exercise.

[S19] Claude Projects tutorial — EnterpriseDNA. <https://enterprisedna.co/resources/guides/guide-claude-projects-tutorial> [S20] Claude Projects for teams — OpenClawReady. <https://openclawready.com/blog/claude-projects-for-teams>

Build your own project in four steps

- 1 Create it** "Replies to inbound enquiries", not "sales"
- 2 Paste the paragraph** The background you retype, as instructions
- 3 Upload two or three files** Not twenty. 30MB per project file [S6]
- 4 Share it, if your plan allows** Team and Enterprise only [S4]; "Can view" or "Can edit"
Check your plan before promising access [S4].

[S4] What are Projects? — Anthropic Help Center. <https://support.claude.com/en/articles/9517075-what-are-projects> [S6] Upload files to Claude — Anthropic Help Center. <https://support.claude.com/en/articles/8241126-upload-files-to-claude>

Instructions describe behaviour; documents carry the content

"Write instructions that describe behavior, not content" — a Q3 strategy summary is content, so it is a document [S19].

Why that is not cosmetic: as your project knowledge approaches how much text Claude can read at one go, the project stops loading everything and starts searching instead [S7]. Nothing on screen announces the switch.

Anthropic's guidance is the smallest possible set of high-signal information [S8]. So more documents is not more accurate answers: replace a superseded document rather than leaving it beside the new one.

[S7] RAG for projects — Anthropic Help Center. <https://support.claude.com/en/articles/11473015-retrieval-augmented-generation-rag-for-projects> [S19] Claude Projects tutorial — EnterpriseDNA. <https://enterprisedna.co/resources/guides/guide-claude-projects-tutorial>

Starting from a slide template — and the file type that will stop you

Accepted

"PDF, DOCX, CSV, TXT, HTML, ODT, RTF, EPUB, JSON, XLSX" [S6]

Not on that list

PPTX. The template file itself will not upload

Instead

Export the template to PDF, or type the slide structure into the instructions

What comes back is the words for each slide against your standing structure. You still assemble the branded file in PowerPoint, which is fine: the writing is the part that took the time.

[S6] Upload files to Claude — Anthropic Help Center. <https://support.claude.com/en/articles/8241126-upload-files-to-claude>

Three million custom GPTs were built; roughly 159,000 are still listed publicly

3,000,000+

custom GPTs built on OpenAI's earlier version of this feature [S16]

~159,000

still public and active in the GPT Store two years later [S17]

About one in twenty — my arithmetic, not a ratio either source reports. Creating one is free; keeping it correct is not. The fix a practitioner write-up names is small: "an owner and a review date" [S20], in the project's instructions.

[S16] *Introducing the GPT Store* — OpenAI (2024). <https://openai.com/index/introducing-the-gpt-store> [S17] *GPT Store business guide 2026* — DigitalApplied. <https://www.digitalapplied.com/blog/gpt-store-custom-gpts-business-guide-2026>

Hands-on — pick one and do it now

- 1 **Rewrite one real prompt** Reader, purpose, what to avoid
- 2 **Set one preference, test it** Re-ask yesterday's question
- 3 **Create one project** Paste the paragraph you retype
Use your own work, not mine: employees "can't transfer a generic demo onto their actual Tuesday" [S18]. For novices, van Gog and colleagues predicted worked examples produce lower cognitive load and better learning than problem-solving alone [S21].

[S18] Enterprise Claude adoption — Correlation One. <https://www.correlation-one.com/blog/enterprise-claude-adoption-7-real-problems-and-the-solutions-that-worked> [S21] Worked examples — van Gog et al. <https://www.sciencedirect.com/science/article/abs/pii/S0361476X1000055X>

What this adds up to — the habits and the three projects do real work this week; your own project is what makes it repeat

- | | | |
|---|---|--|
| 1 | Habits: the wording half | Brief it, role for wording, one message, one check [S1] [S2] [S10] |
| 2 | Projects: the facts half | Prompt Engineering, Personalization, LinkedIn posts, already there |
| 3 | Your own project — owner and review date | The paragraph you retype, written once; stale unless owned [S20] |

[S1] *Playing Pretend* — Wharton GAIL (2025). <https://gail.wharton.upenn.edu/research-and-insights/playing-pretend-expert-personas/> [S2] *Lost in multi-turn conversation* — Laban et al. (2025). <https://arxiv.org/abs/2505.06120> [S10] *AI hallucinations* — IBM. <https://www.ibm.com/think/topics/ai-hallucinations> [S20] *Claude Projects for teams* — OpenClawReady. <https://openclawready.com/blog/claude-projects-for-teams>

References · 1/6

- **[S1]** Playing Pretend: Expert Personas, Prompting Science Report 4 — Wharton Generative AI Labs (2025). <https://gail.wharton.upenn.edu/research-and-insights/playing-pretend-expert-personas/>
- **[S2]** LLMs Get Lost In Multi-Turn Conversation — Laban et al. (2025), arXiv:2505.06120. <https://arxiv.org/abs/2505.06120>
- **[S3]** Be clear and direct — Anthropic. <https://docs.claude.com/en/docs/build-with-claude/prompt-engineering/be-clear-and-direct>
- **[S4]** What are Projects? — Anthropic Help Center. <https://support.claude.com/en/articles/9517075-what-are-projects>
- **[S5]** Understanding Claude's personalization features — Anthropic Help Center. <https://support.claude.com/en/articles/10185728-understanding-claude-s-personalization-features>

References · 2/6

- **[S6]** Upload files to Claude — Anthropic Help Center. <https://support.claude.com/en/articles/8241126-upload-files-to-claude>
- **[S7]** Retrieval augmented generation (RAG) for projects — Anthropic Help Center. <https://support.claude.com/en/articles/11473015-retrieval-augmented-generation-rag-for-projects>
- **[S8]** Effective context engineering for AI agents — Anthropic. <https://www.anthropic.com/engineering/effective-context-engineering-for-ai-agents>
- **[S9]** Claude Cowork — Anthropic. <https://www.anthropic.com/product/claude-cowork>
- **[S10]** What are AI hallucinations? — IBM. <https://www.ibm.com/think/topics/ai-hallucinations>
- **[S11]** To help improve the accuracy of generative AI, add speed bumps — MIT Sloan. <https://mitsloan.mit.edu/ideas-made-to-matter/to-help-improve-accuracy-generative-ai-add-speed-bumps>

References · 3/6

- **[S12]** Chatbot use by non-IT laypersons: a PRISMA review of 28 studies — Universal Access in the Information Society (2026). <https://link.springer.com/article/10.1007/s10209-026-01323-w>
- **[S13]** AI use makes us overestimate our cognitive performance — Aalto University on Fernandes et al. (2026). <https://www.aalto.fi/en/news/ai-use-makes-us-overestimate-our-cognitive-performance>
- **[S14]** AI in your feed — Pangram Labs. <https://www.pangram.com/blog/ai-in-your-feed>
- **[S15]** AI content published on LinkedIn — Originality.ai. <https://originality.ai/blog/ai-content-published-linkedin>
- **[S16]** Introducing the GPT Store — OpenAI (2024). <https://openai.com/index/introducing-the-gpt-store>

References · 4/6

- **[S17]** GPT Store and custom GPTs: business guide 2026 — DigitalApplied. <https://www.digitalapplied.com/blog/gpt-store-custom-gpts-business-guide-2026>
- **[S18]** Enterprise Claude adoption: 7 real problems and the solutions that worked — Correlation One. <https://www.correlation-one.com/blog/enterprise-claude-adoption-7-real-problems-and-the-solutions-that-worked>
- **[S19]** Claude Projects tutorial — EnterpriseDNA. <https://enterprisedna.co/resources/guides/guide-claude-projects-tutorial>
- **[S20]** Claude Projects for teams — OpenClawReady. <https://openclawready.com/blog/claude-projects-for-teams>

References · 5/6

- **[S21]** Effects of worked examples, example-problem, and problem-example pairs on novices' learning — van Gog, Kester & Paas. <https://www.sciencedirect.com/science/article/abs/pii/S0361476X1000055X>
- **[S22]** Non-technical AI adoption — Deel. <https://www.deel.com/deel-works/non-technical-ai-adoption>
- **[S23]** The Prompt Report: A Systematic Survey of Prompting Techniques — Schulhoff et al. (2024), arXiv:2406.06608. <https://arxiv.org/abs/2406.06608>
- **[S24]** Persona Prompting as a Lens on LLM Social Reasoning — Yang, Hechtbauer, Khalilov, Brinkmann, Schmitt & Feldhus, EACL 2026 (main), pp. 1152–1170. <https://aclanthology.org/2026.eacl-long.52/>

References · 6/6

- **[V1]** Role Prompting — author's working notes. Summarised in the printed appendix.
- **[V2]** State of AI in Business 2025 (MIT NANDA) — Overview — author's working notes. Summarised in the printed appendix.

The handout PDF carries the full reference list, and a one-page summary of every [V] source with its own external references.

TAKE IT WITH YOU

Everything from this talk

[Slides \(PDF\)](#)

[Handout \(PDF\)](#)

[Speaker notes \(PDF\)](#)

[Everything \(.zip\) — readable by an AI agent](#)

ondrasek.github.io/talks

me@ondrejkrasicek.com