

Speaker Notes

Contents

- [1 1. Claude on the web: habits, projects, and checking](#)
- [2 2. Agenda — four habits, three shared projects, one thing you build yourself](#)
- [3 3. Takeaways — four sentences, before we start](#)
- [4 4. You have retyped the same background paragraph into a new chat this week](#)
- [5 5. One question before every prompt: short of wording, or short of facts?](#)
- [6 6. Habit 1 — brief Claude the way you brief a new colleague](#)
- [7 7. Habit 2 — "act as an expert" does not make Claude more correct](#)
- [8 8. What adding a role changed: rated depth up, rated clarity down, no overall gain](#)
- [9 9. Habit 3 — let Claude ask, then restate everything in one message](#)
- [10 10. Habit 4 — one deliberate check between the draft and the send button](#)
- [11 11. The four habits — and a study where higher AI literacy came with more overconfidence](#)
- [12 12. A project holds the background paragraph, so you stop retyping it](#)
- [13 13. Shared project 1 — Prompt Engineering, for when you are short of wording](#)
- [14 14. Personalization is three separate places, not one switch](#)
- [15 15. Shared project 2 — Personalization: set one preference, prove it changed something](#)
- [16 16. Shared project 3 — LinkedIn posts, into the feed Pangram measured as the most AI-saturated](#)
- [17 17. What not to put in a shared project](#)
- [18 18. Build your own project in four steps](#)
- [19 19. Instructions describe behaviour; documents carry the content](#)
- [20 20. Starting from a slide template — and the file type that will stop you](#)
- [21 21. Three million custom GPTs were built; about one in twenty is still listed publicly](#)
- [22 22. Hands-on — pick one and do it now](#)
- [23 23. Takeaways — what survives to next Tuesday](#)
- [24 24. References](#)

Generated by the agentic vault — an autonomous research and authoring harness. Every factual claim was gated against a cited source registry before this document was produced.

Slides: 24 · **Duration:** 25 min

11. Claude on the web: habits, projects, and checking

On screen:

AI literacy session · 25 minutes, 8 of them yours

Four habits and three projects already in your account are enough to do real work with Claude this week; building one of your own is what makes it repeatable rather than a one-off. [speculative]

Ondrej (Ondra) Krajicek me@ondrejkrjicek.com · linkedin.com/in/OndrejKrajicek Built at render time

Twenty-five minutes, and eight of them are yours. You will have a keyboard open at the end, and I would rather you spend that time than watch me. [speculative]

Three things: four habits for writing a request to Claude, three projects that already sit in your Projects list — Prompt Engineering, Personalization, and LinkedIn posts — and how to build one of your own. [speculative]

None of it needs you to be technical. The first habit is something you already do every time you brief a colleague. [speculative]

2 2. Agenda — four habits, three shared projects, one thing you build yourself

On screen:

Roadmap

1

Four prompting habits

Brief it [S3]; role: wording not accuracy [S1] [V1]; answer in one message [S2]; check a fact [S10]

2

*Three shared projects**Prompt Engineering, Personalization, LinkedIn posts*

3

*Your own project**Create and share it [S4]; owner and review date [S20]*

4

*Eight minutes hands-on**Practise on your own recurring job [S18]**Not covered: multi-step document work — Anthropic sells Cowork [S9].**[S9] Claude Cowork — Anthropic. <https://www.anthropic.com/product/claude-cowork>**[S18] Enterprise Claude adoption — Correlation One. <https://www.correlation-one.com/blog/enterprise-claude-adoption-7-real-problems-and-the-solutions-that-worked>*

Four habits, then the three projects, then your own, then you practise. [inference]

The last row looks like I am giving you less, so let me defend it. Correlation One run these rollouts for a living, and their finding was that access is not the problem — people had logins and still could not turn a demo into their own Tuesday morning [S18].

One thing I am not covering. If your work is a long multi-step document job, Anthropic has a separate product for that called Cowork, built for non-coding knowledge work [S9]. Today is the chat window and projects.

3 3. Takeaways — four sentences, before we start

On screen:

Write these down

1

*Wording or facts**Rewrite the prompt, or add a document [speculative]*

2

Role prompting

No accuracy gain: six models, two standard exams [S1]

3

Ask, then consolidate

Spread over turns: 39% lower average performance than one message [S2]

4

One check step

No signal fires when Claude is wrong [S10]

[S1] Playing Pretend — Wharton GAIL (2025). <https://gail.wharton.upenn.edu/research-and-insights/playing-pretend-expert-personas/> [S2] LLMs Get Lost In Multi-Turn Conversation — Laban et al. (2025). <https://arxiv.org/abs/2505.06120> [S10] AI hallucinations — IBM. <https://www.ibm.com/think/topics/ai-hallucinations>

Conclusions before the argument. If you catch one slide, this one. [speculative]

One: before every prompt, ask whether you are short of wording or short of facts. Wording you fix in the prompt. Facts you fix by giving Claude the document — that is what a project is for.

Two: telling Claude it is an expert does not make it more correct. Wharton tested that across six models on two benchmarks: no significant positive difference on accuracy [S1].

Three: let Claude ask questions, then put every answer in one message. Two hundred thousand simulated conversations measured a thirty-nine percent drop in average performance when the same information arrived spread across turns instead of in one message [S2].

Four: one check between the draft and the send button, every time, because Claude gives no signal when it is wrong [S10].

4 4. You have retyped the same background paragraph into a new chat this week

On screen:

Part I · Habits

The same paragraph — what we sell, what we never promise — pasted into every new chat, because Claude starts each one knowing nothing about us. [speculative]

Already using AI

"Employees at 90% of companies regularly use AI tools at a personal level" — contested [V2]

What is missing

"Access is not the problem — the bottleneck is fluency" [S18]

"Fluency" here is three specific gaps [S18]: knowing which tool fits the task, how to ask so the answer is usable, and how to check it before you trust it. This talk is those three habits.

[S18] Enterprise Claude adoption — Correlation One. <https://www.correlation-one.com/blog/enterprise-claude-adoption-7-real-problems-and-the-solutions-that-worked> [V2]

State of AI in Business 2025 — author's notes.

Hands up if you opened a new Claude chat this week and typed out who we are and what we sell, because a new chat starts with none of it. Hold on to that paragraph — it is what this session fixes. [speculative]

Most of you already use something. The MIT NANDA survey put it at ninety percent of companies, and my notes file that exact figure under the report's contested findings, so take it as a rough shape rather than a count [V2]. You are not starting from zero; you are starting from unofficial.

And that is the gap. Correlation One found access was never the problem — the bottleneck was fluency [S18]. The useful half is what they name as missing: most of the team did not know which tool or surface to use for which task, how to query effectively, or how to verify an output before relying on it [S18]. Three things — and I give you one habit for the first two and a whole habit for the third. Their other line is the test I am setting for today: employees "can't transfer a generic demo onto their actual Tuesday" [S18], so bring your own Tuesday to the last eight minutes. [speculative]

5 5. One question before every prompt: short of wording, or short of facts?

On screen:

Part I · Habits

*Short of wording**True, but reads wrong — too long, too formal, wrong reader. Fix it in the prompt.**Short of facts**Fluent and wrong, because Claude never had our pricing, our positioning, our customer. Fix it with a document.**The test: could a smart colleague who has never seen our files write this answer? If yes, it is a wording problem — fix the prompt. If no, Claude is short of facts it was never given, and no rewording invents them [speculative]. Then the fix is a document you upload, not a better sentence [S8]. [inference]**[S8] Effective context engineering — Anthropic. <https://www.anthropic.com/engineering/effective-context-engineering-for-ai-agents>*

Every disappointing answer is one of two failures. Either the content is right and the writing is wrong — too long, too stiff, wrong reader. That is wording, and you fix it by changing how you ask.

Or the writing is lovely and the content is wrong, because Claude was never given our price list or our positioning. No amount of rewording touches that. Claude cannot select a fact it was never given — my framing of the two failures, not something a source here measured. [speculative]

Anthropic's guidance is the smallest set of high-signal information that gets the outcome [S8]. Smallest matters.

So: wording, or facts. The habits are the wording half. Projects are the facts half.

6 6. Habit 1 — brief Claude the way you brief a new colleague

On screen:

*Part I · Habits**Anthropic: "a brilliant but new employee who lacks context on your norms and workflows" [S3] — capable, and holding no facts about us. [inference]**Anthropic's test costs nothing: give the prompt to a colleague who does not know the task [S3]. Whatever they would have to ask you is what the prompt is missing. [inference]*

Who reads it

A procurement manager who has never bought from us

What it is for

A reply that wins a meeting, not a sale

What to avoid

No pricing, no dates

[S3] Be clear and direct — Anthropic. <https://docs.claude.com/en/docs/build-with-claude/prompt-engineering/be-clear-and-direct>

Habit one needs no new vocabulary, because you already do it. You brief people all the time. [speculative]

Anthropic's own words: a brilliant but new employee who lacks context on your norms and workflows [S3].

Their test costs nothing. Give the prompt to a colleague who does not know the task and ask them to follow it [S3]. Whatever that person has to ask you is exactly what the prompt left out. [inference]

So three things go into almost every prompt: who reads it, what it must achieve, and what must not be in it. Everybody skips the third, and it is the one that keeps a delivery date out of a draft.

7.7. Habit 2 — "act as an expert" does not make Claude more correct

On screen:

Part I · Habits

Role prompting means naming a role — an "expert persona", in Wharton's words [S1].

no significant

accuracy gain from role prompting — six models, two standard exams (GPQA Diamond, MMLU-Pro) [S1]

So telling Claude it is a senior engineer does not make it more correct [inference] —

Wharton's pre-registered analysis found no significant gain over no role at all [S1].

What it moves is the wording, measured next.

[S1] *Playing Pretend* — Wharton GAIL (2025). <https://gail.wharton.upenn.edu/research-and-insights/playing-pretend-expert-personas/>

This is the habit most of you have, and I want to correct it rather than remove it. [speculative]

Wharton tested six models on two standard knowledge tests, role prompting against none — expert personas is their word for it — with the analysis fixed in advance, so the result could not be chosen after the fact. No significant positive difference in accuracy [S1] — significant meaning the difference could not be told apart from chance. Telling Claude it is a senior engineer does not make it right about engineering. [inference]

8 8. What adding a role changed: rated depth up, rated clarity down, no overall gain

On screen:

Part I · Habits

Another AI model scored the answers to 1,140 questions [V1].

sounds more expert

expertise depth 3.638 $\rightarrow$ 3.923, Cohen's $d = 0.54$ [V1]

reads less clearly

clarity 4.896 $\rightarrow$ 4.550, $p < 0.001$ [V1]

No overall gain: the two effects roughly cancelled [V1]. My note gives no rating scale: read direction, not size [inference]. The study picked the role automatically; you type yours [V1]. That the trade survives typing is my judgement. [speculative] Use role prompting for wording, never for accuracy. [speculative]

[V1] Role Prompting — author's notes.

Role prompting does move something, and here I have to be careful about what was actually measured. In a study of eleven hundred and forty questions across thirty-eight expert roles, another AI model did the scoring: rated expertise depth rose from three point six four to three point nine two, an effect size of nought point five four, which is moderate; rated clarity fell from four point nine oh to four point five five. Both differ-

ences at p below nought point nought nought one [V1]. And the limit, which my own appendix carries: in aggregate the two effects roughly cancelled, with no significant overall difference from answers written with no role at all [V1].

The technique it tested was hybrid role retrieval — the study picked a role for each question automatically — not you typing act as an expert [V1]. That the same trade survives when you type the role is my judgement. [speculative]

I am not going to tell you the clarity loss outweighs the depth gain, because the two ratings sit on different scales and the study reports an effect size for only one of them. What I will say is that the one measured style effect includes a clarity cost, and for most of you clarity is the product. [speculative]

9 9. Habit 3 — let Claude ask, then restate everything in one message

On screen:

Part I · Habits

1

Ask for the questions

"Before you write anything, ask me the five questions you need answered." [speculative]

2

Answer in one message

Not five replies across five turns [inference]

Why restate it all at once: drip-feeding — "write a post", then "make it shorter", then "add our pricing", then "change the tone" — is the trap. Across 200,000+ simulated conversations, spreading the same information over turns instead of one message cost a 39% average performance drop, and one early wrong turn was never recovered [S2].

[S2] LLMs Get Lost In Multi-Turn Conversation — Laban et al. (2025), arXiv: 2505.06120. <https://arxiv.org/abs/2505.06120>

You do not need a perfect prompt from a standing start. Ask for the questions: before you write anything, ask me the five questions you need answered. [speculative]

But there is a second move, and dropping it makes the habit worse than not using it. Once you have the questions, answer them in a single message. Not one reply per question over five turns.

Laban and colleagues simulated over two hundred thousand conversations. Spread across turns, average performance dropped thirty-nine percent, and an early wrong turn was not recovered from later [S2].

That is a preprint with simulated conversations, so treat the number as a direction. [inference]

10 10. Habit 4 — one deliberate check between the draft and the send button

On screen:

Part I · Habits

Claude produces a true sentence and a hallucination by the same mechanism [S10] — nothing signals a wrong answer. Check the costliest fact. [inference]

- *Read the number off the source page, not from Claude. [inference]*
- *Ask the person who owns the fact, not Claude. [inference]*
- *Delete any sentence you cannot trace to a page or person. [inference]*

How much checking? In an MIT Sloan study, 140 professionals reviewed AI output that flagged which claims to verify. Flag everything and they caught the most errors but were slowest; flag the riskiest few and they caught nearly as much, far faster [S11]. So: one deliberate check on the costliest fact — not zero, not an audit. [inference]

[S10] AI hallucinations — IBM. <https://www.ibm.com/think/topics/ai-hallucinations>

[S11] Add speed bumps — MIT Sloan. <https://mitsloan.mit.edu/ideas-made-to-matter/to-help-improve-accuracy-generative-ai-add-speed-bumps>

This is the habit I would keep if I had to cut the other three. [speculative]

The mechanism, in one sentence. IBM: a language model has no way to distinguish a hallucination from the correct answer, because in both cases, in their words, it "makes up" the answer by guessing the next word [S10]. A confident tone tells you nothing.

So the check comes from outside, and it has to be small enough that you do it. Take the fact that would cost most if wrong — the price, the certification, the deadline — and confirm it from the source. Not from Claude in a fresh chat, which is the same mechanism twice. [inference]

MIT Sloan ran a field experiment with a hundred and forty Accenture research professionals. The friction was highlighting that colour-codes the output, so you have to process it rather than accept it. The heaviest highlighting caught more errors and omissions but took longer to finish, and the medium condition is the one they report as balancing detection against time [S11]. So one check, on one fact — not a review board. [inference]

11 11. The four habits — and a study where higher AI literacy came with more overconfidence

On screen:

Part I · Habits

1

Brief it

Reader, purpose, what to avoid [inference]

2

Role for wording

Not for accuracy [S1]

3

Ask, then consolidate

One message, not five turns [S2]

4

One check step

The costliest fact [S10]

Across 500 participants, "higher AI literacy brings more overconfidence" [S13] — never measured on a room leaving a training session, so carrying it here is my step. [speculative] Habit 4 alone puts a person between draft and reader. [inference]

[S13] AI and overconfidence — Aalto University (2026). <https://www.aalto.fi/en/news/ai-use-makes-us-overestimate-our-cognitive-performance>

Four habits: brief it, role for register only, ask then consolidate, one check step. [inference]

Now the uncomfortable part, which I would rather say myself. Fernandes and colleagues ran five hundred participants. Normally the least skilled overestimate themselves most; with AI, the researchers report that pattern vanishes, and higher AI literacy came with more overconfidence [S13].

So I am handing you a risk I can only half evidence. The overconfidence half is measured [S13]; that your output rises this month is my expectation, not a finding. That is not a reason to skip the training — it is why habit four is a takeaway rather than a footnote. [speculative]

12 12. A project holds the background paragraph, so you stop retyping it

On screen:

Part II · Projects

Anthropic's definition: "self-contained workspaces with their own chat histories and knowledge bases" [S4].

Without one

You paste who we are and what we never promise into every new chat, slightly differently each time.

With one

Write it once as instructions, upload the documents once, and every chat inside starts already briefed.

This is the facts half of the wording-or-facts question — not a better way to word things, but a place to put what Claude could not otherwise have. [inference]

[S4] What are Projects? — Anthropic Help Center. <https://support.claude.com/en/articles/9517075-what-are-projects>

Anthropic calls it a self-contained workspace with its own chat history and knowledge base [S4]. Accurate, and not why you would want one.

Here is why. That background paragraph you retype goes in once, as the project's instructions. The documents you keep re-uploading go in once, as project knowledge. Every chat started inside then begins already briefed. [inference]

Notice which half of the question that answers. It is the facts half. Everything you supplied by hand at the top of every conversation now lives in one place, and stops drifting each time you retype it.

13 13. Shared project 1 — Prompt Engineering, for when you are short of wording

On screen:

Part II · Projects

Where

Open your Projects list; sharing needs Team or Enterprise [S4]

Use it when

You cannot get Claude to word what you want. Not when facts are wrong [speculative]

What comes back

Treat it as a draft prompt, not as finished work [speculative]

In a 7,000-person rollout, "people were working inside a Claude Project when what they actually needed was a Skill" [S22]. Read from inside, the wrong surface looks like a tool that cannot do the job [inference].

[S4] What are Projects? — Anthropic. <https://support.claude.com/en/articles/9517075-what-are-projects> [S22] Non-technical AI adoption — Deel. <https://www.deel.com/deel-works/non-technical-ai-adoption>

The first shared project is a prompt helper. Describe the job in ordinary language and it writes the structured prompt habit one asks for. What comes back is a draft prompt you paste into a fresh chat, not the finished piece of work. [speculative]

The important thing is knowing when to reach for it. Reach for it when you know what you want and cannot get Claude to produce it. Not when the answer comes back factually wrong, because better wording does not supply a missing fact. [speculative]

That distinction is worth labouring — a Skill, in that quote, is the third personalization place, and we get to it two slides from now. Deel rolled Claude out to seven thousand people and reported staff sitting inside one feature trying to make it do another's job [S22]. From where you sit it feels like Claude cannot do the thing. It can — you are in the wrong place. [speculative]

14 14. Personalization is three separate places, not one switch

On screen:

Part II · Projects

Anthropic's page names three, differing by how far each reaches: Instructions for Claude, Project instructions, Skills [S5].

1

Instructions for Claude

Account-wide: "applied to all of your conversations" [S5]

2

Project instructions

Only "apply to chats within that project" [S5]

3

Skills

Format and tone: "how Claude formats and delivers" [S5]

Put a preference at the widest place where it is always true, never wider. A tone rule for one campaign, set account-wide, follows you into unrelated chats. [inference]

[S5] Personalization features — Anthropic Help Center. <https://support.claude.com/en/articles/10185728-understanding-claude-s-personalization-features>

If you go looking for a switch labelled personalization, you will not find one. Let me save you the ten minutes. [speculative]

Anthropic's page has three sections, and they differ by how far each one reaches. Instructions for Claude you set once: click your initials in the lower left, then Settings — and anything you put there is applied to all of your conversations [S5]. Project instruc-

tions sit inside one project and only apply to chats within it [S5]. Skills are the third, and the page describes them by what they do rather than by where they live: they "add specific behaviors or capabilities to your conversations with Claude", and the page tells you to use them when you want to customize how Claude formats and delivers its responses [S5]. Formatting and tone, in other words — for most of you the first two are the ones that matter today. [speculative]

One naming trap, and it is the vendor's rather than yours: the same page's summary calls the first one profile instructions, while its heading and the settings screen both say Instructions for Claude [S5]. Same setting, two names — go by the label on the screen, and do not spend five minutes hunting for a section called profile. [inference]

The rule is the one you use for company policy versus team norm versus today's brief. Put a preference at the widest place where it is always true, never wider. Always write in plain English — that is you, so Instructions for Claude. Never mention pricing in this campaign — that is one project. Put the second one account-wide and it follows you into a chat about holiday planning. [inference]

15 15. Shared project 2 — Personalization: set one preference, prove it changed something

On screen:

Part II · Projects

1

Open Personalization

All three places: Instructions for Claude, Project instructions, Skills [S5]

2

Write one line

"Plain English, under 150 words, no bullets unless I ask." [speculative]

3

Test it now

Re-ask yesterday's question and read the difference [inference]

If the answer does not change, check scope [S5]: Instructions for Claude applies to all your conversations, Project instructions only to chats in that project, Skills to how a reply is formatted and delivered. Thirty seconds, instead of assuming it works. [inference]

[S5] Personalization features — Anthropic Help Center. <https://support.claude.com/en/articles/10185728-understanding-claude-s-personalization-features>

The second shared project walks you through all three places — Instructions for Claude, Project instructions, Skills — so you do not have to remember which is which. [speculative]

Do one thing today, not five. Write one line. Mine would be: plain English, under a hundred and fifty words, no bullet points unless I ask — my own line, not one any source recommends. [speculative] Yours will differ and it does not matter what it is.

Then test it within thirty seconds. Take a question you asked yesterday, ask it again, read the difference. If the answer looks identical, the likeliest cause is scope: the line went into Project instructions or Skills when the chat needed it in Instructions for Claude, or the other way round [S5]. Pick a question whose answer would visibly change, or you will not be able to tell either way. [inference]

16 16. Shared project 3 — LinkedIn posts, into the feed

Pangram measured as the most AI-saturated

On screen:

Part II · Projects

40%+

of longform LinkedIn posts flagged fully AI-generated, 1,002,627 analysed [S14]

–45%

engagement on likely-AI posts versus likely-original, 2,726 posts [S15]

Both numbers come from companies selling AI detection, so treat them as a direction, not a measurement to bet on [inference]. The direction is enough: take the structure and the first draft from the project, and keep the specific — the customer, the number, the thing that happened — in your own words.

[S14] AI in your feed — Pangram Labs. <https://www.pangram.com/blog/ai-in-your-feed>

[S15] AI content on LinkedIn — Originality.ai. <https://originality.ai/blog/ai-content-published-linkedin>

The third project drafts LinkedIn posts, and this is where I slow you down. [speculative]

Two numbers. Pangram Labs analysed just over a million posts and found LinkedIn the most AI-saturated platform they measured, more than forty percent of longform posts flagged as fully AI-generated [S14]. Originality dot AI, on two thousand seven hundred posts: the likely-AI ones got forty-five percent less engagement [S15].

Both companies sell AI detection, and the labels come from their own classifiers, so treat these as a direction. [inference]

The direction still changes how you use it. Take the structure and the first draft. Keep the specifics yourself — the customer, the number, the mistake you made.

17 17. What not to put in a shared project

On screen:

Part II · Projects

A shared project's instructions and documents are "visible to everyone with access" [S19].

Fine to upload

Published material, product pages, the tone-of-voice guide, an approved case study.

Not in a shared project

A customer contract, a quotation, a named complaint.

Two more hazards: two people editing one project "will overwrite each other" [S19], and stale files never announce themselves [S20]. If you are not certain your organisation's data policy covers customer material in Claude, ask before the exercise. [speculative]

[S19] Claude Projects tutorial — EnterpriseDNA. <https://enterprisedna.co/resources/guides/guide-claude-projects-tutorial> [S20] Claude Projects for teams — OpenClawReady. <https://openclawready.com/blog/claude-projects-for-teams>

This slide before anybody opens a keyboard, because the first thing that happens after a session like this is somebody dropping a customer contract into a shared project. [speculative]

Here is the distinction nobody asks out loud. Your conversations stay yours. The files and instructions do not — they are visible to everyone the project is shared with [S19]. The chat is private and the knowledge is not, which is the wrong way round from how it feels.

My test: would I send this file to everyone on that distribution list? If not, it does not go in. [inference]

Two more. Two people editing the same project overwrite each other — there is no merge [S19]. And stale files never announce themselves [S20].

18 18. Build your own project in four steps

On screen:

Part III · Your own

1

Create it

"Replies to inbound enquiries", not "sales" [speculative]

2

Paste the paragraph

The background you retype, as instructions [inference]

3

Upload two or three files

Not twenty [speculative]. 30MB per project file [S6]

4

Share it, if your plan has sharing

"Team and Enterprise plans only" [S4]

Check your own plan before promising a colleague access [S4].

[S4] *What are Projects?* — Anthropic Help Center. <https://support.claude.com/en/articles/9517075-what-are-projects> [S6] *Upload files to Claude* — Anthropic Help Center. <https://support.claude.com/en/articles/8241126-upload-files-to-claude>

Four steps, three minutes. [speculative]

Create the project and name it after a job you do repeatedly rather than a topic: replies to inbound enquiries, not sales. Nothing in the documentation says to do that — it is my practice, because one set of instructions and files can only be right for one repeating job. [speculative]

Paste that background paragraph in as the instructions. You already wrote it several times this week. [inference]

Upload two or three documents, not twenty — next slide is why. The project-file limit is thirty megabytes, different from the chat upload limit on the same page [S6]. If a file bounces, check that first. [speculative]

Then share it. One caution: Anthropic says sharing is Team and Enterprise plans only, with two permission levels [S4]. Check what your account has before you promise anyone access.

19 19. Instructions describe behaviour; documents carry the content

On screen:

Part III · Your own

"Write instructions that describe behavior, not content" — a Q3 strategy summary is content, so it is a document [S19].

Why that is not cosmetic: as your project knowledge approaches how much text Claude can read at one go, the project stops loading everything and starts searching instead [S7]. Nothing on screen announces the switch. [speculative]

Anthropic's guidance is the smallest possible set of high-signal information [S8]. So more documents is not more accurate answers: replace a superseded document rather than leaving it beside the new one. [inference]

[S7] RAG for projects — Anthropic Help Center. <https://support.claude.com/en/articles/11473015-retrieval-augmented-generation-rag-for-projects> [S19] Claude Projects tutorial — EnterpriseDNA. <https://enterprisedna.co/resources/guides/guide-claude-projects-tutorial>

The rule: instructions describe behaviour, documents carry content [S19]. Never quote a delivery date — behaviour, so instructions. Here is our Q3 positioning — content, so a document. People invert this constantly, because pasting content into an instructions box feels like configuring.

Now the reason. There is a limit to how much text Claude can read at one go. As your uploaded files get close to it, Anthropic changes what the project does: instead of reading everything you uploaded, it searches your files and reads back the parts it finds. Their name for that is RAG mode — retrieval-augmented generation — and their docs say it expands capacity by up to ten times and needs no setup from you [S7]. That is the problem: I have seen nothing on screen that says it happened. [speculative]

So resist uploading everything. Anthropic's own guidance is the smallest set of high-signal material that gets the outcome [S8]. More documents is therefore not more accurate answers: when a document is superseded, replace it rather than leaving both in. [inference]

20 20. Starting from a slide template — and the file type that will stop you

On screen:

Part III · Your own

Accepted

"PDF, DOCX, CSV, TXT, HTML, ODT, RTF, EPUB, JSON, XLSX" [S6]

Not on that list

PPTX. The template file itself will not upload [inference]

Instead

Export the template to PDF, or type the slide structure into the instructions [inference]

What comes back is the words for each slide against your standing structure. You still assemble the branded file in PowerPoint, which is fine: the writing is the part that took the time. [speculative]

[S6] Upload files to Claude — Anthropic Help Center. <https://support.claude.com/en/articles/8241126-upload-files-to-claude>

This is the one I would have got wrong live, so let me save you two minutes. [speculative]

The idea is good. You have a template you use for every customer presentation. Turn it into a project and slide prep stops starting from nothing.

The obvious first move is to upload the PowerPoint file. Anthropic's supported list is PDF, DOCX, CSV, text, HTML, ODT, RTF, EPUB, JSON and XLSX [S6]. PPTX is not on it, so that upload fails, and it fails in front of you. [inference]

Two ways round. Export the template to PDF. Or type the structure into the instructions — slide one the problem, slide two our approach, slide three the evidence. The second is often better: a fixed sequence is behaviour. [inference]

21 21. Three million custom GPTs were built; about one in twenty is still listed publicly

On screen:

Part III · Your own

3,000,000+

custom GPTs built on OpenAI's earlier version of this feature [S16]

~159,000

still public and active in the GPT Store two years later [S17]

Creating one is free; keeping it correct is not, and nothing tells you when a project has gone stale [inference]. One practitioner write-up names the fix, and it is small — "an owner and a review date" [S20], written into the project's own instructions [inference].

[S16] Introducing the GPT Store — OpenAI (2024). <https://openai.com/index/introducing-the-gpt-store> [S17] GPT Store business guide 2026 — DigitalApplied. <https://www.digitalapplied.com/blog/gpt-store-custom-gpts-business-guide-2026>

We are not the first people to do this, and the previous round is worth thirty seconds. [speculative]

OpenAI shipped custom GPTs in late 2023 — the same idea, one generation earlier. By January 2024 they reported over three million built [S16]. Two years on, an independent count puts roughly a hundred and fifty-nine thousand as public and active [S17]. That second figure comes from a site with no published counting method, so take it as an order of magnitude — which is the point. [inference]

Creating one is free. Keeping it correct is not, and nothing tells you when a project has gone out of date — it just quietly answers with last year's price. One practitioner write-up names the fix, and it is small: a name and a review date in the project's own instructions [S20].

22 22. Hands-on — pick one and do it now

On screen:

Part IV · Now you

1

Rewrite one real prompt

Reader, purpose, what to avoid [inference]

2

Set one preference, test it

Re-ask yesterday's question [inference]

3

Create one project

Paste the paragraph you retype [inference]

Use your own work, not mine: employees "can't transfer a generic demo onto their actual Tuesday" [S18]. For novices, van Gog and colleagues predicted worked examples produce lower cognitive load and better learning than problem-solving alone [S21].

[S18] Enterprise Claude adoption — Correlation One. <https://www.correlation-one.com/blog/enterprise-claude-adoption-7-real-problems-and-the-solutions-that-worked> [S21]

Worked examples — van Gog et al. <https://www.sciencedirect.com/science/article/abs/pii/S0361476X1000055X>

Eight minutes. One of these three, not all of them. [speculative]

Rewrite a real prompt you sent this week, adding the reader, the purpose and what to avoid. Or set one preference and test it against yesterday's question. Or create one project and paste that paragraph into it.

One condition, and it is why this block exists: use your own work. Correlation One found people could not transfer a generic demo onto their actual Tuesday [S18]. If your own Tuesday is a customer email, strip the customer out of it first — the previous slide is why. [speculative]

And the order matters, which is why I demonstrated before asking. For novices, the research predicts that studying a worked example first leads to better outcomes than attempting first [S21]. You have seen four. Go.

23 23. Takeaways — what survives to next Tuesday

On screen:

Part IV

1

Wording or facts

Before every prompt [speculative]

2

Role prompting

No accuracy gain [S1]; wording only [speculative]

3

Ask, then one message

Consolidate; do not dribble [S2]

4

One check step, every draft

Nothing signals a wrong answer [S10]

5

An owner and a review date

In the project instructions [inference]

One practitioner: unowned projects answer from stale files [S20].

[S10] AI hallucinations — IBM. <https://www.ibm.com/think/topics/ai-hallucinations>

[S20] Claude Projects for teams — OpenClawReady. <https://openclawready.com/blog/claude-projects-for-teams>

Five things, and the fifth decides whether any of this still works in six months. [speculative]

Wording or facts, before every prompt. No significant accuracy gain from role prompting, across six models [S1] — so role prompting for wording, never for accuracy. [speculative] Let Claude ask, then answer in one message [S2]. One check step on every draft, on the fact that would cost most [S10].

And if you build a project anyone else will use, put your name and a review date in it [S20]. One quarter is my default, because a price list or a certification can go stale inside one. [speculative] Three million custom assistants were built on the last generation of this idea, and about a hundred and fifty-nine thousand are still listed publicly. Nothing measured why. The one thing a practitioner write-up points at is whether somebody owned it [S20].

If you want to read further, the field calls this prompt engineering — the systematic survey of it reviewed 1,565 papers and catalogues 58 text-based techniques [S23]. That is the phrase to search.

24 24. References

On screen:

- **[S1]** *Playing Pretend: Expert Personas, Prompting Science Report 4* — Wharton Generative AI Labs (2025). <https://gail.wharton.upenn.edu/research-and-insights/playing-pretend-expert-personas/>
- **[S2]** *LLMs Get Lost In Multi-Turn Conversation* — Laban et al. (2025), arXiv: 2505.06120. <https://arxiv.org/abs/2505.06120>
- **[S3]** *Be clear and direct* — Anthropic. <https://docs.claude.com/en/docs/build-with-claude/prompt-engineering/be-clear-and-direct>
- **[S4]** *What are Projects?* — Anthropic Help Center. <https://support.claude.com/en/articles/9517075-what-are-projects>

- **[S5]** Understanding Claude's personalization features — Anthropic Help Center. <https://support.claude.com/en/articles/10185728-understanding-claude-s-personalization-features>
- **[S6]** Upload files to Claude — Anthropic Help Center. <https://support.claude.com/en/articles/8241126-upload-files-to-claude>
- **[S7]** Retrieval augmented generation (RAG) for projects — Anthropic Help Center. <https://support.claude.com/en/articles/11473015-retrieval-augmented-generation-rag-for-projects>
- **[S8]** Effective context engineering for AI agents — Anthropic. <https://www.anthropic.com/engineering/effective-context-engineering-for-ai-agents>
- **[S9]** Claude Cowork — Anthropic. <https://www.anthropic.com/product/claude-cowork>
- **[S10]** What are AI hallucinations? — IBM. <https://www.ibm.com/think/topics/ai-hallucinations>
- **[S11]** To help improve the accuracy of generative AI, add speed bumps — MIT Sloan. <https://mitsloan.mit.edu/ideas-made-to-matter/to-help-improve-accuracy-generative-ai-add-speed-bumps>
- **[S12]** Chatbot use by non-IT laypersons: a PRISMA review of 28 studies — Universal Access in the Information Society (2026). <https://link.springer.com/article/10.1007/s10209-026-01323-w>
- **[S13]** AI use makes us overestimate our cognitive performance — Aalto University on Fernandes et al. (2026). <https://www.aalto.fi/en/news/ai-use-makes-us-overestimate-our-cognitive-performance>
- **[S14]** AI in your feed — Pangram Labs. <https://www.pangram.com/blog/ai-in-your-feed>
- **[S15]** AI content published on LinkedIn — Originality.ai. <https://originality.ai/blog/ai-content-published-linkedin>
- **[S16]** Introducing the GPT Store — OpenAI (2024). <https://openai.com/index/introducing-the-gpt-store>
- **[S17]** GPT Store and custom GPTs: business guide 2026 — DigitalApplied. <https://www.digitalapplied.com/blog/gpt-store-custom-gpts-business-guide-2026>
- **[S18]** Enterprise Claude adoption: 7 real problems and the solutions that worked — Correlation One. <https://www.correlation-one.com/blog/enterprise-claude-adoption-7-real-problems-and-the-solutions-that-worked>
- **[S19]** Claude Projects tutorial — EnterpriseDNA. <https://enterprisedna.co/resources/guides/guide-claude-projects-tutorial>

- **[S20]** *Claude Projects for teams — OpenClawReady.* <https://openclawready.com/blog/claude-projects-for-teams>
- **[S21]** *Effects of worked examples, example-problem, and problem-example pairs on novices' learning — van Gog, Kester & Paas.* <https://www.sciencedirect.com/science/article/abs/pii/S0361476X1000055X>
- **[S22]** *Non-technical AI adoption — Deel.* <https://www.deel.com/deel-works/non-technical-ai-adoption>
- **[S23]** *The Prompt Report: A Systematic Survey of Prompting Techniques — Schulhoff et al. (2024), arXiv:2406.06608.* <https://arxiv.org/abs/2406.06608>
- **[V1]** *Role Prompting — author's working notes. Summarised in the printed appendix.*
- **[V2]** *State of AI in Business 2025 (MIT NANDA) — Overview — author's working notes. Summarised in the printed appendix.*

The handout PDF carries the full reference list, and a one-page summary of every [V] source with its own external references.

Every source is on this slide. The credibility grade I gave each one, and why, is in the printed version rather than here. [inference]

Seven of the twenty-three are vendor or consultancy write-ups rather than research — the two LinkedIn measurements, the GPT Store count, and the four practitioner accounts of running Claude projects. I graded three of those Low. The two marked V are my own working notes; the printed version summarises each and lists every external source it draws on, because you cannot open my notes. [inference]