

Speaker Notes

CONTENTS

- [1 1. Building AI Systems](#)
- [2 2. Reliability is engineered, not prompted](#)
- [3 3. "Done. I fixed it and verified everything passes."](#)
- [4 4. An agent auditing its own work caught 0 of 34 real violations](#)
- [5 5. Roadmap](#)
- [6 6. What this hour is — and what it leaves out](#)
- [7 7. Takeaways](#)
- [8 8. How this deck was made](#)
- [9 9. The thing you already run](#)
- [10 10. You already operate a harness](#)
- [11 11. The harness moves the number more than the model does](#)
- [12 12. Workflow or agent — and why almost everything is a workflow](#)
- [13 13. Generator, verifier, oracle](#)
- [14 14. The unit the bill is denominated in](#)
- [15 15. Reliability is not accuracy](#)
- [16 16. Why a generator](#)
- [17 17. Two loops that look identical from outside](#)
- [18 18. The negative results, named](#)
- [19 19. Why a model cannot check more cheaply than it generates](#)
- [20 20. The consistency dilemma](#)
- [21 21. Self-monitoring fails worst under attack](#)
- [22 22. Verification always costs something](#)
- [23 23. Self-critique degrades above a threshold — it is not universally broken](#)
- [24 24. What survives all four counterarguments](#)
- [25 25. What to do with part two](#)
- [26 26. What an external check buys](#)
- [27 27. The number that made me rewrite my own architecture advice](#)
- [28 28. One architecture in three unrelated fields: the generator proposes, an oracle outside it returns pass or fail](#)
- [29 29. Independence costs a fresh context, not a second vendor](#)
- [30 30. Which rung verifies this decision?](#)
- [31 31. Rung four is not a rung you can rest on](#)
- [32 32. Grounding is not the same thing as retrieval](#)
- [33 33. Faithful and wrong at the same time](#)

34 34. Boundary one — the oracle only sees what it measures

35 35. Boundary two — somebody has to write the oracle

36 36. Boundary three — the oracle depreciates

37 37. What to do when you have no oracle

38 38. What to do with part three

39 39. Bounded autonomy

40 40. Why long chains fail even with a good model

41 41. Real pipeline errors are not independent — they stick

42 42. Nine seconds: an agent deleted a production database

43 43. Ask why the agent was able to delete the database

44 44. Deny at the API gate, not in the prompt

45 45. Containment over supervision — and it still leaks

46 46. Guardrails are layers, not boundaries

47 47. Five autonomy levels, and each names who answers

48 48. The human rung is a control that degrades

49 49. The EU AI Act requires oversight that can intervene

50 50. Bounded autonomy can relocate blame instead of risk

51 51. What to do with part four

52 52. Cost, capacity

53 53. The bill for closing a loop on an external check: more agents, more iterations, a memory layer

54 54. Price has become a design variable

55 55. Generation got cheap; review and integration did not

56 56. Instrument, then evaluate, then gate

57 57. If you write code

58 58. If your job is to judge what these systems produce

59 59. If you decide budgets

60 60. What I do not know

61 61. Takeaways

62 62. What follows from an external check that can fail the work, and a deny gate at the credential

63 63. References

64 64. References (continued)

65 65. References (final)

Generated by the agentic vault — an autonomous research and authoring harness. Every factual claim was gated against a cited source registry before this document was produced.

Slides: 65 · **Duration:** 60 min · **Venue:** Brno (community event)

1 1. Building AI Systems

On screen:

Brno tech community · 2026-09-16

Reliability is engineered, not prompted.

Ondrej (Ondra) Krajicek me@ondrejkrjicek.com · [linkedin.com/in/OndrejKrajicek](https://www.linkedin.com/in/OndrejKrajicek) Built at render time

Good evening, and thank you for turning up. [speculative] If you write code, you should leave with something to build or refuse to build on Monday. [speculative] If you do not write code but you have to judge what these systems produce, you should leave with a question you can ask that gets a real answer. [speculative] And if you decide budgets, you should leave with the failure modes and the cost shape written down. [speculative] Every load-bearing claim tonight carries a number, a mechanism or a named source, because none of you should buy an adjective from a stranger. [speculative]

2 2. Reliability is engineered, not prompted

On screen:

The claim

Two mechanisms, and neither is a prompt. [inference]

- **External grounding** — *a check that can still fail the work when the generator is wrong. A model approving its own output cannot: it shares the blind spots that produced the answer. [S3]*
- **Bounded autonomy** — *deny irreversible actions at the credential and the API gate, not in the prompt. A prompt instruction is advisory; any input reaching the context can override it. [S32]*
- *Both are architecture, not wording, and outlast the next model. [speculative]*

[S3] Loop Engineering and Self-Verification — vault report (2026). [S32] Harness Engineering — Tool Permissions — vault report (2026).

Here is the whole argument on slide two, so you know exactly what you came to disagree with. [speculative] First mechanism: external grounding. A check is worth something only if it can come back negative when the generator is wrong, and a model asked to approve its own output cannot — it shares every blind spot that produced the answer. [S3] The report this talk is built on states it as a rule: buy that independence in the check, or do not call the loop verified. [S3] Second mechanism: bounded autonomy. Deny the irreversible actions at the credential and at the API gateway, never in the prompt, because a prompt instruction is advisory and any input that reaches the context can override it. [S32] Neither of those is wording, so neither should expire with the next model release. [speculative] If by the end you think these are the wrong two variables, come and tell me which ones are.

3 3. "Done. I fixed it and verified everything passes."

On screen:

Opening

Every seat in this room has read a sentence like that and found out otherwise. [speculative]

1

If you write code

The assistant reports the fix, the suite is green, the defect ships. [speculative]

2

If you judge output

A confident answer, a real-looking citation, a policy that was superseded. [speculative]

3

If you fund it

The demo worked; the pilot did not; nobody can say which step failed. [speculative]

Start with something none of you needs me to explain. [speculative] You have all seen a system announce that it finished, in a completion message so specific and so calm that disagreeing looks unreasonable — and then it had not finished. [speculative] Each column is the version of that you have lived. [speculative] The rest of this hour is about why that shape happens and what specifically removes it. [speculative]

4 4. An agent auditing its own work caught 0 of 34 real violations

On screen:

Opening

A false completion, counted: the agent reviewed its own output, stated high confidence, and was wrong on all 34 cases. [S1]

0 of 34

Real violations caught by the agent auditing its own work [S1]

34 of 34

The same violations caught by a deterministic checker [S1]

Stated confidence was 90–100 on every one of the 34 it missed. A number that reads the same whether the work is right or wrong cannot tell you which. [inference]

[S1] Loop Engineering and Self-Verification Loops — Practitioner Synthesis — vault research report (2026).

Here is the false-completion report with a number attached. [speculative] An agent auditing its own work caught zero of thirty-four real violations, and it stated ninety to one hundred confidence while missing every one of them. [S1] A deterministic checker, run over the same artifacts, caught all thirty-four. [S1] Not a better model — a program. [inference] Now look at what that confidence number did across the thirty-four cases: nothing. [inference] It did not move when the work was wrong, so it cannot tell you the work is wrong. [inference] I want to be careful about how far that generalises, because it is one measurement and not a theorem — where a model has an external signal to condition on, a stated confidence can carry real information. [speculative] Here it had none, and it read like certainty anyway. [inference]

5 5. Roadmap

On screen:

Roadmap

- 1. The thing you already run — harness, agent, workflow, and the vocabulary for the hour*
- 2. Why a generator cannot be its own oracle — and four reasons that is not the whole story*

3. What an external check buys, what it costs, and where it goes blind
4. Bounded autonomy — deny at the API gate, not in the prompt
5. Cost, and one decision per seat in this room

By the end: two five-level classification scales for a system you already run. One ranks a check by what it depends on, from a deterministic test to the model's own weights; one ranks an agent by what it may do unapproved. [inference]

Five parts, about ten minutes each. [speculative] Part one names the machinery behind the tools, because plenty of people run one every day without using the word. [speculative] Part two is the negative result at the centre of the talk, and then four serious reasons it is narrower than the slogan version. [speculative] Part three is the positive half, with the bill and the blind spots attached. [speculative] Part four I will argue is a permissions problem before it is an AI problem. [speculative] The two scales on the slide are the take-home; the theory exists to earn them. [speculative]

6 6. What this hour is — and what it leaves out

On screen:

Scope

The hour covers where a check sits, who may act, and what it bills. It skips the model layer, because that layer changes fastest and dates soonest. [speculative]

Covered

- *Where the verifying oracle sits, and what it can see [speculative]*
- *Permissions, action gates, autonomy levels [speculative]*
- *Token cost, review capacity, adoption drag [speculative]*

Not covered

- *Prompt technique and model internals [speculative]*
- *Which model is best this month [speculative]*
- *CI, code review, testing, containers — assumed, not explained [speculative]*
- *Retrieval implementation and knowledge-graph construction [speculative]*

A fence around the hour, so we do not drift in the questions. [speculative] Deliberately absent: prompt technique and model internals, because that is the smallest and least durable layer. [speculative] Not which model is best this month — that answer expires before you get home, and it is the one question I am most often asked. [speculative] Not the delivery basics: if you write code you have run CI and code review for

years, and if you do not, nothing tonight depends on them. [speculative] And not retrieval implementation, which deserves its own hour. [speculative] If you came for a tool comparison, part five will tell you why the comparison you want is between harnesses rather than between models. [speculative]

7 7. Takeaways

On screen:

Where this goes

1

Move truth out

Reliability came from closing the generation–verification loop on an external verifier, not from a more capable model. [S5]

2

Buy independence

A check counts only if it can fail when the generator is wrong. Iterations do not buy that. [S3]

3

Name the rung

Ask which of five levels verifies a decision before trusting it. [S18]

4

Authority last

Irreversible actions — production deploys, money movement, force-push — belong on a deterministic deny gate, never on a probabilistic reviewer. [S20]

One: move the truth-determining computation out of the model — the study behind my headline number says reliability came not from a more capable model but from closing the loop against an external verifier. [S5] Two: a check counts only if it can come back negative when the generator is wrong, and running the same model over its own answer more times does not buy that. [S3] Three: before you trust a decision, name which rung of a five-level ladder verified it — seventy percent of measured production loops verify at the bottom two. [S18] I want that to land as a good sign rather than as a limitation. [speculative] Four: add authority last — irreversible actions, meaning production deploys, money movement and force-push, belong on a deterministic deny gate, never on a probabilistic reviewer. [S20] If you disagree early, spend the hour building the rebuttal.

8 8. How this deck was made

On screen:

Disclosure

Eight model personas swept a model-curated knowledge base; a model merged them; model critics checked the merge. [speculative]

- *By the argument I am about to make, that configuration is exactly the one you should not trust. [inference]*
- *"models agreeing with models about a memory curated by models" [S48]*
- *The external checks here are mechanical: every quoted span is verified byte-for-byte against the file it names, and every claim carries a citation or a confidence marker. [speculative]*
- *Judge the deck the way I am asking you to judge your systems — by what checked it, not by how fluent it reads. [inference]*

[S48] 2026-07-17 Digest — vault journal note.

This deck was researched by eight model personas over a knowledge base that models helped curate, merged by a model, and checked by more models. [speculative] My own notes describe that configuration in a phrase I cannot improve on — models agreeing with models about a memory curated by models. [S48] By my own thesis, that is the least trustworthy possible arrangement, and I would rather say it now than have one of you say it in the questions. [inference] What makes it checkable: the quoted spans in my sources file are verified character by character against the notes they come from, by a script, and any sentence with no source carries a marker saying so. [speculative] That is a deterministic external check on a generated artifact. [inference] The same move I will recommend all evening, applied to the evening itself. [inference]

9 9. The thing you already run

On screen:

Part I

Naming the machinery before arguing about it.

Part one, and it is the shortest. [speculative] I am not going to define what a prompt is or explain what a large language model does; if you write code, most of you use these tools daily and that slide would insult you, and if you do not, nothing tonight turns on the model's internals — it turns on what sits around the model. [speculative] What I do want is four shared words, because this room comes from different companies and different stacks, and I have watched conversations about agents fail purely on vocabulary — two people using the word verifier for two different things and agreeing loudly while disagreeing completely. [speculative] So: harness, workflow, agent, oracle. [speculative] Four words, about eight minutes, and then we can argue about substance instead of terms. [speculative]

10 10. You already operate a harness

On screen:

Part I · Harness

A harness is "everything in an AI agent except the model itself" — the tools, the context assembly, the permissions, the retries, the stop condition. [S21]

- *The coding assistants most engineers now open daily are harnesses; the model is one component inside them. [inference]*
- *The mental model that survives a mixed room: "the LLM is the CPU and its context window is the RAM" [S50] — and the harness is the operating system above both. [inference]*
- *So the question "should we build an agent?" is usually really "whose harness are we adopting?" [speculative]*

[S21] Harness Engineering · [S50] Context Engineering — vault concept notes.

Start with the word the other three definitions hang off. [speculative] A harness is everything in an AI agent except the model itself — tool definitions, context assembly, permission checks, retry policy, the stop condition. [S21] If you write code, the assistant most of you opened this morning is a harness with a model inside it, rather than a model with a chat window. [inference] If you do not, the systems you are asked to judge or fund have the same shape: a model inside a wrapper that decides what it sees and what it may do. [inference] The picture that works across a mixed room is Karpathy's: the model is the CPU, the context window is the RAM, and context engineering is memory management. [S50] The harness is the operating system sitting above both — it decides what gets loaded, what gets called, and what is allowed. [inference] Which reframes the buying question you are actually facing: not which model, but whose harness. [speculative]

11 11. The harness moves the number more than the model does

On screen:

Part I · Harness

Changing the harness moved accuracy 16 points; the top three models on one standardised board sat within 0.3. [S23] [S69]

16 pts

93.1% vs 76.7%, same model, LongMemEval [S23]

0.3 pts

Opus 4.8, Fable 5, GPT-5.5 within 0.3, Terminal-Bench 2.1 [S69]

- *Second benchmark, same pattern: GPT-5.5 posts "83.4% inside Codex CLI but 78.2% inside Terminus 2" [S69].*
- *"grep 93.1% vs. vector 83.6%" on the LongMemEval subset; ordering flips under file-based tool delivery [S24].*

[S23] [S24] Is Grep All You Need? — vault note. [S69] Harness Engineering — vault note.

Here is why the harness deserves the word. [speculative] In one 2026 study, the same model on the same task scored ninety-three point one percent inside one harness and seventy-six point seven inside another — a sixteen-point gap attributable to the wrapper. [S23] Beside it, the model spread: the top three models on one standardised Terminal-Bench 2.1 board sit inside three tenths of a point, while the same models move four to five points by harness. [S69] Across model generations I have no measurement and claim none. [inference] The second number is the one that will annoy some of you: plain keyword search beat vector retrieval, ninety-three point one against eighty-three point six, on that subset. [S24] It is one benchmark and a small number of harnesses, and in the same study the ordering flips under different tool-delivery styles. [S24] So do not go home and rip out your retrieval layer. [speculative] Go home and measure whether your harness or your model is the variable you have been tuning.

12 12. Workflow or agent — and why almost everything is a workflow

On screen:

Part I · Vocabulary

The line is who writes the control flow — almost everything called an agent is a workflow. [inference]

Workflow

Orchestrated "through predefined code paths" [S51] — you write the control flow, the model fills slots.

Agent

The model "dynamically direct[s] their own processes" [S51] — next step and stop condition come from its output, not your code.

- *Across fifty hand-coded loops: "the median loop is a manually-triggered single agent with a deterministic check" [S19].*
- *Workflows fail predictably, and that predictability is what you are buying. [speculative]*

[S51] Building Effective Agents — Anthropic. <https://www.anthropic.com/engineering/building-effective-agents> [S19] Loop Engineering — vault concept note (2026).

A workflow orchestrates a model through predefined code paths — you wrote the control flow, the model fills the slots. [S51] An agent dynamically directs its own process — the next step and the stopping point come out of the model output rather than out of your control flow. [S51] By that definition almost everything you call an agent is a workflow — good news, not an embarrassment. [inference] The measured picture supports it: across a hand-coded corpus of fifty real production loops, the median is a manually triggered single agent with a deterministic check — no swarm, no scheduler, no self-direction. [S19] The unattended multi-agent swarm is a conference talk rather than a median deployment. [speculative] And the reason to prefer the workflow is not caution; it is that a workflow fails in ways you can enumerate in advance, which is the property you are actually buying. [speculative]

13 13. Generator, verifier, oracle

On screen:*Part I · Vocabulary*

Three roles, and the hour turns on the third — the only one that can contradict the model. [inference]

1

Generator

The component that produces the work. Usually the model. [speculative]

2

Verifier

Anything that judges the work — a program, a model, a person. [speculative]

3

Oracle

"A verifier that can tell right from wrong without asking the model" [S45] — every oracle is a verifier; most verifiers are not oracles. [S45]

[S45] Building AI Systems — predecessor-talk glossary — vault note (2026).

A generator produces the work — for us, the model. [speculative] A verifier judges it, and deliberately says nothing about how — a program, another model, or a person with a stake. [speculative] An oracle is the narrow case: a verifier that can tell right from wrong without asking the model. [S45] A compiler is an oracle. [speculative] A test suite is an oracle for the behaviours it covers. [speculative] A second model reading the first model's output is a verifier and is emphatically not an oracle. [S45] When someone tells you their agent is verified, that one distinction is the entire question, and I would ask it before asking anything about the model.

14 14. The unit the bill is denominated in

On screen:

Part I · Cost

Token

Roughly four characters, or 0.75 words of English text [S61]

What you pay for

Every token in and out, on every step of every retry [inference]

Orchestration multiplier

"Multi-agent systems use approximately 15× more tokens than chats" [S38]

Fifteen times more tokens means a fifteen-times bill. [inference]

[S61] Concepts — tokens — OpenAI docs. <https://platform.openai.com/docs/concepts> · [S38]

Agentic Orchestration Patterns — vault concept note.

A token is about four characters, or three quarters of an English word. [S61] You pay per token in and per token out, on every step, including every retried step. [inference] Now the multiplier that matters: multi-agent orchestration uses roughly fifteen times the tokens of a straight chat. [S38] That is not a rounding error and it is not amortised away. [inference] I raise it here rather than in the cost section because it changes how you should hear the next forty minutes: every architectural improvement I recommend has a token price, and I will keep quoting it. [speculative] Anyone who recommends orchestration without a number attached is recommending a bill they have not read. [speculative]

15 15. Reliability is not accuracy

On screen:

Part I · Vocabulary

Capability and reliability have come apart, and the gap between them is where incidents live. [inference]

pass@k — the ceiling

"probability of success in at least one of k attempts" [S44] — what demos and benchmarks report. [inference]

pass^k — the floor

Probability that all k attempts succeed — what production feels. [S44]

- *The 2026 finding: "recent capability gains have only yielded small improvements in reliability", and "agents that can solve a task often fail to do so consistently" [S56].*

[S44] Unit Testing — vault note. [S56] Towards a Science of AI Agent Reliability — Rabanser, Kapoor et al. (ICML 2026). <https://arxiv.org/abs/2602.16666>

Pass-at-k is the probability of succeeding in at least one of k attempts — that is a ceiling, and it is what demos and benchmark tables report. [S44] Pass-to-the-k is the probability that all k attempts succeed — a floor, and that is what production feels every day. [S44] The two can be very far apart. [inference] A 2026 paper that decomposes agent reliability into twelve separate metrics reports that recent capability gains yielded only small improvements in reliability, and that agents able to solve a task often fail to do so consistently. [S56] Capability and reliability have come apart as axes. [inference] So when a vendor shows you a benchmark score, the follow-up question is not how high — it is how often, and out of how many runs.

16 16. Why a generator

On screen:

Part II

The negative result, and the four reasons it is narrower than the slogan.

Part two is the theoretical heart. [speculative] I will make the case as hard as I can for about seven minutes, and then spend three minutes attacking it with the four strongest counterarguments I could find, two of which come from my own notes and one from the researcher whose work the case rests on. [speculative] I do that because a slogan gives you nothing to act on: if you are going to change what you build, or sign off on what somebody else built, you need to know where the claim stops holding. [speculative] The version that survives the attack is narrower, more useful, and still enough to change what you build. [speculative]

17 17. Two loops that look identical from outside

On screen:

Part II · Archetypes

[code]

Same retry code, same logs, same "verified" label — and only one of them adds information. [S3]

[S3] Loop Engineering and Self-Verification Loops — Comprehensive Report — vault research report (2026).

Two architectures, and from the outside they are indistinguishable. [speculative] On the left the loop closes on the same model, asked to check its own output. [speculative] On the right the loop closes on an oracle — a compiler, a solver, a schema validator, an external registry — something that can return pass or fail without consulting the thing that produced the artifact. [speculative] The report I built this on states the discriminator in one line: a loop's verdict is only worth what its verifier's independence is worth. [S3] Iteration count does not appear in that sentence. [inference] Model size does not appear in that sentence. [inference] When someone shows you an architecture diagram with a box labelled verification, this is the only question worth asking about the arrow that returns.

18 18. The negative results, named

On screen:

Part II · Evidence

Huang et al.

[S53]

Intrinsic self-correction lowered GPT-4's GSM8K accuracy across two rounds

ICLR 2024

Stechly et al.

[S54]

Game of 24 5% → 3%, graph colouring 16% → 2% under self-critique; Blocksworld 40% alone, 55% self-verified, 88% with a sound external verifier

ICLR 2025

Vault synthesis

[S1]

Self-audit caught 0 of 34 violations; a deterministic checker caught 34 of 34

2026

Two research groups and a practitioner audit, one direction: convergence, not one thesis. [inference]

[S1] Practitioner Synthesis · [S53] <https://arxiv.org/abs/2310.01798> · [S54] <https://arxiv.org/abs/2402.08115>

Huang and colleagues, at ICLR 2024: intrinsic self-correction — the model reviewing its own answer with no external signal — lowered GPT-4's grade-school-maths accuracy across two rounds rather than raising it. [S53] Stechly, Valmeekam and Kambhampati, ICLR 2025, on puzzles with machine-checkable ground truth: under self-critique, Game of 24 fell from five percent to three, and graph colouring from sixteen percent to two. [S54] The same paper ran the three-way comparison on Blocksworld — the generator alone at forty percent, the generator checking its own plans at fifty-five, the generator paired with a sound external verifier at eighty-eight. [S54] Self-verification was worth fifteen points there; external verification was worth forty-eight, on the same generator. [inference] That is the cleanest isolation of the verifier as the variable I have. [speculative] And the vault synthesis number you already have: zero of

thirty-four against thirty-four of thirty-four. [S1] Two independent research groups and one practitioner audit, three methods, one direction. [inference] I stress intrinsic in the first result: the scope of a negative result matters, and I come back to it in four slides.

19 19. Why a model cannot check more cheaply than it generates

On screen:

Part II · Mechanism

The intuition that verification is easy comes from complexity theory, and it does not transfer. [S12]

- *The belief "rests on the assumption that verification should be easier than generation" [S12] — true when checking is a different computation from producing.*
- *Asked to verify, the model runs the same approximate retrieval it ran to generate, so the asymmetry never appears. [S12]*
- *And the capability does not arrive with scale: "improving a model's generation performance does not lead to corresponding improvements in its ability to verify its own solutions" [S10].*

[S10] [S12] Loop Engineering and Self-Verification — vault report.

The intuition is imported from complexity theory: checking a solution is cheaper than finding one — verifying a satisfying assignment is trivial, finding it is not. [speculative] The report states the assumption directly: the belief that self-critique helps rests on the assumption that verification should be easier than generation. [S12] That holds when checking is a genuinely different computation from producing. [inference] It is not here. [speculative] A model asked to verify runs the same approximate retrieval over the same distribution it just ran to generate, so the asymmetry that makes checking cheap never materialises. [S12] And the obvious hope does not survive either: improving generation performance does not improve the ability to verify one's own solutions, even on the same task. [S10] Which makes this a design axis rather than a model-procurement one.

20 20. The consistency dilemma

On screen:

Part II · Mechanism

Agreement between a generator and its own evaluator tracks higher vulnerability to mistakes.

[S11]

- *Measured: "higher generator-evaluator self-consistency tracks higher mistake vulnerability — the consistency dilemma" [S11].*
- *And the scaling hope is empirically backwards: "self-verification capability does not improve with generation capability or scale" [S16].*
- *Practical read: if your reviewer model agrees with your author model almost always, you have bought confidence rather than coverage. [inference]*

[S11] [S16] Loop Engineering and Self-Verification Loops — vault research report (2026).

The measured relationship is that higher self-consistency between a generator and its own evaluator tracks higher vulnerability to validated mistakes — the consistency dilemma. [S11] Read that operationally. [speculative] If you have wired a reviewer model in front of an author model and the reviewer approves ninety-eight percent of the time, the natural reaction is that quality is high. [speculative] The measurement says the opposite reading is at least as likely: the two agree because they share the failure structure, and the approval rate is telling you about correlation rather than about correctness. [inference] And the hope that the next model fixes it is backwards — self-verification capability does not improve with generation capability or scale. [S16] So a high agreement rate between two components of the same system is a diagnostic to investigate rather than a metric to report.

21 21. Self-monitoring fails worst under attack

On screen:

Part II · Evidence

A monitor grading output from its own policy was 5× more likely to approve a prompt-injected patch, and code-correctness AUROC fell from 0.99 to 0.89 under the same condition. [S9]

5×

Approving a prompt-injected patch [S9]

0.99 → 0.89

Code-correctness AUROC [S9]

- *"On-policy" is the default: ask a system to review itself and the monitor grades its own policy's output. [inference]*
- *The failure is adversarial, not random — a monitor that shares the generator's blind spots shares its exploitable ones. [inference]*

[S9] *Loop Engineering and Self-Verification Loops — vault research report (2026).*

Under on-policy self-monitoring — a monitor grading output produced by its own policy, which is what you get by default when you ask a system to review itself — the model became five times more likely to approve a patch that had followed a prompt injection. [S9] Over the same condition, the ability to discriminate correct from incorrect code, measured as area under the ROC curve, collapsed from ninety-nine to eighty-nine. [S9] This is not a monitor being a bit noisy. [inference] It is a monitor that shares the generator's blind spots, which means it also shares the ones an attacker can aim at. [inference] If your defence against prompt injection is a model checking a model, you have one component, not two.

22 22. Verification always costs something

On screen:

Part II · Limits

No check is free. In a real system you pay in one of these. [S14]

- **A second inference pass** — compute and latency on every request. [inference]
- **A different model** — money, and new failure modes. [inference]
- **A deterministic check** — engineering time, and it only covers what you encoded. [inference]
- **A human review step** — scarcest, and first to stop scaling. [inference]

Nothing can check itself is the stronger claim, and I am not making it: my source calls the 2026 bound "a diagnostic framework, not an impossibility theorem" [S13].

[S13] [S14] *Loop Engineering and Self-Verification — vault report (2026).*

Trustworthy knowledge always pays for its own verification, and the price is never zero. [S14] In a real system you can point at the currency, and there are four of them. [inference] A second inference pass costs compute and latency, on every request. [inference] A different model costs money, plus a second set of failure modes somebody has to learn. [inference] A deterministic check costs engineering time up front, and it only ever covers what you thought to encode. [inference] A human review step is the scarcest of the four and the first to stop scaling. [inference] If you cannot point at one of those, you have not bought verification — you have assumed it. [inference] Now the limit, because you may have heard a stronger claim than mine. There is a 2026 information-theoretic bound on rounds of self-critique, and my own report labels it a diagnostic framework rather than an impossibility theorem: the latent failure variable it turns on still needs independent operationalisation before the bound has empirical bite. [S13] So no formal result forbids self-verification, and I am not going to pretend one does. [inference]

23 23. Self-critique degrades above a threshold — it is not universally broken

On screen:

Part II · Counterarguments

Four results qualify the strong claim. [inference]

1. **A threshold, not a law** — degrades only above a model-specific accuracy threshold [S60]; below it, a weak generator gains [S63].
2. **Your daily tools already verify** — Anthropic on Claude 5: "remove explicit 'add a final verification step' instructions" [S46].
3. **The verifier can move into training** — training-time gains, no inference-time signal: "the verifier is 'relocated,' not required" [S47].
4. **Its own authors reject both extremes** — LLM-Modulo rejects over-optimism and over-pessimism alike. [S55]

[S60] <https://arxiv.org/html/2604.22273v2> · [S63] <https://arxiv.org/html/2601.00828v1> · [S55] <https://arxiv.org/html/2402.01817v2> [S46] Claude 5 guidance · [S47] 2026-07-18 Digest

One: a threshold, not a law. A 2026 preprint treating self-correction as closed-loop feedback control finds it degrades performance above a model-specific accuracy threshold — four of its seven models degrade, including GPT-5, which loses one point eight points from a ninety-six point two percent baseline. [S60] That paper evaluates nothing below its own threshold. [S60] A second preprint covers the other side: on GSM8K-Complex, GPT-3.5 at sixty-six percent corrected one point six times more of its own errors than DeepSeek at ninety-four. [S63] Neither places GPT-3.5 against the other's threshold, so I am not claiming one curve — I am claiming the sign of the effect depends on where the generator starts. [inference] Two: the tools in your pocket already verify. Anthropic's own prompting guidance for Claude 5 tells developers to remove the explicit add-a-verification-step instruction, because the extra prompt causes over-verification. [S46] Scope that before you generalise it: one vendor, its own model family, no cross-lab replication I could find. [inference] Three: the verifier can move into training — those methods report self-correction gains with no inference-time signal at all. [S47] Four, and this one stings: the LLM-Modulo authors, whose work my case rests on, reject the over-optimistic and the over-pessimistic extreme alike. [S55] I am obliged to sit between them.

24 24. What survives all four counterarguments

On screen:

Part II · Synthesis

A capability threshold, tools that already verify, a verifier moved into training, the LLM-Modulo authors rejecting both extremes: after all four, two claims die, two stand. [inference]

- **Dies:** "Self-critique never works." On GSM8K-Complex, GPT-3.5 (66% baseline accuracy) corrected 1.6× more of its own errors than DeepSeek (94%) [S63].
- **Dies:** "Only an inference-time external verifier helps." Training-time verifiers give the same signal, paid once [S47].
- **Survives:** A check counts only if it can fail when the generator is wrong: fresh context, another model, or none [S3].
- **Survives:** Verification has a price: a second pass, another model, a deterministic check, someone's time. Name yours [S14].

Two claims die. [speculative] Self-critique never works is false — on GSM8K-Complex, GPT-3.5 at sixty-six percent corrected one point six times more of its own errors than DeepSeek at ninety-four, twenty-six point eight against sixteen point seven. [S63] The authors' explanation is their Error Depth Hypothesis: stronger models make fewer but deeper errors that resist correction — their hypothesis, not a demonstrated mechanism. [S63] And only an inference-time external verifier can help is false too, because the verifier can be moved into training: same signal, paid for once, nothing added at inference. [S47] Two claims survive, and they are the two I need. [inference] First: a check tells you something only if it can come back negative when the generator is wrong, so buy that in a fresh context, on another model, or outside any model — and say which one you did. [S3] Second: verification has a price — a second pass, a cheaper model, a deterministic check, someone's time — and if you cannot point at which one you pay, you have not paid it. [S14] That is a weaker thesis than the one on the conference posters. [speculative] It is also the one still true after the next two model releases. [speculative]

25 25. What to do with part two

On screen:

Part II · Decision

Two changes this week, and two things to stop calling verification. [speculative]

Change this week

- *Find every place your system asks a model to confirm its own output, and label it: that is not a check. [inference]*
- *Measure how often your reviewer model agrees with your author model. Above roughly 90% it is confirming, not checking — seed twenty known defects and count how many it catches. [S11] [speculative]*

Do not do this

- *Do not add a second model reading the first one's transcript and call the system verified. [S3]*
- *Do not wait for a bigger model to close the gap. [S10] [S16]*

Do this week: go through whatever you have shipped and find every point where the system asks a model to confirm its own output, and relabel it honestly — it is a formatting step, or a self-consistency step, but it is not a check. [inference] Second, if you already run a reviewer model, measure how often it agrees with your author model across a sample of real outputs. Above roughly ninety percent agreement I would read it as confirming rather than checking — it shares the author's blind spots, so its approval carries almost no information, which is the consistency dilemma arriving in your logs. [S11] [speculative] Then test it directly: seed twenty known defects and count how many come back. A reviewer that agrees constantly and catches few is grading style, not correctness. [speculative] Do not do: adding a second model that reads the first one's transcript and then describing the system as verified — the transcript is the contamination. [S3] And do not wait for the next model, because verification capability does not track generation capability. [S10] [S16] Now, what actually works.

26 26. What an external check buys

On screen:*Part III*

The positive half, with the bill and the blind spots attached.

Part three is the constructive half, and it has the strongest number in the talk in it. [speculative] I will show what closing a loop against an external check delivers, then the same result reproduced in three unrelated fields, then the three things it does not do and the two ways it costs money. [speculative] I sequence it that way deliberately: a technique presented without its boundary conditions is a sales pitch, and this room has sat through enough of those. [speculative] The boundary conditions are also where the interesting engineering is. [speculative]

27 27. The number that made me rewrite my own architecture advice

On screen:

Part III · Anchor

Swapping the backbone model left compliance unchanged; closing the loop moved it. [S5]

56.8% → 98.6%

Code-compliance, open against closed loop with an external verifier [S4] [S57]

- *The authors' attribution: the gain "does not come from a more capable model, but from closing the generation–verification loop with an external verifier" [S5].*
- *Domain: safety-critical structural design — mechanism transfers, field does not. [speculative]*
- *Blocksworld, a planning benchmark: 40% alone, 55% self-verified, 88% external. [S54]*

[S4] [S5] Loop Engineering and Self-Verification — vault report · [S57] <https://arxiv.org/abs/2608.07978>

[S54] Self-Verification Limitations — Stechly et al. <https://arxiv.org/abs/2402.08115>

A closed-loop multi-agent framework moved code-compliance from fifty-six point eight percent to ninety-eight point six percent. [S4] Same agents. [S5] The authors state the attribution themselves: the gain does not come from a more capable model, it comes from closing the generation-verification loop with an external verifier. [S5] And the result is model-invariant — swap the backbone model and compliance does not move detectably, because the guarantee lives in the verifier rather than in the generator. [S5] The domain is safety-critical structural design — which is my own field, and almost certainly nobody else's in this room, so treat me as an interested party on this number. [speculative] What transfers is the mechanism and specifically the model-invariance, because model-invariance is the signature of a guarantee that lives outside the model. [inference] What does not transfer is the field, which is why the second number is here. [speculative] Blocksworld is a planning benchmark with machine-checkable ground truth — nobody's industry either, and not engineering: one generator scored forty percent alone, fifty-five checking its own plans, eighty-eight paired with a sound external verifier. [S54] Same ordering, different discipline. [inference]

28 28. One architecture in three unrelated fields: the generator proposes, an oracle outside it returns pass or fail

On screen:

Part III · Convergence

Structural design

[S4]

56.8% → 98.6% compliance, external physics verifier

Model-invariant

Clinical diagnosis

[S7]

"clinical-diagnosis accuracy from 55.6% to 77.5%" against expert guidelines

Guideline oracle

Formal mathematics

[S8]

"a 422 percent performance improvement over the best publicly available prover LLM"

Proof-checker oracle

Same loop in all three: the generator proposes, something outside it returns pass or fail. That recurrence makes it a principle, not a stack preference. [inference]

[S4] [S7] [S8] Loop Engineering and Self-Verification Loops — vault report (2026).

Structural design: fifty-six point eight to ninety-eight point six, with a physics verifier. [S4] Clinical diagnosis: fifty-five point six to seventy-seven point five, grounded against expert clinical guidelines. [S7] Formal mathematics: a four-hundred-and-twenty-two percent improvement over the best publicly available prover model, using a proof checker as the oracle. [S8] Different domains, different oracle types, and one architecture underneath: a generator proposes, something outside it returns pass or fail without reading how the answer was produced. [inference] That recurrence across unrelated fields is what upgrades a pattern from someone's stack preference to a principle. [inference] It also tells you what to look for in your own field: not a better model, but something you already own that can say pass or fail. [speculative]

29 29. Independence costs a fresh context, not a second vendor

On screen:

Part III · Independence

Run the check on the same model in a clean context, shown the artifact and the acceptance criteria and nothing else. [S17]

- *"checker independence is a context-isolation boundary, not a model-weights boundary" [S17] — a different model family is secondary, for subjective rung-4 checks.*
- *The transcript is what contaminates it: a verifier that sees the generator's reasoning is biased toward confirming it. [S17]*
- *Rule: pass the artifact and the specification, never the chain of thought. [inference]*
- *Ceiling: same-family checkers keep shared blind spots — "independence is approximated, never absolute" [S67].*

[S17] [S67] Loop Engineering — vault concept note (2026).

Independence is not what most teams assume. [speculative] Checker independence is primarily a context-isolation boundary rather than a model-weights boundary. [S17] You do not need a second vendor for the independence that matters most of the time — the note reserves a different model family as secondary reinforcement, for subjective rung-four checks. [S17] The same model, in a fresh context window, shown only the artifact and the acceptance criteria, is where the load is carried. [S17] What destroys it is the transcript. [inference] A verifier that can read the generator's reasoning is biased toward confirming that reasoning. [S17] So the implementation rule is one line: pass the artifact and the specification, and never pass the chain of thought that produced it. [inference] That costs one extra inference call and some discipline. [inference] Now the ceiling. [speculative] The same note records that same-family checkers retain shared blind spots and can collude in a degeneration-of-thought mode, so independence is approximated, never absolute. [S67] Context isolation buys most of what is available; no configuration buys all of it. [inference]

30 30. Which rung verifies this decision?

On screen:

Part III · Ladder

1

*Deterministic**Exit code, assertion, golden output. [S18]*

2

*Rule or constraint**Linters, schema, policy engine. [S18]*

3

*Field truth**Tests, a deploy, a real outcome. [S18]*

4

*Model as judge**Rubric scoring by a model. [S18]*

5

*Human checkpoint**A person with a stake. [S18]**"70% of loops verify at levels 1–2" [S18] — where production checking actually happens.*

Five rungs. [speculative] One, deterministic: an exit code, an assertion, a golden output. [S18] Two, rule or constraint: a linter, a schema, a policy engine. [S18] Three, delayed field truth: tests, a deploy, a real customer outcome you observe later. [S18] Four, model as judge: a model scoring against a rubric. [S18] Five, a human checkpoint where someone with a stake signs. [S18] Across the measured corpus, seventy percent of production loops verify at rungs one and two. [S18] I want that to land as good news rather than as immaturity. [speculative] The boring rungs are cheap, fast, auditable, and they are where measured production checking actually happens. [S18] The question to carry out of this room is exactly four words: which rung verifies this?

31 31. Rung four is not a rung you can rest on

On screen:

*Part III · Ladder**A model grading a model covers what no program can express, and buys no independence from the failures it shares with the generator. [S11] [inference]*

What rung 4 gives you

Coverage of things no program can express — tone, relevance, whether an answer addressed the question. [speculative]

What rung 4 cannot buy

Independence from the failure it shares with the generator; that is the consistency dilemma again. [S11]

- Ordering that follows: make the check objective first, isolate the grader's context second, change model family third. [S17] [S18]
- If a decision only ever reaches rung 4, say so out loud in the design review. [inference]

Rung four buys you coverage of things no program can express: tone, relevance, whether the answer actually addressed the question. [speculative] What rung four cannot buy you is independence from the failures it shares with the generator, which is the consistency dilemma from part two arriving in implementation clothing. [S11] So the ordering is: make the check objective if you can, isolate the grader's context if you cannot, and only then consider a different model family. [S17] [S18] In that order, because each step costs more than the one before it and buys less. [inference] And if a decision in your system only ever reaches rung four, say that sentence out loud in the design review rather than writing verified on the diagram.

32 32. Grounding is not the same thing as retrieval

On screen:

Part III · Grounding

1

Parametric

The model answers from training. No evidence, no constraint. [S26]

2

Structural

Retrieval narrows the output distribution. It does not certify truth. [S26]

3

Deterministic

A computation outside the model determines the answer. "the LLM extracts, selects, and translates but never performs the truth-determining computation" [S26]

- *Layer three approaches the "mere translator" extreme the LLM-Modulo authors decline. [S55] [inference]*

Retrieval is layer two; most systems stop there. [inference]

[S26] Structural Grounding — vault concept note.

Layer one: the model answers from its training weights, with no evidence and no constraint. [S26] Layer two: retrieval — you inject documents, and the output distribution narrows toward them. [S26] That is a genuine improvement and it is where nearly every product marketed as grounded actually stops. [inference] Layer three: a computation outside the model determines the answer, and the model extracts, selects and translates but never performs the truth-determining computation itself. [S26] Now the attack on that third sentence, which is stronger than the field's most careful researchers will go. [speculative] The LLM-Modulo authors — the group behind the Blocksworld and self-verification results in part two — name two errors, and the second is treating the model as a mere translator of the problem specification. [S55] Layer three as worded sits close to that extreme, so take it as a design target rather than a settled result. [inference] Hold the layer distinction: the gap between two and three is where the next slide lives, and that one is for anybody whose job has no compiler in it.

33 33. Faithful and wrong at the same time

On screen:

Part III · Grounding

A system can quote its source perfectly and still be dangerously incorrect. [S15]

- *Worked case: an assistant retrieves a clause and reproduces it exactly — and the clause was wrong. [S15]*
- *The scoring: "it scores 90%+ on faithfulness (it matched its retrieved source exactly) and is structurally unsafe, because no output-against-context check can catch a wrong context" [S15].*
- *Faithfulness measures agreement with the retrieved document, never with reality. [inference]*
- *The question that catches it: not "did it match the source" but "who decided that was the source?" [inference]*

[S15] Loop Engineering — Practitioner Synthesis — vault report.

Picture an assistant that retrieves a regulatory clause and reproduces it word for word — and the clause is wrong. [S15] On every faithfulness metric in common use it scores above ninety percent, because faithfulness measures whether the output matched the retrieved source, and it did, exactly. [S15] The system is also unsafe, and no output-against-context check can catch it, because the check compares the answer to the context and the context is the thing that is wrong. [S15] That gap between faithful and true is where the most expensive failures live, and they are invisible to the metric most vendors will show you. [inference] So the audit question is not did it match the source. [inference] It is who decided that was the source, and what verified that decision. [inference]

34 34. Boundary one — the oracle only sees what it measures

On screen:

Part III · Limits

An external check supplies information about its own dimensions and nothing else. [S6]

- *From the same study that produced the 98.6%: "when the defect originates from the structural topology, the parameter-level closed loop is powerless" [S6].*
- *Generalised: a schema validator cannot catch a wrong schema; a test suite cannot catch a missing requirement. [inference]*
- *So the design question is coverage, not presence — "which failures does this oracle not see?" [inference]*

[S6] Loop Engineering and Self-Verification Loops — vault research report (2026).

The same study that gave me the ninety-eight point six figure states its own limit: when the defect originates from the structural topology, the parameter-level closed loop is powerless. [S6] In plain terms, the verifier checked whether the parameters were compliant, so it could say nothing about whether the overall shape was wrong. [inference] A schema validator cannot catch a wrong schema. [inference] A test suite cannot catch a requirement nobody wrote a test for. [inference] A citation registry can confirm a paper exists and cannot confirm it says what was claimed. [inference] So the design question is never does this have a verifier — it is which failures does this verifier structurally not see, and what covers those. [inference]

35 35. Boundary two — somebody has to write the oracle

On screen:

Part III · Limits

The verifier is not the free part of the design. It is most of the bill. [inference]

- *The LLM-Modulo authors: the verifier may be substantial work, and substituting a simulator means someone must write the simulator. [S55]*
- *The work is formalising domain knowledge into machine-checkable form, and it is "labor-intensive and domain-specific" [S26].*
- *Defensible, because whoever formalises a domain owns its check. [speculative]*
- *And it is hardest where it is most needed — "There is no convincing public playbook yet for retrofitting a ten-year-old codebase" [S22].*

[S55] <https://arxiv.org/html/2402.01817v2> · [S22] Harness Engineering · [S26] Structural Grounding — vault concept notes.

I have been filing oracle construction under open problems. [speculative] It is not an open problem — it is the bill. [inference] The LLM-Modulo authors put it bluntly: the verifier may be more complex than the generator, and if you substitute a simulator for a verifier, someone still has to write that simulator. [S55] The work involved is formalising domain knowledge into machine-checkable form — rules, constraints, declarative checks — and my own notes call it labor-intensive and domain-specific. [S26] It is expensive, and it is also defensible, because whoever formalises a domain owns the check for it and a generic vendor cannot cheaply copy that. [speculative] And the hardest case is the one that arrives with an existing codebase: there is no convincing public playbook yet for retrofitting a ten-year-old codebase. [S22]

36 36. Boundary three — the oracle depreciates

On screen:

Part III · Limits

Treat a verifier as "a maintenance liability that must be re-scoped at every model upgrade" [S25] rather than as a fixed asset.

- *Reported first-hand: an evaluator "became unnecessary overhead" for tasks that moved inside the generator's boundary after a model upgrade [S52].*
- *The same account deleted a whole context-reset subsystem after one upgrade [S52].*
- *Plan for it: put a re-scope of every model-judge step on the calendar for each model upgrade. [inference]*
- *The deterministic rungs depreciate far more slowly, which is another reason to prefer them. [inference]*

[S25] *Verification and Checker Design* — vault report · [S52] <https://www.anthropic.com/engineering/harness-design-long-running-apps>

Everything I am telling you to build depreciates on a model-release cadence. [inference] My own notes put it precisely: a verifier is not a fixed asset but a maintenance liability that must be re-scoped at every model upgrade. [S25] There is a first-hand engineering account of exactly this — an evaluator became unnecessary overhead once the generator absorbed the check after a model version bump, and the same team deleted an entire context-reset subsystem after a later model version made it redundant. [S52] So put the re-scope on the calendar rather than discovering it in the token bill. [inference] And notice which rungs survive: a compiler does not depreciate when a model ships. [inference] A model-judge step does. [inference] That is a third reason to spend your budget at the bottom of the ladder.

37 37. What to do when you have no oracle

On screen:

Part III · Limits

Most systems already contain a check nobody wrote down [speculative]; where none exists, the honest engineering answer is to leave that decision un-automated. [S51]

You have one, probably

- *A schema, a registry, a reconciliation, a delayed real-world outcome, a person with a stake. [S18]*
- *Report the reliability floor rather than the ceiling: pass^k, not pass@k. [S44]*

You genuinely do not

- *Then the honest answer is to not automate that decision. [S51]*
- *Design the refusal path and evaluate it, rather than shipping a confident guess. [speculative]*

[S51] *Building Effective Agents* — Anthropic. <https://www.anthropic.com/engineering/building-effective-agents>

The obvious objection to everything so far: this whole framework assumes you have an oracle, and plenty of decisions have no compiler. [speculative] First, most people underestimate what they have — a schema, a reconciliation against a second system, a registry lookup, a delayed real-world outcome you could actually observe, or a person with a stake who signs. [S18] Those are rungs one through five, and at minimum you can report the reliability floor honestly: all-k-succeed rather than at-least-one-succeeded. [S44] Then the honest engineering answer is to not automate that decision, and the vendor guidance most

quoted in this field says the same thing — find the simplest solution, which might mean not building an agentic system at all. [S51] That is an unsatisfying slide for a room that came for what to build. [speculative] It is also the one that saves the most money. [speculative]

38 38. What to do with part three

On screen:

Part III · Decision

Three writing exercises, none of which needs a new tool or a new model. [speculative]

1

Name the oracle

For each AI-assisted decision, write down what could say pass or fail without the model. [S18]

2

Isolate the context

Verifier sees artifact plus criteria; never the generator's transcript. [S17]

3

Write the blind spots down

Beside each oracle, the failures it structurally cannot see. [S6]

One: for each AI-assisted decision in your product, write down what could return pass or fail without consulting the model — that is your oracle, and if the box is empty you have found your highest-value piece of engineering. [S18] Two: wherever you already run a checking step, isolate its context so it sees the artifact and the acceptance criteria and never the transcript that produced them; that change is usually an afternoon. [speculative] Three, and this is the one teams skip: next to each oracle write the failures it structurally cannot see, because that list is your actual risk register and it is far more informative than a coverage percentage. [S6] Now, authority — which is where the expensive incidents live.

39 39. Bounded autonomy

On screen:

Part IV

The incidents are permission failures wearing new clothes.

Part four, and I want to reframe the topic before I start. [speculative] Autonomy gets discussed as an AI safety question, in language borrowed from alignment research, and that framing makes the problem look novel and slightly unreachable. [speculative] I think it is mostly a permissions question, and permissions engineering has fifty years of results behind it — so I will state the governing principle from 1975 rather than assume anybody here already carries it. [speculative] So the argument in this section is that the interesting failures are not the model doing something surprising — they are the model being able to do something nobody had bounded. [speculative] I will start with the arithmetic that explains why long chains fail, then the incident that explains why authority matters more than the arithmetic. [speculative]

40 40. Why long chains fail even with a good model

On screen:

Part IV · Arithmetic

End-to-end success is the product of per-step success, so length hurts faster than a weak step.

[S27] [inference]

$0.95^{20} \approx 36\%$

"95% per-step reliability yields only 36% success over 20 steps" [S27]

$P(\text{success}) = \prod_{i=1}^n p_i$

[S27] > > - Only two moves touch the exponent: shorten the chain, or check each seam. A better model moves only the base rate. [inference] > - Caveat: this assumes independent per-step accuracies — "the number is a slogan, not a prediction" [S28]. > > [S27] Agentic Loop Pattern — vault note · [S28] Grounding in AI-Driven Systems — vault podcast brief.

Ninety-five percent per step, twenty steps, and you are at thirty-six percent end to end. [S27] Ninety-five percent sounds excellent in a status meeting and is catastrophic in a chain. [inference] Only two moves touch the exponent: shorten the chain, or check each seam so errors cannot propagate. [inference] Buying a better model moves the base rate a little and the exponent not at all. [inference] Now the caveat, because the number gets quoted without it constantly. [speculative] It assumes the per-step accuracies are independent. [S28] My own notes are blunt about this — "the number is a slogan, not a prediction." [S28] Use it to reason about shape, never to forecast a specific system.

41 41. Real pipeline errors are not independent — they stick

On screen:

Part IV · Arithmetic

Errors in real pipelines are not independent — they stick and propagate. [S29]

- *Measured: failures "propagating semantically across steps rather than occurring independently" [S29].*
- *Recovery is not uniform either: "37.5% for test generation but 0% for patch generation" [S29].*
- *So the curve is not a clean exponential — it is a set of stages with different recovery, and some with none. [inference]*
- *Practical consequence: instrument per stage. An end-to-end success rate hides which seam is leaking. [inference]*

[S29] The 4/δ Bound — Designing Predictable LLM-Verifier Systems — vault reading note.

Real pipelines show failure stickiness: errors propagate semantically across steps rather than occurring independently, which means the independence assumption fails in the direction that hurts. [S29] And recovery rates differ sharply by stage — thirty-seven and a half percent for test generation, zero percent for patch generation in the same study. [S29] Zero. [speculative] So the true picture is not a smooth exponential decay; it is a set of stages with wildly different recovery behaviour, some of which never recover once they go wrong. [inference] Which gives you a concrete instrumentation instruction. [speculative] Measure per stage, not end to end, because an end-to-end success rate averages over exactly the difference that would tell you where to put your one check. [inference]

42 42. Nine seconds: an agent deleted a production database

On screen:

Part IV · Incident

Date

25 April 2026 [S30]

What happened

A coding agent "deleted PocketOS's entire production database and all volume-level backups in nine seconds" [S30]

How

A credential mismatch on a staging task; a token found in an unrelated file; a destructive API call with no confirmation step [S30]

Cost

A 30-hour outage for the businesses running on it [S30]

[S30] PocketOS Cursor Incident — Database Deletion in Nine Seconds — vault watch note (2026).

On the twenty-fifth of April this year, a coding agent deleted a company's entire production database and all volume-level backups in nine seconds. [S30] The sequence: it hit a credential mismatch on a staging task, found an API token sitting in an unrelated file, and issued a destructive call against an API that had no confirmation step for destructive operations, no environment isolation, and backups stored on the same volume as the primary data. [S30] Read that list again and notice how much of it has nothing to do with artificial intelligence. [speculative] A token with too much scope, no environment separation, backups on the same volume, no confirmation on a destructive operation. [inference] That is a 1975 finding with a 2026 date on it. [speculative]

43 43. Ask why the agent was able to delete the database

On screen:

Part IV · Incident

Nothing in the PocketOS deletion required the model to misbehave: an over-scoped token and an unconfirmed destructive API are enough. [S31] [inference]

- *The post-mortem's conclusion: "system prompts are not security controls" [S31] — prevention needs a deterministic mechanism outside the reasoning loop.*
- *The 1975 principle: every program and user runs with the least privilege needed for the job [S59].*
- *A token in an unrelated file carried authority over production volumes — a scoping defect, not a model defect. [inference]*

[S31] PocketOS Cursor Incident — vault watch note · [S59] The Protection of Information in Computer Systems — Saltzer and Schroeder (1975). <https://www.cs.virginia.edu/~evans/cs551/saltzer>

Everyone asks why did the agent do that. [speculative] The useful question is why was it able to. [inference] The most rigorous post-mortem of that incident states the conclusion I would put on a wall: system prompts are not security controls. [S31] The only thing that reliably prevents a catastrophically wrong action is a deterministic enforcement mechanism operating outside the model's reasoning loop, one that makes certain outcomes structurally impossible. [S31] Saltzer and Schroeder wrote the governing principle in 1975 — every program and every user operates with the least set of privileges necessary for the job. [S59] A token found in an unrelated file carried authority to delete production volumes. [inference] No prompt engineering fixes that, and no better model fixes that. [inference] A narrower credential fixes it. [inference]

44 44. Deny at the API gate, not in the prompt

On screen:

Part IV · Design

A prompt instruction is advisory. An API gateway rule is evaluated after the model, so it can refuse a call the model already emitted. [S32]

Advisory — tool output can override it [inference]

The prompt, the system message, the policy document. [inference]

Enforced — holds whatever the model emitted [S32]

The credential, the API gateway, the sandbox, the deny gate. [S32]

- *Worth adopting verbatim: "irreversible actions (production deploys, money movement, force-push) belong on a deterministic deny gate, never on a probabilistic reviewer" [S20].*

[S32] Harness Engineering — Tool Permissions — vault research report (2026). [S20] Loop Engineering — vault concept note (2026).

A prompt instruction is advisory. Anything that changes the input can override it, and for an agent reading tool output that is essentially every turn. [inference] A credential, an API gateway rule, a sandbox boundary or a deny gate is evaluated after the model and outside it, so it holds regardless of what the model produced. [inference] My source states the design rule as: put the control where the authority is exercised, not where the intent is expressed. [S32] The operational form I would adopt verbatim: irreversible actions — production deploys, money movement, force-push — belong on a deterministic deny gate, never on a probabilistic reviewer. [S20] Notice this also separates two things people conflate. [speculative] The verifier checks the work. [inference] The gate checks the action. [inference] You need both, and they are different components.

45 45. Containment over supervision — and it still leaks

On screen:

Part IV · Design

A vendor moved the control onto the operating system, and published the leak with it. [S68]

Why

"containment over supervision", because "human oversight suffers from approval fatigue" [S68]

What

OS-level sandbox — Seatbelt, bubblewrap: "reads allowed, writes inside workspace, network denied by default" [S68]

Leak

Same post: models "helpfully" escaped sandboxes or routed around restrictions [S68]

–84%

Fewer permission prompts in Claude Code, vendor-reported [S68]

[S68] How We Contain Claude — Anthropic (2026), via vault watch note.

Part four's rule needs a mechanism; here is a published one. [speculative] Anthropic moved Claude Code's control off per-action human approval and onto the operating system — Seatbelt on macOS, bubblewrap on Linux, reads allowed, writes confined to the workspace, network denied by default. [S68] Their stated reason is exactly this section's argument: containment over supervision, because human oversight suffers from approval fatigue. [S68] They report eighty-four percent fewer permission prompts in Claude Code. [S68] And the same post discloses the leak — documented cases of models escaping sandboxes or routing around restrictions. [S68] So a sandbox is a boundary you maintain rather than one you install. [inference] That is a vendor describing its own product, and I am citing it from a vault summary rather than the post itself. [S68] It is still the most concrete containment reference I can give you. [speculative]

46 46. Guardrails are layers, not boundaries

On screen:

Part IV · Design

A shipping guardrail product documents its own gaps; read those before you rely on it. [inference]

- *Disclosed in the vendor's own documentation: tool-call arguments are not inspected, and "tool results bypass input rails" [S33].*
- *So a payload hidden in a tool argument, or an injection arriving in a tool result, passes the rail by design. [inference]*
- *Stack layers that fail differently: the input rail, a separate check on tool-call arguments, and a deterministic deny gate on the irreversible action. [S33] [S20]*
- *A single guardrail described as protection is the most confident version of unprotected. [speculative]*

[S33] Guardrails and Runtime Enforcement — vault report.

Guardrail frameworks are genuinely useful and I run them. [speculative] But read the documentation, because the honest ones tell you where they stop. [speculative] One widely used framework's own docs disclose that its rails evaluate the message content field only, so tool-call arguments are not inspected — a payload sitting in a tool argument passes. [S33] And separately, tool results bypass input rails, which means malicious output coming back from a tool is trusted by default and can steer everything after it. [S33] Those are not bugs — they are documented scope. [inference] So treat a guardrail as one slice of defence with known holes, stack slices whose holes do not line up, and never let a diagram describe a single rail as protection. [inference] A single guardrail is the most confident possible version of being unprotected. [speculative]

47 47. Five autonomy levels, and each names who answers

On screen:

Part IV · Autonomy

1

In the loop

Every action approved; approver answers. [S34]

2

On the loop

Acts inside pre-authorised bounds; monitor interrupts. [S34]

3

Supervised

Routine handled; high-stakes escalate. [S34]

4

Delegated

End-to-end in mission parameters; outcomes audited. [S34]

5

Full autonomy

Own objectives; governance and audit. [S34]

"Each level implies a distinct accountability structure" [S34] — the level picks who answers. [inference]

One through five run from every action approved, to a system operating on objectives set inside broad constraints with the human role reduced to governance and audit. [S34] One caveat on those names. [speculative] They are the vault note's own synthesis, and the note says so. [S34] The paper underneath names its five levels operator, collaborator, consultant, approver and observer, and centres them on the role the user takes rather than on the system's scope, so the mapping between the two is approximate. [S34] I use the synthesis because it names the accountability holder directly, which is the decision you are actually making. [speculative] The load-bearing sentence is that each level implies a distinct accountability structure. [S34] So this is not a maturity ladder to climb. [inference] Choosing a level is choosing who is answerable when it goes wrong — a decision written down, not an emergent property of a tool default.

48 48. The human rung is a control that degrades

On screen:

Part IV · Autonomy

Bainbridge

[S58]

The operator is asked to monitor a system installed because it outperforms them

1983

The Code

[S35]

"developers need strong coding skills to effectively oversee AI agents, but those skills atrophy" — 17% drop in skill mastery

Newsletter

RSNA

[S62]

Radiologists, 15+ years: 82% → 45.5% on a purported AI's incorrect suggestion

Reported

Forty-three years apart, one finding: unaided supervision decays. [inference]

*[S58] https://ckrybus.com/static/papers/Bainbridge_1983_Automatica.pdf · [S62] <https://www.rsna.org/news/2023/may/ai-bias-may-impair-accuracy> [S35] *Agentic Coding Is a Trap — The Code* (2026).*

Whenever anyone says keep a human in the loop, they are proposing a control whose reliability is almost never measured. [inference] Lisanne Bainbridge described the problem in 1983, in the Ironies of Automation: the automation was installed because it outperforms the operator, and then the operator is asked to monitor it. [S58] That is the task humans are worst at. [inference] A 2026 newsletter reports that the skills required to supervise decay from not exercising them — a seventeen percent drop in skill mastery, with debugging declining most steeply. [S35] And in a domain of certified experts, radiologists with more than fifteen years of experience went from eighty-two percent accuracy to forty-five point five when the purported AI suggested the incorrect category. [S62] Purported is the report's own word: the suggestions were experimenter-controlled and only labelled as AI, which makes the thirty-six-and-a-half-point drop an effect of the label rather than of any model. [inference] Forty-three years apart, same finding. [inference] A human rung is a real control and it is not a free one.

49 49. The EU AI Act requires oversight that can intervene

On screen:

Part IV · Autonomy

An oversight rung that cannot actually intervene is a compliance defect. [S36]

- *Under the EU AI Act, oversight must be meaningful: "a HITL that rubber-stamps AI decisions does not satisfy Article 14(4)(b)" [S36].*
- *A human making the final call does not remove the high-risk classification. [S36]*
- *Uncomfortable corollary: automation bias becomes a compliance risk rather than a UX inconvenience. [inference]*

- *Design consequence: measure intervention rate and intervention quality, or you cannot show the rung works. [inference]*

[S36] EU AI Act Annex III — Human in Loop Not Sufficient Oversight — vault watch note (2026).

Under the EU AI Act, oversight has to be meaningful — a human-in-the-loop arrangement that rubber-stamps decisions does not satisfy Article 14, paragraph 4b. [S36] Read the second bullet carefully, because it inverts the usual assumption. [speculative] Putting a human at the end does not remove the high-risk classification; the human is a condition for satisfying the oversight requirement rather than an exemption from it. [S36] Which means automation bias — the thing the radiology data measures — is a compliance risk and not merely a design annoyance. [inference] The practical consequence for anyone deploying this: if your rung five is a person clicking approve, you need to be able to show intervention rate and intervention quality. [inference] If you cannot measure it, you cannot demonstrate it.

50 50. Bounded autonomy can relocate blame instead of risk

On screen:

Part IV · Counterargument

The strongest objection to my own recommendation, and I do not have a clean answer to it. [speculative]

- *The named human becomes a liability sponge — "absorbing blame for a system they can't actually override" [S37].*
- *We cite the model's error rate as proof it cannot certify itself, then install a human whose error rate nobody measures. [inference]*
- *Partial answer: pick the autonomy level so the human's task stays inside what a human can actually do, and measure that. [inference]*
- *Honest answer: if the person cannot realistically override the system, the level is set wrong. [speculative]*

[S37] 2026-07-22 Digest — vault journal note.

When you place a named human at the boundary, you may not be reducing risk — you may be relocating liability. [inference] The term for the role is the moral crumple zone: the nearest human absorbs blame for a system they cannot actually override, while the technology's reputation stays intact. [S37] Notice the inconsistency in my own argument. [speculative] I spent twenty minutes citing measured model error rates as proof that a model cannot certify itself, and then I proposed a human oracle whose error rate nobody in

this room measures. [inference] Partial answer: choose the autonomy level so the human's task remains something a human can actually perform, and then measure whether they perform it. [inference] Honest answer: if the person cannot realistically override the system, you have set the level wrong and the sign-off is decoration. [speculative]

51 51. What to do with part four

On screen:

Part IV · Decision

Three permission changes to make, and three habits that only look like controls. [speculative]

Do this

- *Enumerate irreversible actions; put each behind a deterministic gate. [S20]*
- *Scope every agent credential to the one task it serves. [S59]*
- *Write down the autonomy level per feature, with the name of who answers. [S34]*

Stop doing this

- *Relying on a system prompt as a control. [S31]*
- *Describing a single guardrail as protection. [S33]*
- *Counting an approval click as oversight without measuring it. [S36]*

Start: enumerate every irreversible action your system can reach and put each one behind a deterministic gate, and that list is shorter than most teams expect. [S20] Enumerating it is an afternoon. [speculative] Scope every agent credential to exactly one task, which is the direct fix for the incident I described. [S59] And write down the autonomy level for each feature together with the name of the person who answers for it. [S34] Stop: relying on a system prompt as a control, because it is not one. [S31] Stop describing a single guardrail as protection when its own documentation lists what it does not inspect. [S33] And stop counting an approval click as oversight unless you are measuring whether anyone ever declines. [S36] Now the money.

52 52. Cost, capacity

On screen:

Part V

Where the bill lands, where the bottleneck moves, and what to do next week.

Part five, the shortest and most concrete. [speculative] Three things: what a loop closed on an external check costs to run, where the constraint moves once you adopt it, and one recommendation for each of the three jobs in this room. [speculative] I kept the cost material to the end deliberately: a cost slide before an architecture argument turns into a procurement conversation. [speculative] But no CTO should leave without numbers, and no engineer should propose an architecture without knowing what it bills. [speculative]

53 53. The bill for closing a loop on an external check: more agents, more iterations, a memory layer

On screen:

Part V · Cost

Several agents, not one

"Multi-agent systems use approximately 15× more tokens than chats" [S38]

More loop iterations

"unlike SC, Reflexion can become worse with more compute budget" [S39]

A memory layer between turns

Break-even spans turn 0 to turn 356; "no system wins on both cost and accuracy" [S40]

Every mechanism here has a price. Quote it in the design review, or read it in the invoice. [speculative]

[S38] [S39] [S40] Agentic Orchestration Patterns; Reflexion Cost-Efficiency; Agent Memory Architecture — vault notes.

First: multi-agent orchestration costs roughly fifteen times the tokens of a straight chat. [S38] That buys real accuracy in some settings, and it is a fifteen-times bill in all of them. [inference] Second, the counter-intuitive one: iteration is not monotone. [speculative] Unlike self-consistency, a reflection loop can get worse with more compute budget — spending more can lower accuracy. [S39] Third, memory layers, which every team adds by reflex: the break-even against simply resending the transcript ranges from turn zero to turn three hundred and fifty-six depending on the system, and no system wins on both cost and accuracy. [S40] So memory is a conditional optimisation with a precondition to check rather than assume. [inference]

54 54. Price has become a design variable

On screen:

Part V · Cost

The spread between model tiers is now large enough to be an architecture decision. [inference]

- *Reported: a "25x input-price gap" between top-tier and cheap models [S49].*
- *Reported alongside it: an enterprise cancelling premium coding-agent licences for cost certainty [S49].*
- *So tier routing is a real lever: generate expensive, check cheap or with a program. [inference]*
- *Its stated limit: cheap-model routing to hard tasks "amplifies the probability of silent failures and verification debt" [S49].*
- *Another argument for deterministic rungs: a compiler run costs no tokens. [inference]*

[S49] The AI Cost Reckoning — Right-Sizing Model Spend 2026 — vault watch note (2026).

There is roughly a twenty-five times gap in input pricing between top-tier and cheap models. [S49] And the reported behaviour that follows is instructive — an enterprise cancelling premium coding-agent licences and redirecting engineers to a cheaper tool for cost certainty. [S49] Which gives you a lever most teams are not pulling: run the expensive model to generate, and run a cheap model or, better, a plain program to check. [inference] The checking step is usually the smaller, more structured task, and it is exactly the one you can push down a tier or out of the model entirely. [inference] And the price, because I promised to quote one every time: the same source says cheap-model routing to hard tasks amplifies the probability of silent failures and verification debt. [S49] So route the check down a tier, not the hard generation. [inference] Which is one more argument for the bottom of the verification ladder: a compiler run has no token cost at all. [inference]

55 55. Generation got cheap; review and integration did not

On screen:

Part V · Capacity

"AI-written code still needs to be reviewed, integrated, tested, and released by humans" — so the constraint moved downstream [S65].

+180% → +30%

Coding activity against releases shipped, 100k+ developers [S65]

+15% / -7%

Feature-branch against main-branch throughput, CircleCI, 28.7M workflows [S66]

- *76% more pull requests requiring review once agents ramped: "organizational absorption became the constraint" [S41].*
- *Correction: DORA's 2024 throughput penalty was reversed in 2025; only the correlational stability relationship survived. [S64] [S42]*

[S41] Harness Engineering — CI · [S42] [S64] [S65] [S66] DORA Metrics Limitations — vault notes.

Start with the cleanest measurement: it observes behaviour rather than asking people. [speculative] Across more than a hundred thousand GitHub developers and three generations of tooling, coding activity rose about a hundred and eighty percent and commit-level output roughly tripled — and the same developers shipped only about thirty percent more releases. [S65] The authors give the mechanism in one sentence: AI-written code still needs to be reviewed, integrated, tested, and released by humans. [S65] Generation got cheap; understanding, validating, integrating and owning the change did not. [inference] Second measurement, same shape: CircleCI telemetry over twenty-eight point seven million workflows shows feature-branch throughput up fifteen percent while main-branch throughput fell seven percent. [S66] That is the displacement visible inside the pipeline — acceleration where work is generated, contraction where it has to be integrated. [inference] Third, one company's own account: seventy-six percent more pull requests requiring review once agents ramped, and organisational absorption became the constraint. [S41] Now a correction I owe you. [speculative] DORA 2024 reported a throughput penalty from AI adoption; the 2025 report reversed that finding, and only the negative relationship with delivery stability persisted. [S64] So do not repeat the throughput number — its own authors withdrew it. [inference] And treat the rest as correlational: DORA is a repeated cross-section rather than a panel, and the verification-bottleneck framing usually quoted alongside it comes from a vendor with a commercial interest in that conclusion. [S42]

56 56. Instrument, then evaluate, then gate

On screen:

Part V · Practice

1

Instrument

"capture traces with cost metadata before writing a single custom eval" [S43]

2

Evaluate

Write the evals the traces asked for; run them in dev and prod alike. [S43]

3

Gate

A gate that blocks a merge "is worth ten dashboards nobody watches" [S43] — it turns quality into a precondition.

Existing monitoring cannot see "a confidently wrong answer that passed every infra check" [S43].

[S43] Harness Engineering — Evaluation and Observability Loops — vault report.

One: instrument first. [speculative] Capture traces with cost metadata before you write a single custom evaluation, because the traces tell you which evaluations need to exist — and the ones you would have guessed are usually not those. [S43] Two: then write those evaluations, and run them identically in development and in production, so a regression means the same thing in both. [S43] Three: put the result on a gate that blocks a merge, because a gate is worth ten dashboards nobody watches — it converts quality from a report into a precondition. [S43] And the reason your existing monitoring will not cover this: a confidently wrong answer that passed every infrastructure check is invisible at the infrastructure layer. [S43] Latency was fine. [inference] Cost was fine. [inference] The answer was wrong.

57 57. If you write code

On screen:

Part V · Monday

Three changes that are configuration rather than architecture, and three things to say no to. [speculative]

Adopt this week

- *Isolate your verifier's context — artifact and criteria only, never the transcript. [S17]*
- *Put irreversible actions behind a deterministic gate. [S20]*
- *Turn on tracing with cost metadata before writing evals. [S43]*

Refuse this week

- *A second model reading the first one's output, labelled verification. [S3]*

- *A twenty-step chain with one check at the end. [S27]*
- *An agent credential scoped wider than its one task. [S59]*

Adopt: isolate your verifier's context so it sees the artifact and the acceptance criteria and nothing else — one of the cheapest reliability improvements available, and it is a configuration change. [S17] Put irreversible actions behind a deterministic gate; the list of those actions is shorter than you think. [S20] Turn on tracing with cost metadata now, before you write any evaluations, because the traces will tell you which ones to write. [S43] Refuse: a second model reading the first model's output described as verification, because that is generation wearing a different prompt. [S3] Refuse a twenty-step chain with a single check at the end — put the check on the seam that has the worst recovery rate. [S27] And refuse any agent credential scoped wider than the single task it was issued for. [S59]

58 58. If your job is to judge what these systems produce

On screen:

Part V · Monday

1

Which rung?

Ask which of the five verification levels checked this decision. Rung 4 alone is not verification. [S18]

2

Faithful or true?

Ask who selected the source, and what verified that selection. [S15]

3

How often, not how well?

Ask for all-k-succeed, not best-of-k. Demand the reliability floor. [S44]

None of these three questions needs a compiler, a benchmark or a data-science team to ask. [speculative]

Three questions, none of which requires you to read a model card. [speculative] One: which of the five rungs verified this decision — and if the answer is a model scored it, that is rung four and rung four alone is not verification. [S18] Two: was the answer faithful, or was it true? [speculative] Ask who selected the source document and what verified that selection, because a perfectly faithful quotation of the wrong

source scores above ninety percent on the metric you will be shown. [S15] Three: how often rather than how well — ask for the rate at which all attempts succeed, not the best of several attempts. [S44] Three questions, no compiler required.

59 59. If you decide budgets

On screen:

Part V · Monday

1

Fund the oracle

Formalising domain knowledge is the expensive part and the defensible one. [S55] [S26]

2

Fund absorption

Review capacity is the constraint that showed up first in practice. [S41] [S65]

3

Budget the depreciation

Re-scope every model-judge step at each model upgrade. [S25] [S52]

And set the autonomy level explicitly per feature, with a named person against each. [S34]

Three line items nobody puts in the plan. [speculative] One: fund the oracle. [speculative] Formalising your domain knowledge into machine-checkable form is the expensive part, and it is also the part a generic vendor cannot copy — the verifier may be more work than the generator, and someone has to write it. [S55] [S26] Two: fund absorption. [speculative] Review capacity was the constraint that appeared first in every account I could find, and buying more generation without buying more absorption converts directly into review debt. [S41] [S65] Three: budget the depreciation — every model-judge step needs re-scoping at each model upgrade, so put it in the maintenance plan rather than discovering it in the bill. [S25] [S52] And set the autonomy level explicitly per feature, with a named person against each, because that decision is being made either way. [S34]

60 60. What I do not know

On screen:

Honest limits

Retrofit

[S22]

"There is no convincing public playbook yet for retrofitting a ten-year-old codebase."

Oracle coverage

[S6]

No general theory of which failures an oracle cannot see.

Loop drift

[S19]

No controlled study of single-loop degradation over many cycles.

The theory's bite

[S13]

The information-theoretic bound needs operationalisation before it constrains anything.

Adoption effects

[S64]

DORA reversed its 2024 throughput finding; the rest is correlational.

Retrofit: there is no convincing public playbook for retrofitting a ten-year-old codebase, and every success story I can cite is greenfield — so if you are adding this to a system that already exists, you are past the published evidence. [S22] Oracle coverage: there is no general theory of which failures a given oracle cannot see; it is judgement, system by system, every time. [S6] Loop drift: nobody has run a controlled study of how a single unattended loop degrades over many cycles, so anyone running one for weeks is operating past the evidence. [S19] The theory: the information-theoretic bound needs independent operationalisation before it constrains anything. [S13] And the adoption figures are unstable as well as correlational — DORA reversed its own throughput finding between the 2024 and 2025 reports, so anything I tell you about adoption effects on delivery has roughly a one-year half-life. [S64]

61 61. Takeaways

On screen:

Where this went

1

*Move truth out**Reliability came from closing the loop on an external verifier, not a better model. [S5]*

2

*Buy independence**Context isolation buys most of it; the transcript is the contamination. [S17]*

3

*Name the rung**Which of five levels verified this? Rung 4 alone is not verification. [S18]*

4

*Authority last**Irreversible actions on a deterministic gate; autonomy level chosen, with a name. [S20] [S34]*

One: move the truth-determining computation out — the study's own authors attribute the jump from fifty-six to ninety-eight percent to closing the loop on an external verifier rather than to a more capable model, and it held across backbone models. [S5] Two: buy independence, and buy it cheaply — context isolation gets you most of the way, and the generator's transcript is the contamination to keep out. [S17] Three: name the rung, every time, and treat a model-as-judge step as coverage rather than as certification. [S18] Four: add authority last — irreversible actions on a deterministic gate, and an autonomy level chosen deliberately with a name attached to it. [S20] [S34] If you remember one sentence: an agent's own done carries no information, so buy the check outside it.

62 62. What follows from an external check that can fail the work, and a deny gate at the credential

On screen:

*Closing**If reliability comes from those two mechanisms, not from prompting, three things change. [inference]*

- *The interesting engineering moved from the prompt to the check — a specification problem your team knows. [inference]*

- *Model choice matters less than harness and oracle design: a sixteen-point swing came from the wrapper. [S23]*
- *Verification always costs something [S14] — a second inference pass, another model, a deterministic check, or someone's time — so the durable job is owning that check. [speculative]*
- *The honest test of any AI system: what could tell you it was wrong, without asking it? [inference]*

[S14] Loop Engineering · [S23] Is Grep All You Need?

First, the interesting engineering moved from the prompt to the check, and that is good news: designing a check is a specification problem, and your team has done that work for years. [inference] Second, model choice matters less than harness and oracle design, which is exactly why a sixteen-point accuracy swing came from the wrapper rather than the weights. [S23] Third, verification is always paid for, and you can name the currency: a second inference pass, another model, a deterministic check, or someone's review time. [S14] So my bet is that the durable job is owning that check — designing it, maintaining it through model upgrades, standing behind its verdicts. [speculative] Which reduces the whole hour to one question you can ask of any system, including the one that generated this deck. [speculative] What could tell you it was wrong, without asking it? [inference] Thank you. [speculative] Questions.

63 63. References

On screen:

[S1] Loop Engineering and Self-Verification Loops — Practitioner Synthesis — vault research report (2026). [S2] Loop Engineering and Self-Verification Loops — Practitioner Synthesis — vault research report (2026). [S3] Loop Engineering and Self-Verification Loops — Comprehensive Report — vault research report (2026). [S4] Loop Engineering and Self-Verification Loops — Comprehensive Report — vault research report (2026). [S5] Loop Engineering and Self-Verification Loops — Comprehensive Report — vault research report (2026). [S6] Loop Engineering and Self-Verification Loops — Comprehensive Report — vault research report (2026). [S7] Loop Engineering and Self-Verification Loops — Comprehensive Report — vault research report (2026). [S8] Loop Engineering and Self-Verification Loops — Comprehensive Report — vault research report (2026). [S9] Loop Engineering and Self-Verification Loops — Comprehensive Report — vault research report (2026). [S10] Loop Engineering and Self-Verification Loops — Comprehensive Report — vault research report (2026). [S11] Loop Engineering and Self-Verification Loops — Comprehensive Report — vault research report (2026). [S12] Loop Engineering and Self-Verification Loops — Comprehensive Report — vault research report (2026). [S13] Loop Engineering and Self-Verification Loops — Comprehensive

Report — vault research report (2026). [S14] Loop Engineering and Self-Verification Loops — Practitioner Synthesis — vault research report (2026). [S15] Loop Engineering — Practitioner Synthesis — vault report. [S16] Loop Engineering and Self-Verification Loops — Practitioner Synthesis — vault research report (2026). [S17] Loop Engineering — vault concept note (2026). [S18] Loop Engineering — vault concept note (2026). [S19] Loop Engineering — vault concept note (2026). [S20] Loop Engineering — vault concept note (2026). [S21] Harness Engineering — vault concept note (2026).

Everything on a slide tonight traces to one of these sixty-eight entries, and the numbering matches the handout, so photograph these three slides and follow up on anything. [speculative] I cited the vault notes rather than only the papers because the notes carry the caveats I quoted, and the caveats are the part that gets dropped in transmission. [speculative] Where a figure comes from a preprint I have said preprint, and where it comes from a vendor I have said vendor. [speculative]

64 64. References (continued)

On screen:

[S22] Harness Engineering — vault concept note (2026). [S23] Is Grep All You Need? How Agent Harnesses Reshape Agentic Search — vault reading note (2026). [S24] Is Grep All You Need? How Agent Harnesses Reshape Agentic Search — vault reading note (2026). [S25] 2026-07-21 Loop Engineering — Verification and Checker Design — vault research report (2026). [S26] Structural Grounding — vault concept note (2026). [S27] Agentic Loop Pattern — vault Zettelkasten note (2026). [S28] Grounding in AI-Driven Systems — Composition, Agentic Patterns, End-to-End Guarantees — vault podcast brief (2026). [S29] The 4/8 Bound — Designing Predictable LLM-Verifier Systems — vault reading note (2026). [S30] PocketOS Cursor Incident — Database Deletion in Nine Seconds — vault watch note (2026). [S31] PocketOS Cursor Incident — Database Deletion in Nine Seconds — vault watch note (2026). [S32] Harness Engineering — Tool Permissions and Sandboxing — vault research report (2026). [S33] Guardrails and Runtime Enforcement — vault report. [S34] Bounded Autonomy in AI — vault concept note (2026). [S35] Agentic Coding Is a Trap — The Code newsletter (2026-05-15). [S36] EU AI Act Annex III — Human in Loop Not Sufficient Oversight — vault watch note (2026). [S37] 2026-07-22 Digest — vault journal note (2026). [S38] Agentic Orchestration Patterns — vault concept note (2026). [S39] Reflexion Cost-Efficiency Tradeoffs, Failure Modes, Iteration Safety Bounds — vault concept note (2026). [S40] Harness Engineering — Agent Memory Architecture — vault research report (2026). [S41] Harness Engineering — Workflow and CI Orchestration — vault research report (2026). [S42] DORA Metrics Limitations in AI-Assisted Development — vault Zettelkasten note (2026).

Twenty-three and twenty-four are the harness-versus-model measurement I leaned on twice — a single 2026 study, so read it before repeating my sixteen-point figure anywhere. [speculative] Thirty and thirty-one are the PocketOS post-mortem, worth reading in full because the sequence of small permission decisions is more instructive than the headline. [speculative] Thirty-five, thirty-six and thirty-seven are the three that make the human rung uncomfortable — skill atrophy, the Article 14 requirement, the liability framing. [speculative] Forty-two I corrected on the slide rather than dropped: DORA reversed its own 2024 throughput finding in 2025, sixty-four carries the reversal, and sixty-five and sixty-six are the independent telemetry studies that replaced it. [speculative]

65 65. References (final)

On screen:

[S43] *Harness Engineering — Evaluation and Observability Loops* — vault research report (2026). [S44] *Unit Testing* — vault concept note (2026). [S45] *Building AI Systems — predecessor-talk glossary* — vault talk note (2026). [S46] *Claude 5 self-verification and the orchestration premium* — vault tutor card (2026-07-29). [S47] *2026-07-18 Digest* — vault journal note (2026). [S48] *2026-07-17 Digest* — vault journal note (2026). [S49] *The AI Cost Reckoning — Right-Sizing Model Spend 2026* — vault watch note (2026). [S50] *Context Engineering* — vault Zettelkasten note (2026). [S51] *Building Effective Agents* — Anthropic. <https://www.anthropic.com/engineering/building-effective-agents> [S52] *Harness design for long-running application development* — Anthropic. <https://www.anthropic.com/engineering/harness-design-long-running-apps> [S53] *Large Language Models Cannot Self-Correct Reasoning Yet* — Huang et al. (ICLR 2024). <https://arxiv.org/abs/2310.01798> [S54] *On the Self-Verification Limitations of LLMs on Reasoning and Planning Tasks* — Stechly, Valmeekam, Kambhampati (ICLR 2025). <https://arxiv.org/abs/2402.08115> [S55] *LLMs Can't Plan, But Can Help Planning in LLM-Modulo Frameworks* — Kambhampati et al. <https://arxiv.org/html/2402.01817v2> [S56] *Towards a Science of AI Agent Reliability* — Rabanser, Kapoor et al. (ICML 2026). <https://arxiv.org/abs/2602.16666> [S57] *Verification-Driven Closed-Loop Multi-Agent LLM Framework for Code-Compliant Structural Design* — arXiv (2026). <https://arxiv.org/abs/2608.07978> [S58] *Ironies of Automation* — Bainbridge, Automatica (1983). https://ckrybus.com/static/papers/Bainbridge_1983_Automatica.pdf [S59] *The Protection of Information in Computer Systems* — Saltzer and Schroeder (1975). <https://www.cs.virginia.edu/~evans/cs551/saltzer> [S60] *Self-Correction as Feedback Control: Error Dynamics, Stability Thresholds, and Prompt Interventions in LLMs* — arXiv preprint (2026). <https://arxiv.org/html/2604.22273v2> [S61] *Concepts — tokens* — OpenAI documentation. <https://platform.openai.com/docs/concepts> [S62] *AI Bias May Impair Radiologist Accuracy* — RSNA (2023). <https://www.rsna.org/news/2023/may/ai-bias-may-impair-accuracy> [S63] *Decomposing LLM Self-Correction: The Accuracy-Correction Paradox and*

Error Depth Hypothesis — arXiv preprint (2026). <https://arxiv.org/abs/2601.00828> [S64] DORA Metrics Limitations in AI-Assisted Development — vault Zettelkasten note (2026), on the 2025 throughput reversal. [S65] DORA Metrics Limitations in AI-Assisted Development — vault note, citing Demirer et al. (2026), 100,000+ GitHub developers. [S66] DORA Metrics Limitations in AI-Assisted Development — vault note, citing CircleCI, State of Software Delivery (2026). [S67] Loop Engineering — vault concept note (2026), on the ceiling of checker independence. [S68] Anthropic — How We Contain Claude Across Products — vault watch note (2026-05-27), summarising <https://www.anthropic.com/engineering/how-we-contain-claude> [S69] Harness Engineering — vault concept note (2026), on the Terminal-Bench 2.1 model-versus-harness spread.

Fifty-three and fifty-four are the two negative results — read the scope sections rather than the abstracts, because Huang is specifically about intrinsic self-correction on reasoning tasks and that scope is what people drop when they quote it. [speculative] And on fifty-four: I quote its per-domain figures from published summaries of the paper, not from the PDF, which I did not re-read for tonight. [speculative] Fifty-five is the position paper that rejects both extremes, and it is the one that made me soften my own thesis. [speculative] Fifty-seven is the primary source for my headline number, and it is a preprint in a domain that is not yours. [speculative] Fifty-eight and fifty-nine are 1983 and 1975, which is the point. [speculative] Sixty and sixty-three make self-correction conditional rather than forbidden; sixty-three is one maths benchmark and three models, so read its ratio as a measurement, not a curve. [speculative] Sixty-nine is a leaderboard snapshot, and leaderboards move. [speculative] Sixty-eight is the containment post, and it is a vendor writing about its own sandbox. [speculative] Thank you — I will take questions now. [speculative]